3DF

Reactive Datalog for Datomic

by Nikolas Göbel

“We want to make reactive systems that don't poll. And we want those systems to get a consistent view of the world.”

– *Rich Hickey, “Deconstructing the Database”*

Nikolas Göbel

niko@clockworks.io

in collaboration with:

Frank McSherry (ETH)

David Bach (Clockworks)

Malte Sandstede (Clockworks)

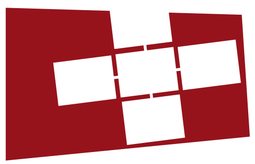

Systems Group

Reactive Systems?

- spreadsheets, chats, stock tickers, real-time dashboards
- alerting and rule engines
- collaboration tools
- stream processors

Reactive Principles

- data coordinates code
- change propagates
- a relationship, once established, is maintained
- results always reflect the current state of the domain

fx | =C4+C5

	A	B	C	D	E	F	G
1							
2							
3							
4		A	12				
5		B	3				
6		**(+ A B)**	**15**				
7							
8							
9							
10							
11							
12							
13							
14							

Data coordinates code, not the other way around.

~~evaluating~~

subscribing to a query

~~re-computing from scratch~~

incrementally updating a previous result

<demo>

3DF (github.com/comnik/clj-3df)

- reactive Datalog queries
- feeds from durable sources (Datomic, Datahike, Kafka)
- scales from browsers to clusters

github.com/comnik/clj-3df github.com/comnik/declarative-dataflow

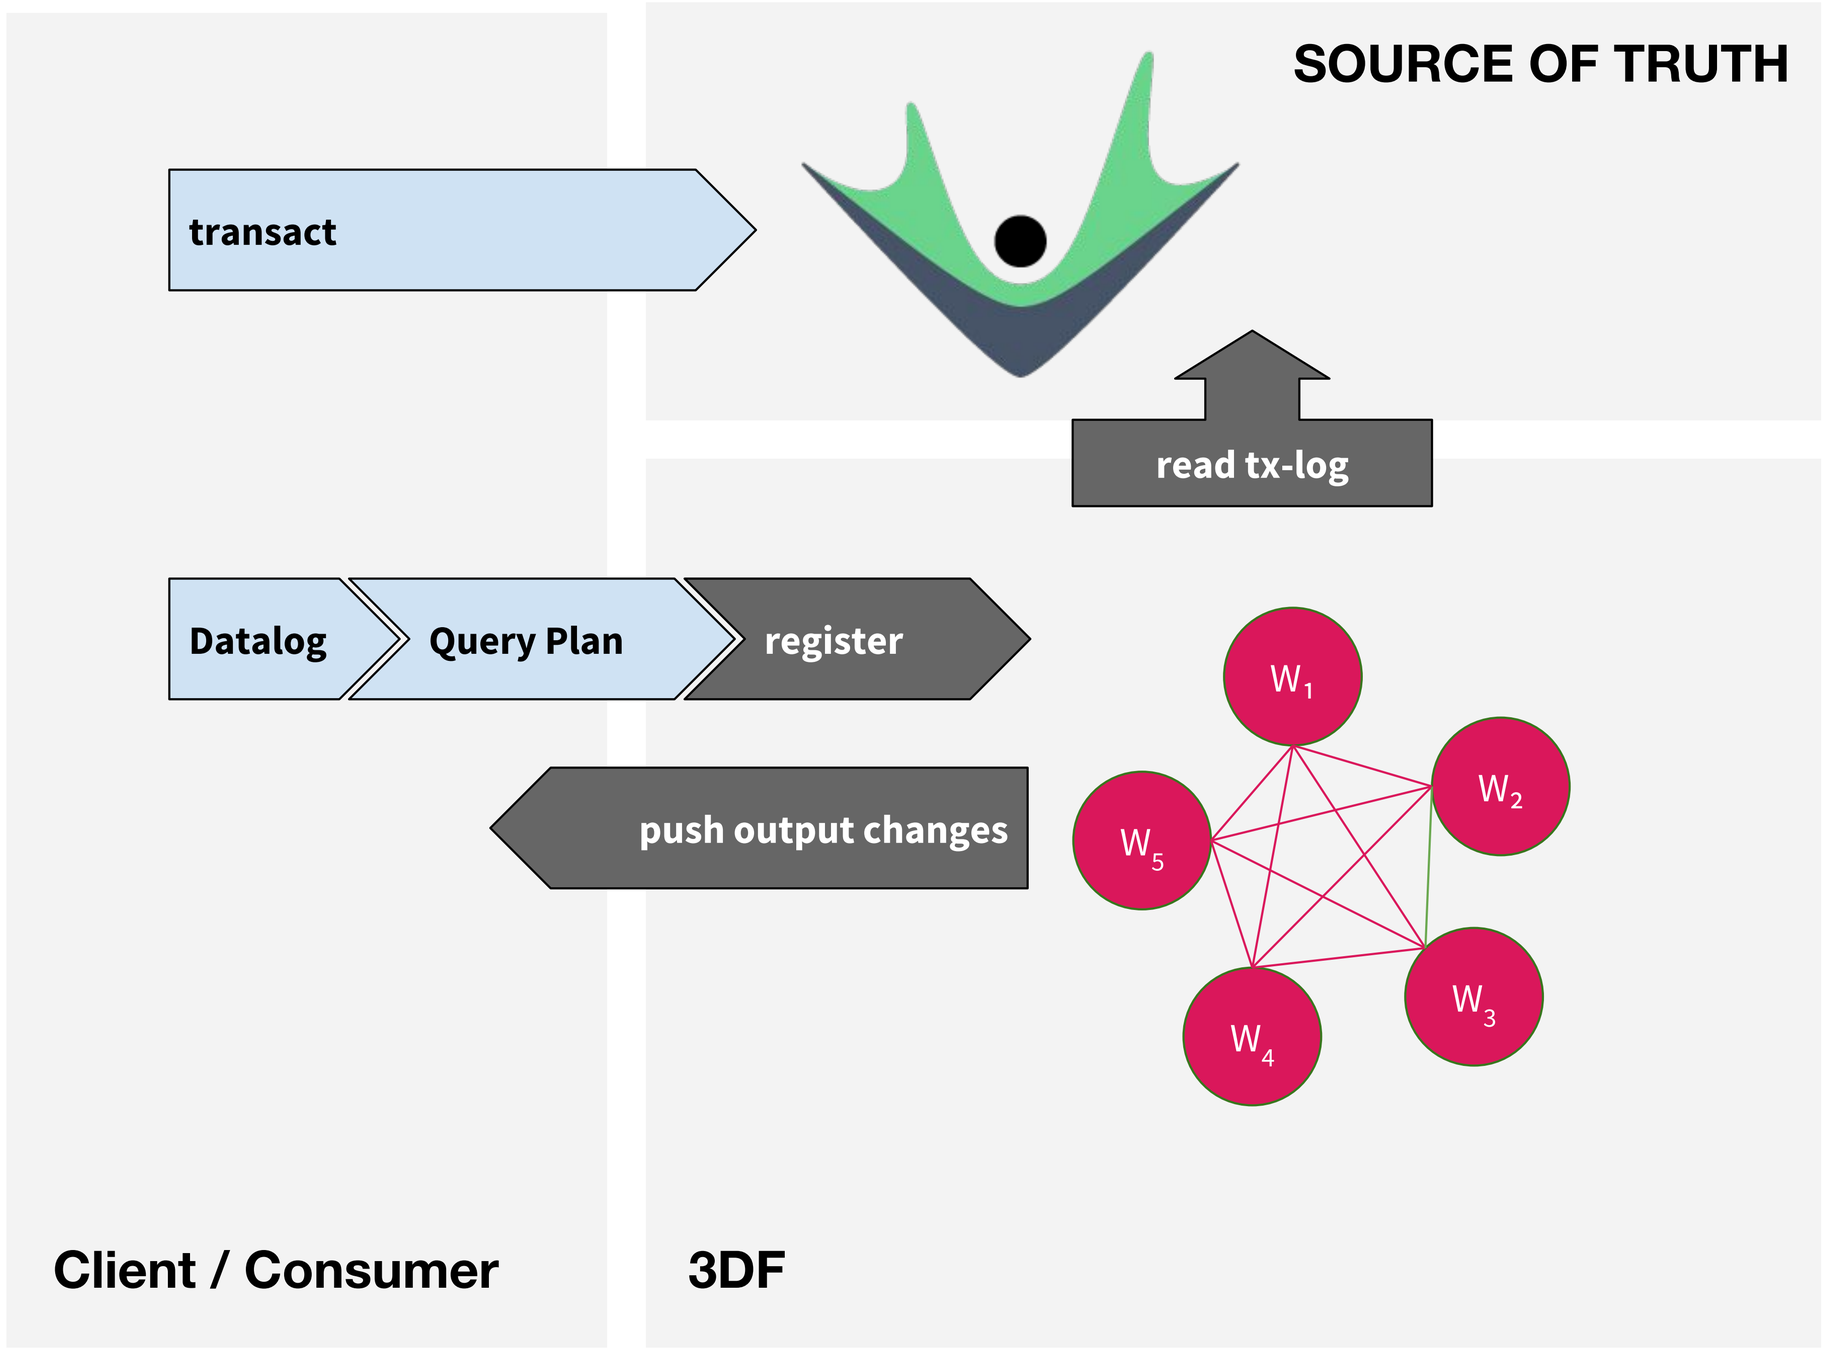
SOURCE OF TRUTH
transact
read tx-log
Datalog
Query Plan
register
push output changes
W1
W2
W3
W4
W5
Client / Consumer
3DF

Making Reactive Practical (I)

data coordinates code
work like a spreadsheet…

declarative
…without writing joins by hand

Making Reactive Practical (II)

transparent
notion of change shouldn’t leak into Datalog…

efficient
…but update results incrementally

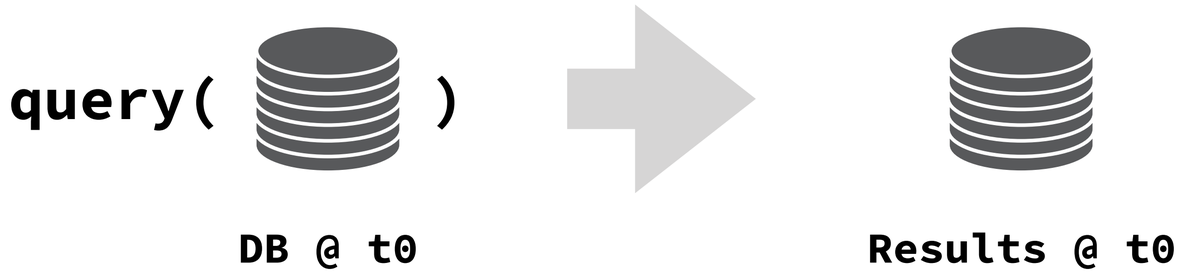

Computing with Snapshots

```
[:find ?device
 :where
 [?device :settings/speed ?target]
 [?device :device/speed ?speed]
 [(< ?speed ?target)]]
```

✓ transparent

✓ declarative

Computing with Differences

```
[:find ?device
 :in $ ?tx
 :where
 [$ ?device :settings/speed ?target]
 [?tx ?device :device/speed ?speed]
 [(< ?speed ?target)]]
```

incremental

Differential Dataflow

github.com/TimelyDataflow/differential-dataflow

“a data-parallel programming framework designed to quickly respond to arbitrary changes in input”

Time-varying Collections

`:edge`

(a b)	+1	**t0**
(c d)	+1	**t0**
(c d)	-1	**t1**
(c e)	+1	**t1**

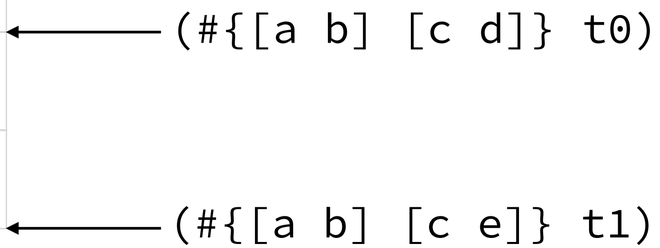

Partially-ordered Times

```
(bfs :edge a)
```

(a b) | +1 | (t0 0)		
(c d) | +1 | (t0 0)		
(b c) | +1 | (t1 0)	(a c) | +1 | (t1 1)	(a d) | +1 | (t1 2)
(c d) | -1 | (t2 0)	-	(a d) | -1 | (t2 2)

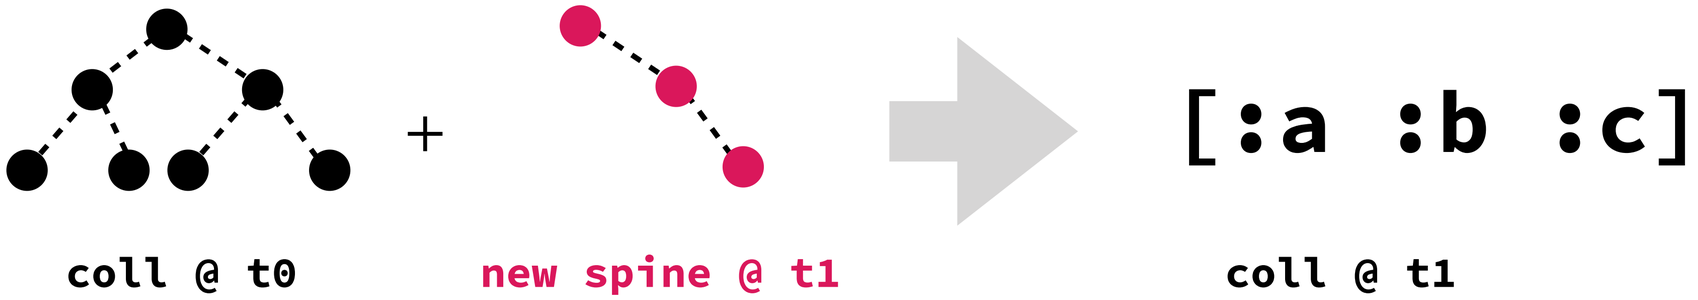

Computing with Differences

Reactive Made Practical

- ✓ data coordinates code
- ✓ declarative joins, filters, and aggregations
- ✓ write Datalog without thinking about change
- ✓ all queries update incrementally / only propagate change

~~rule the world~~

Everybody wants to be a database peer.

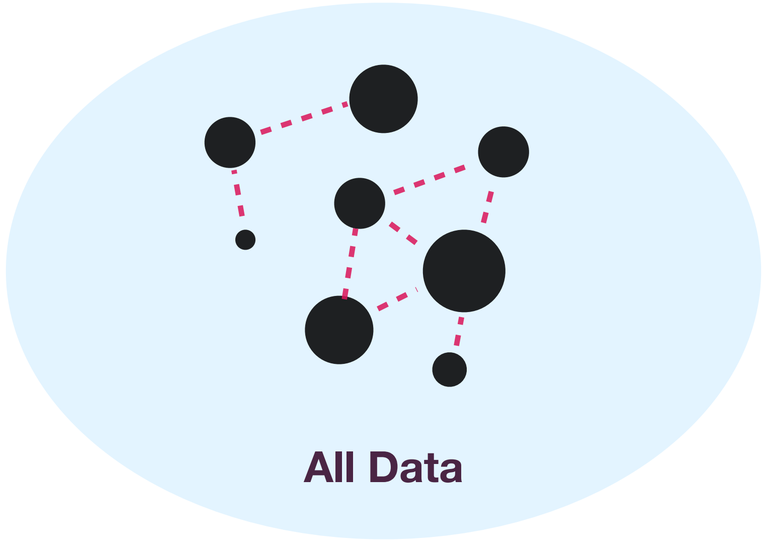
All Data

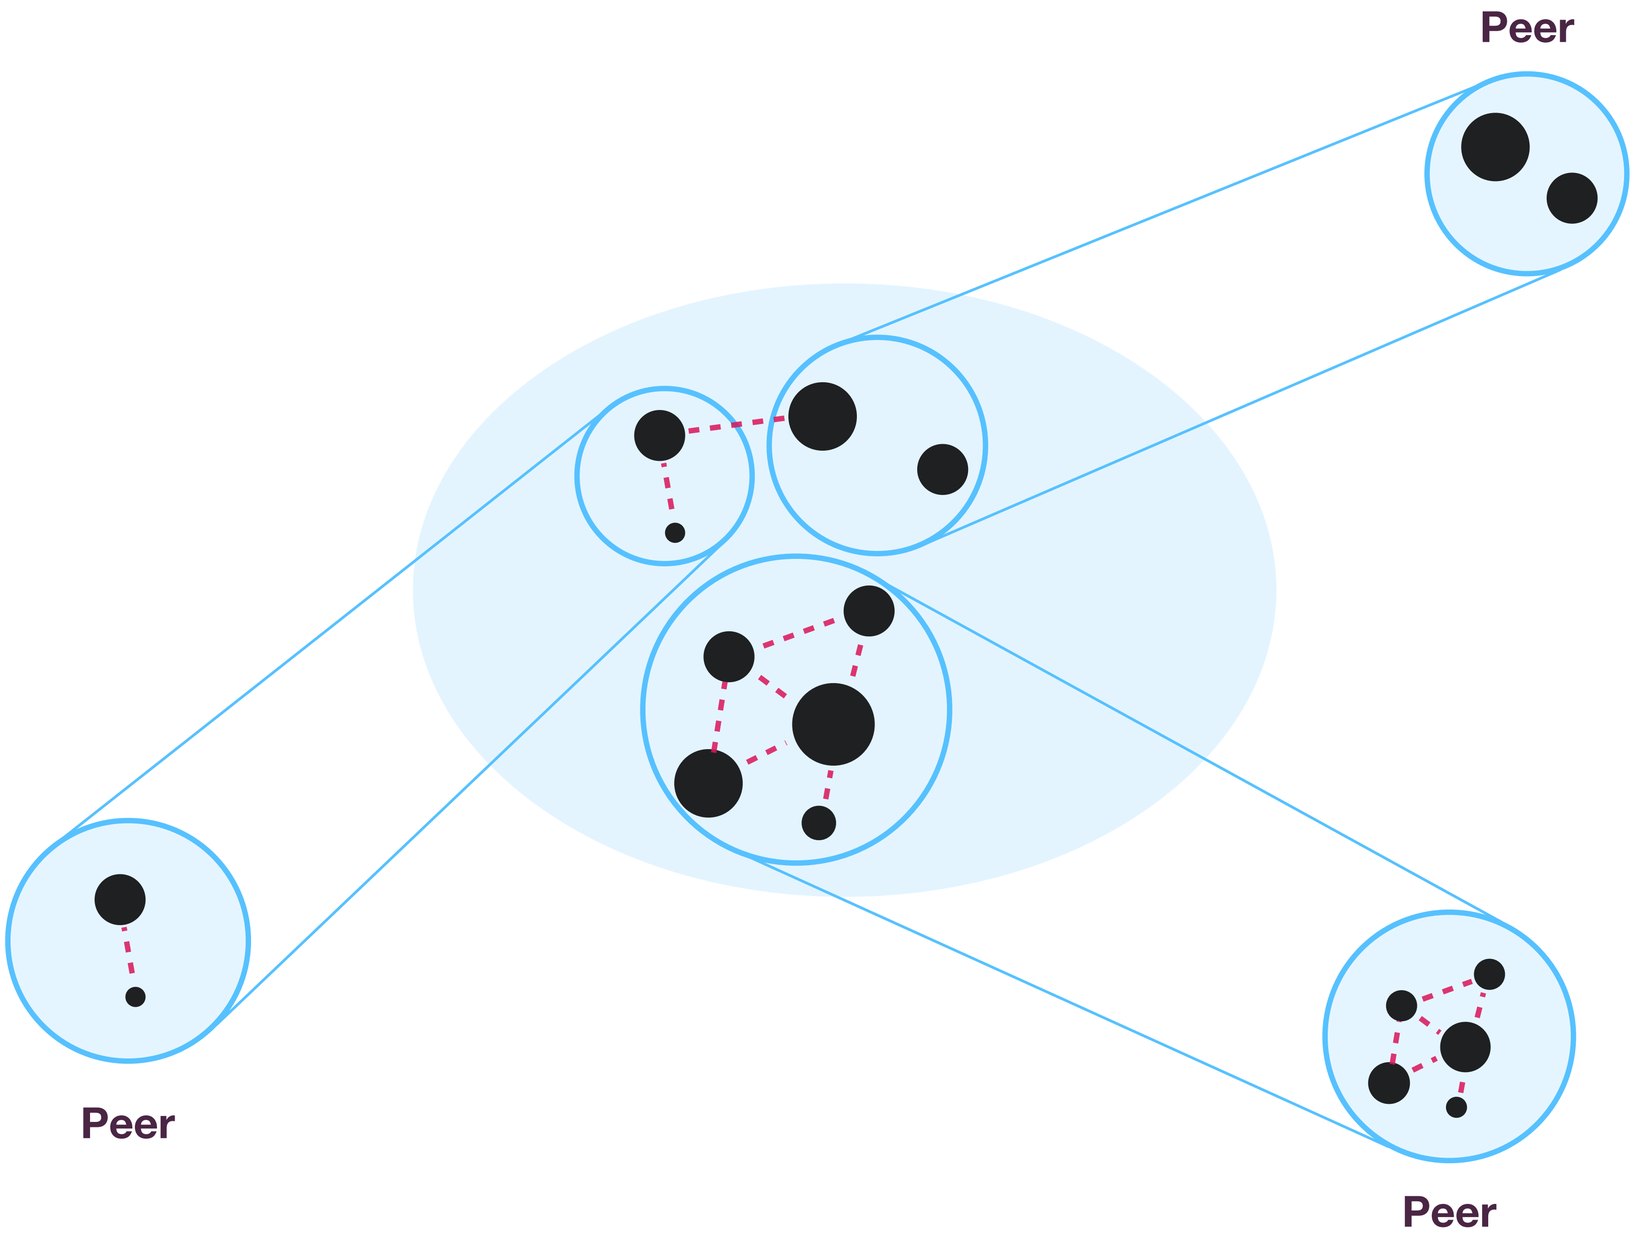
Peer
Peer
Peer

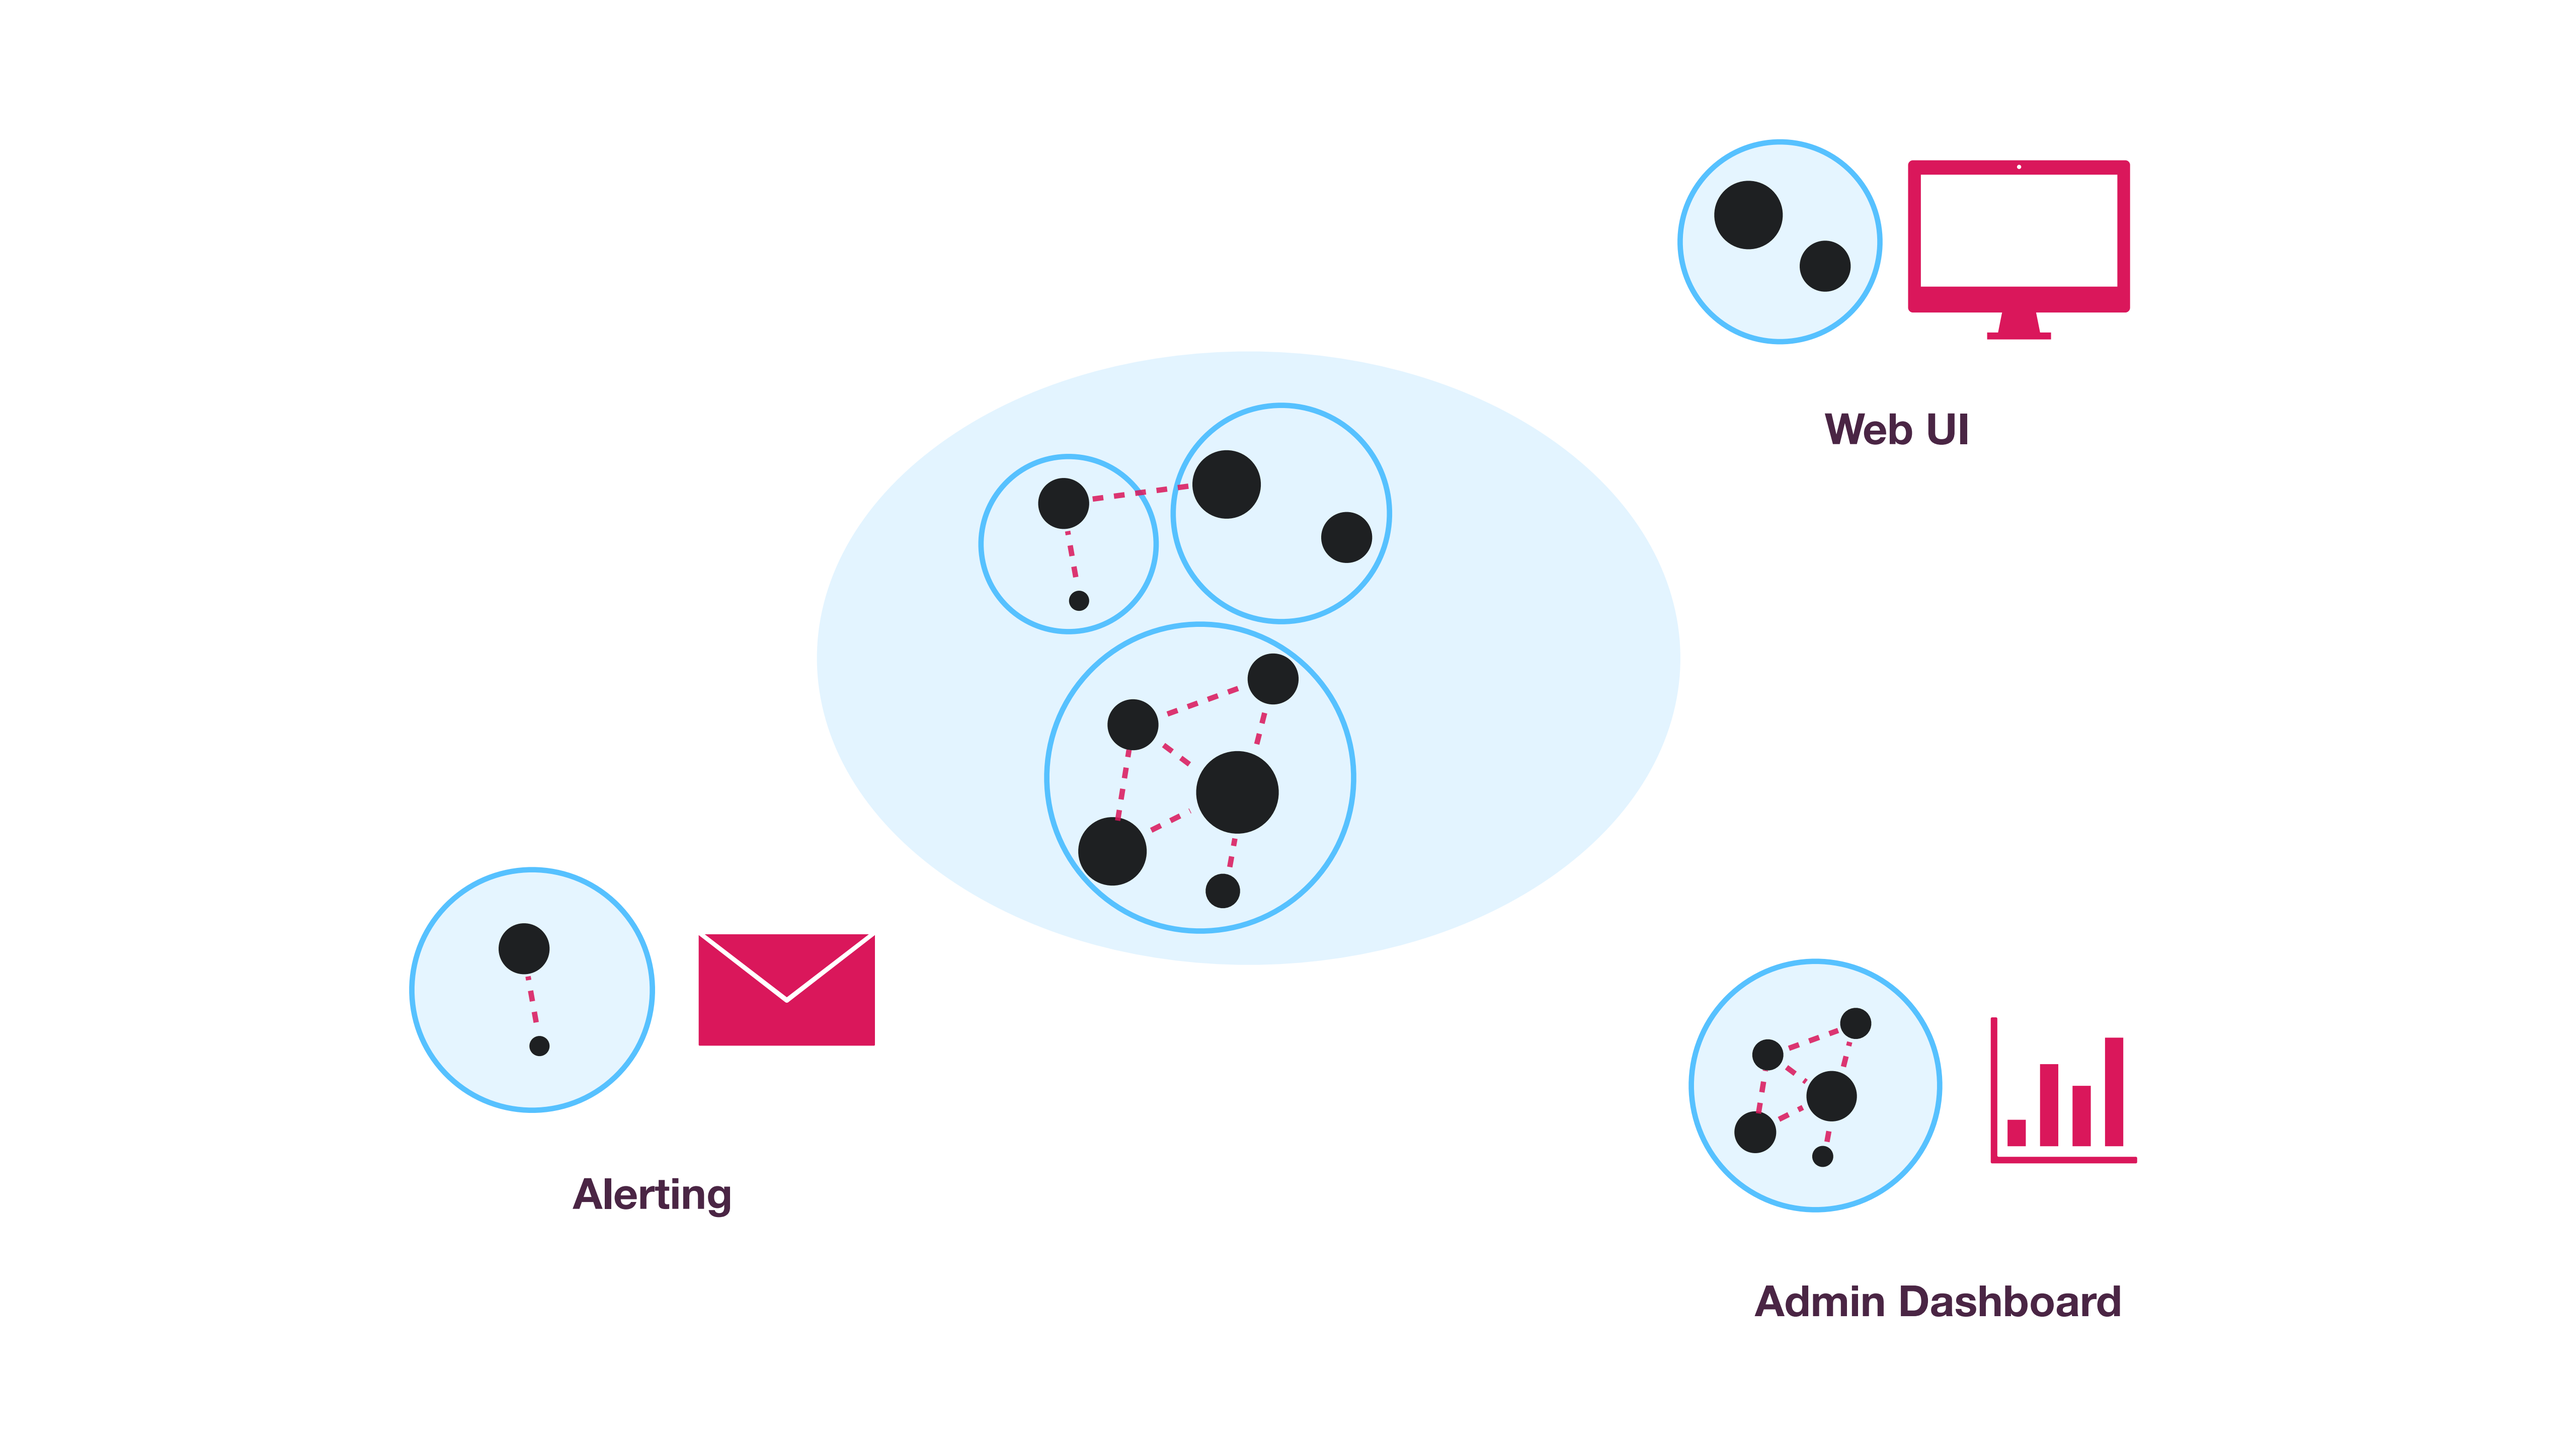
Web UI
Alerting
Admin Dashboard

Datomic:
Broadcast All Novelty To All Peers

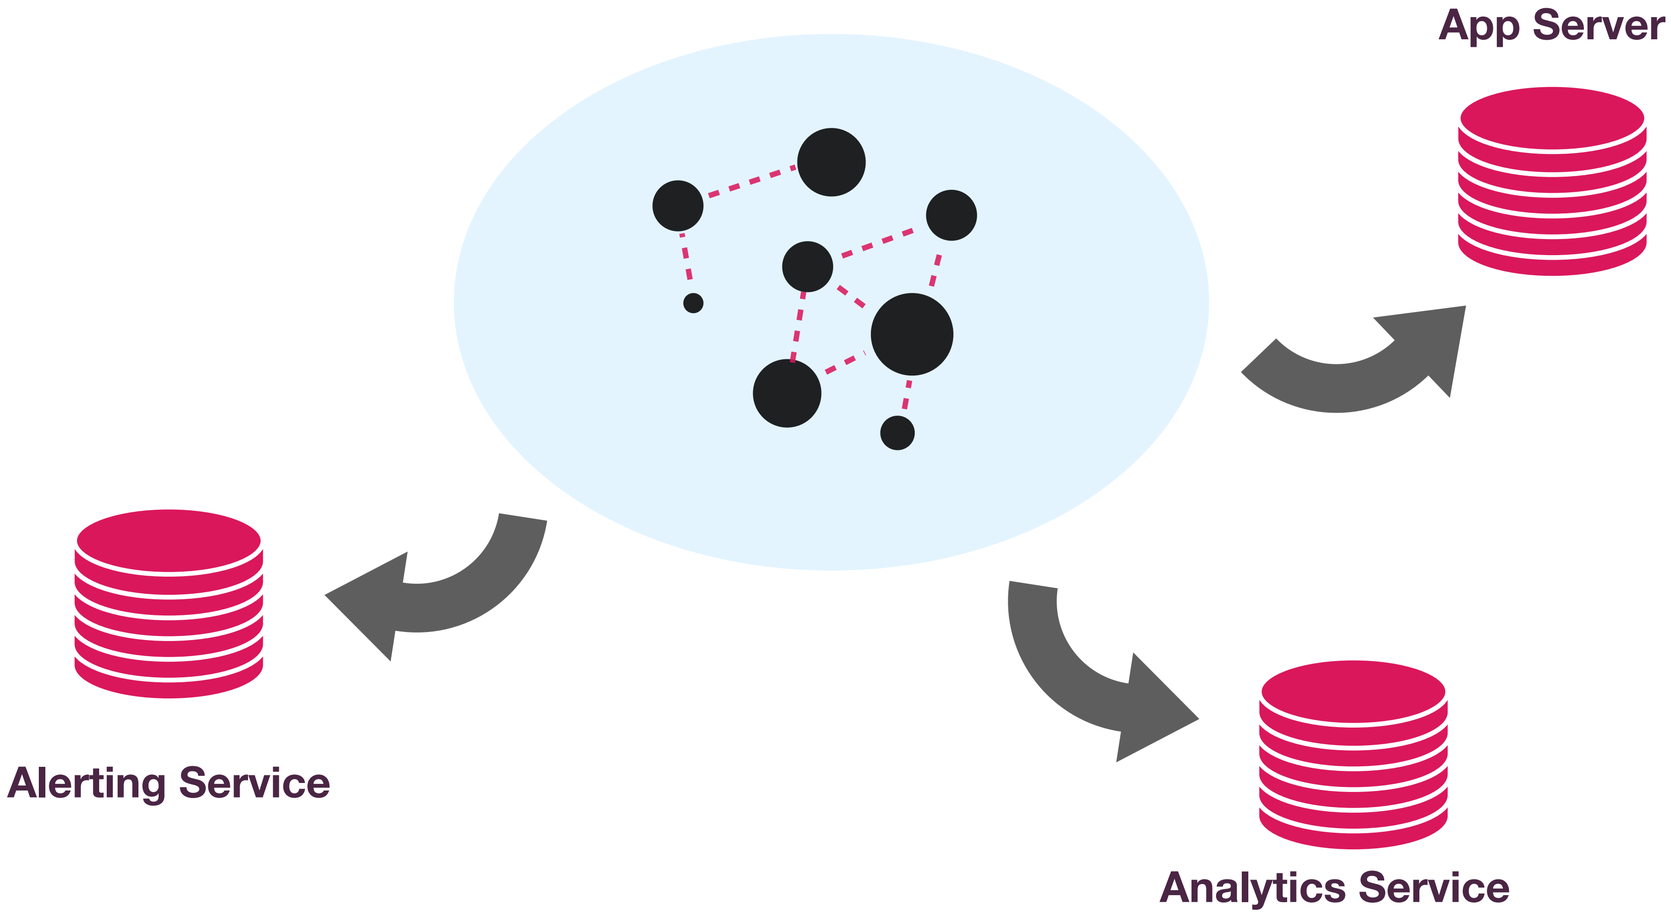

Ions:
Move Consumers Into The DB

Selective Replication

multiple data sources?
Datomic shards, Kafka, external services…

only send relevant novelty to peers?
for browsers or mobile clients, a LRU cache is not enough

runtimes other than the JVM?
browsers, ML, data science…

3DF:
Declarative Replication via Datalog

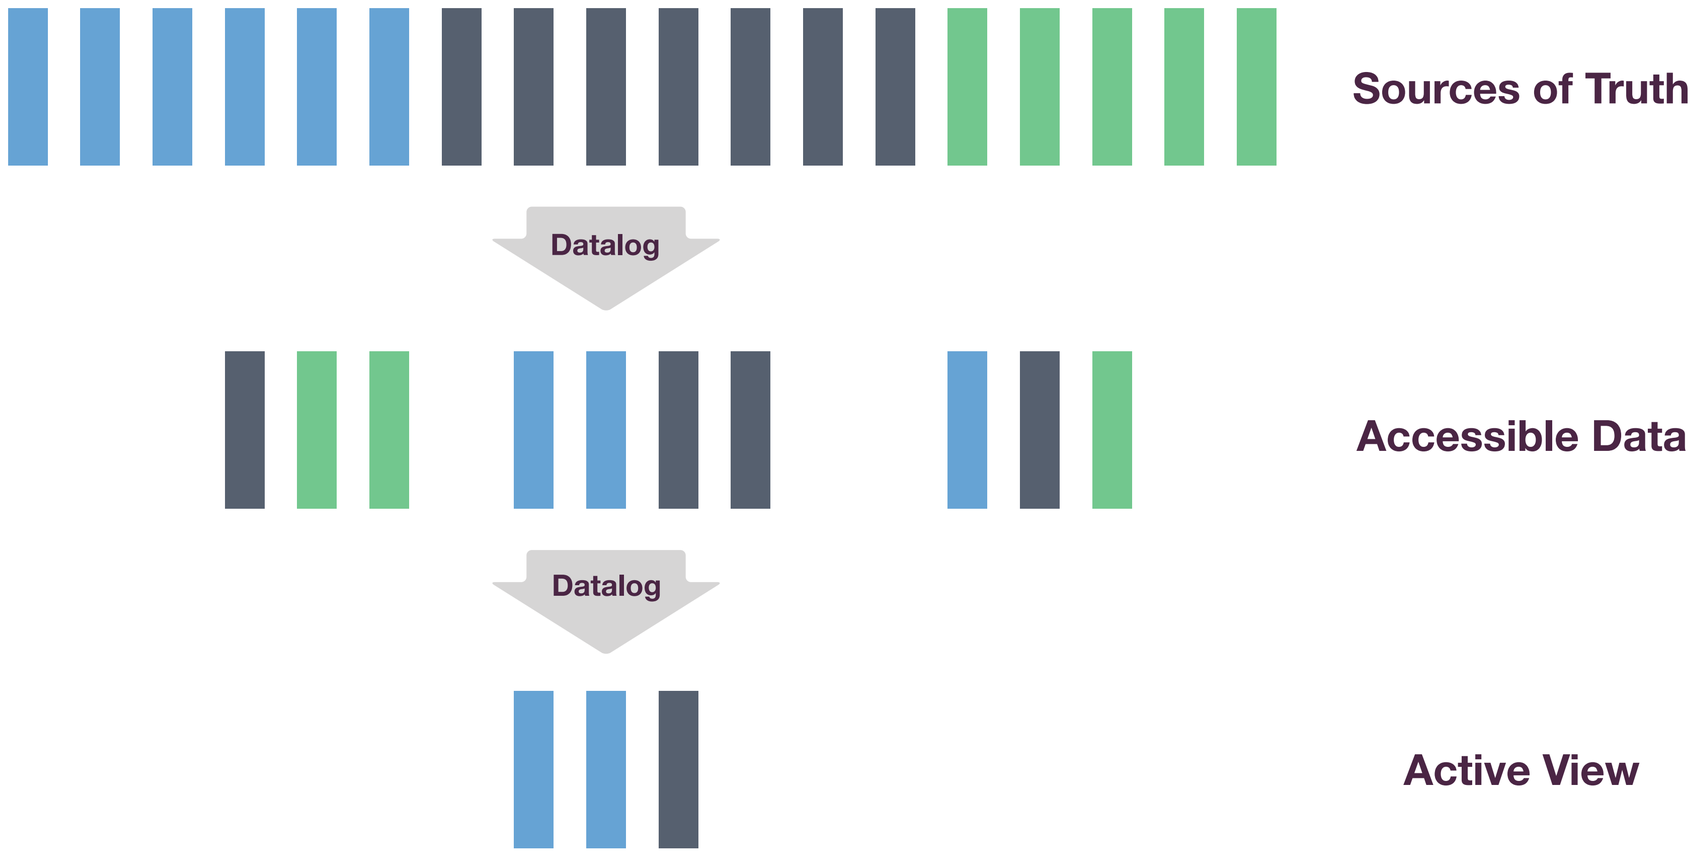

Selective Replication

- ✓ open to anyone who speaks Datalog
- ✓ only relevant novelty propagated to each consumer
- ✓ scales to deal with multiple sources

“We want to make reactive systems that don't poll. And we want those systems to get a consistent view of the world.”

– *Rich Hickey, “Deconstructing the Database”*

3DF lets you…

- …leverage Datalog to build reactive systems
- …treat data consumers as database peers
- …scale queries beyond the limits of a single peer
- …integrate with hand-written Differential Dataflow

Interested? Let's talk!

@NikolasGoebel | niko@clockworks.io

github.com/comnik/clj-3df

github.com/TimelyDataflow