

It's About **Time**: An Introduction to Timely Dataflow

Data Council, October '19



Malte Sandstede

malte@clockworks.io / [@MalteSandstede](https://twitter.com/MalteSandstede)

Nikolas Göbel

niko@clockworks.io / [@NikolasGoebel](https://twitter.com/NikolasGoebel)

David Bach

david@clockworks.io

Moritz Moxter


moritz@clockworks.io

In collaboration with:

Frank McSherry


Vasia Kalavri (ETH)

***ETH* zürich**



Systems Group


Stream Processing's Trifecta



Stream Processing's Trifecta

Naive Stateless Processing


- Low latency
- Issue: Late arrivals
- Issue: Complex computations



Stream Processing's Trifecta

MapReduce

- No late arrivals (by definition)
- Easy to scale
- Issue: Complex computations
- Issue: High latency



Stream Processing's Trifecta

Database

- No late arrivals
- High expressivity
- ACID
- Issue: Not realtime!


Timeliness

Consistency

Expressivity




Stream Processing's Trifecta

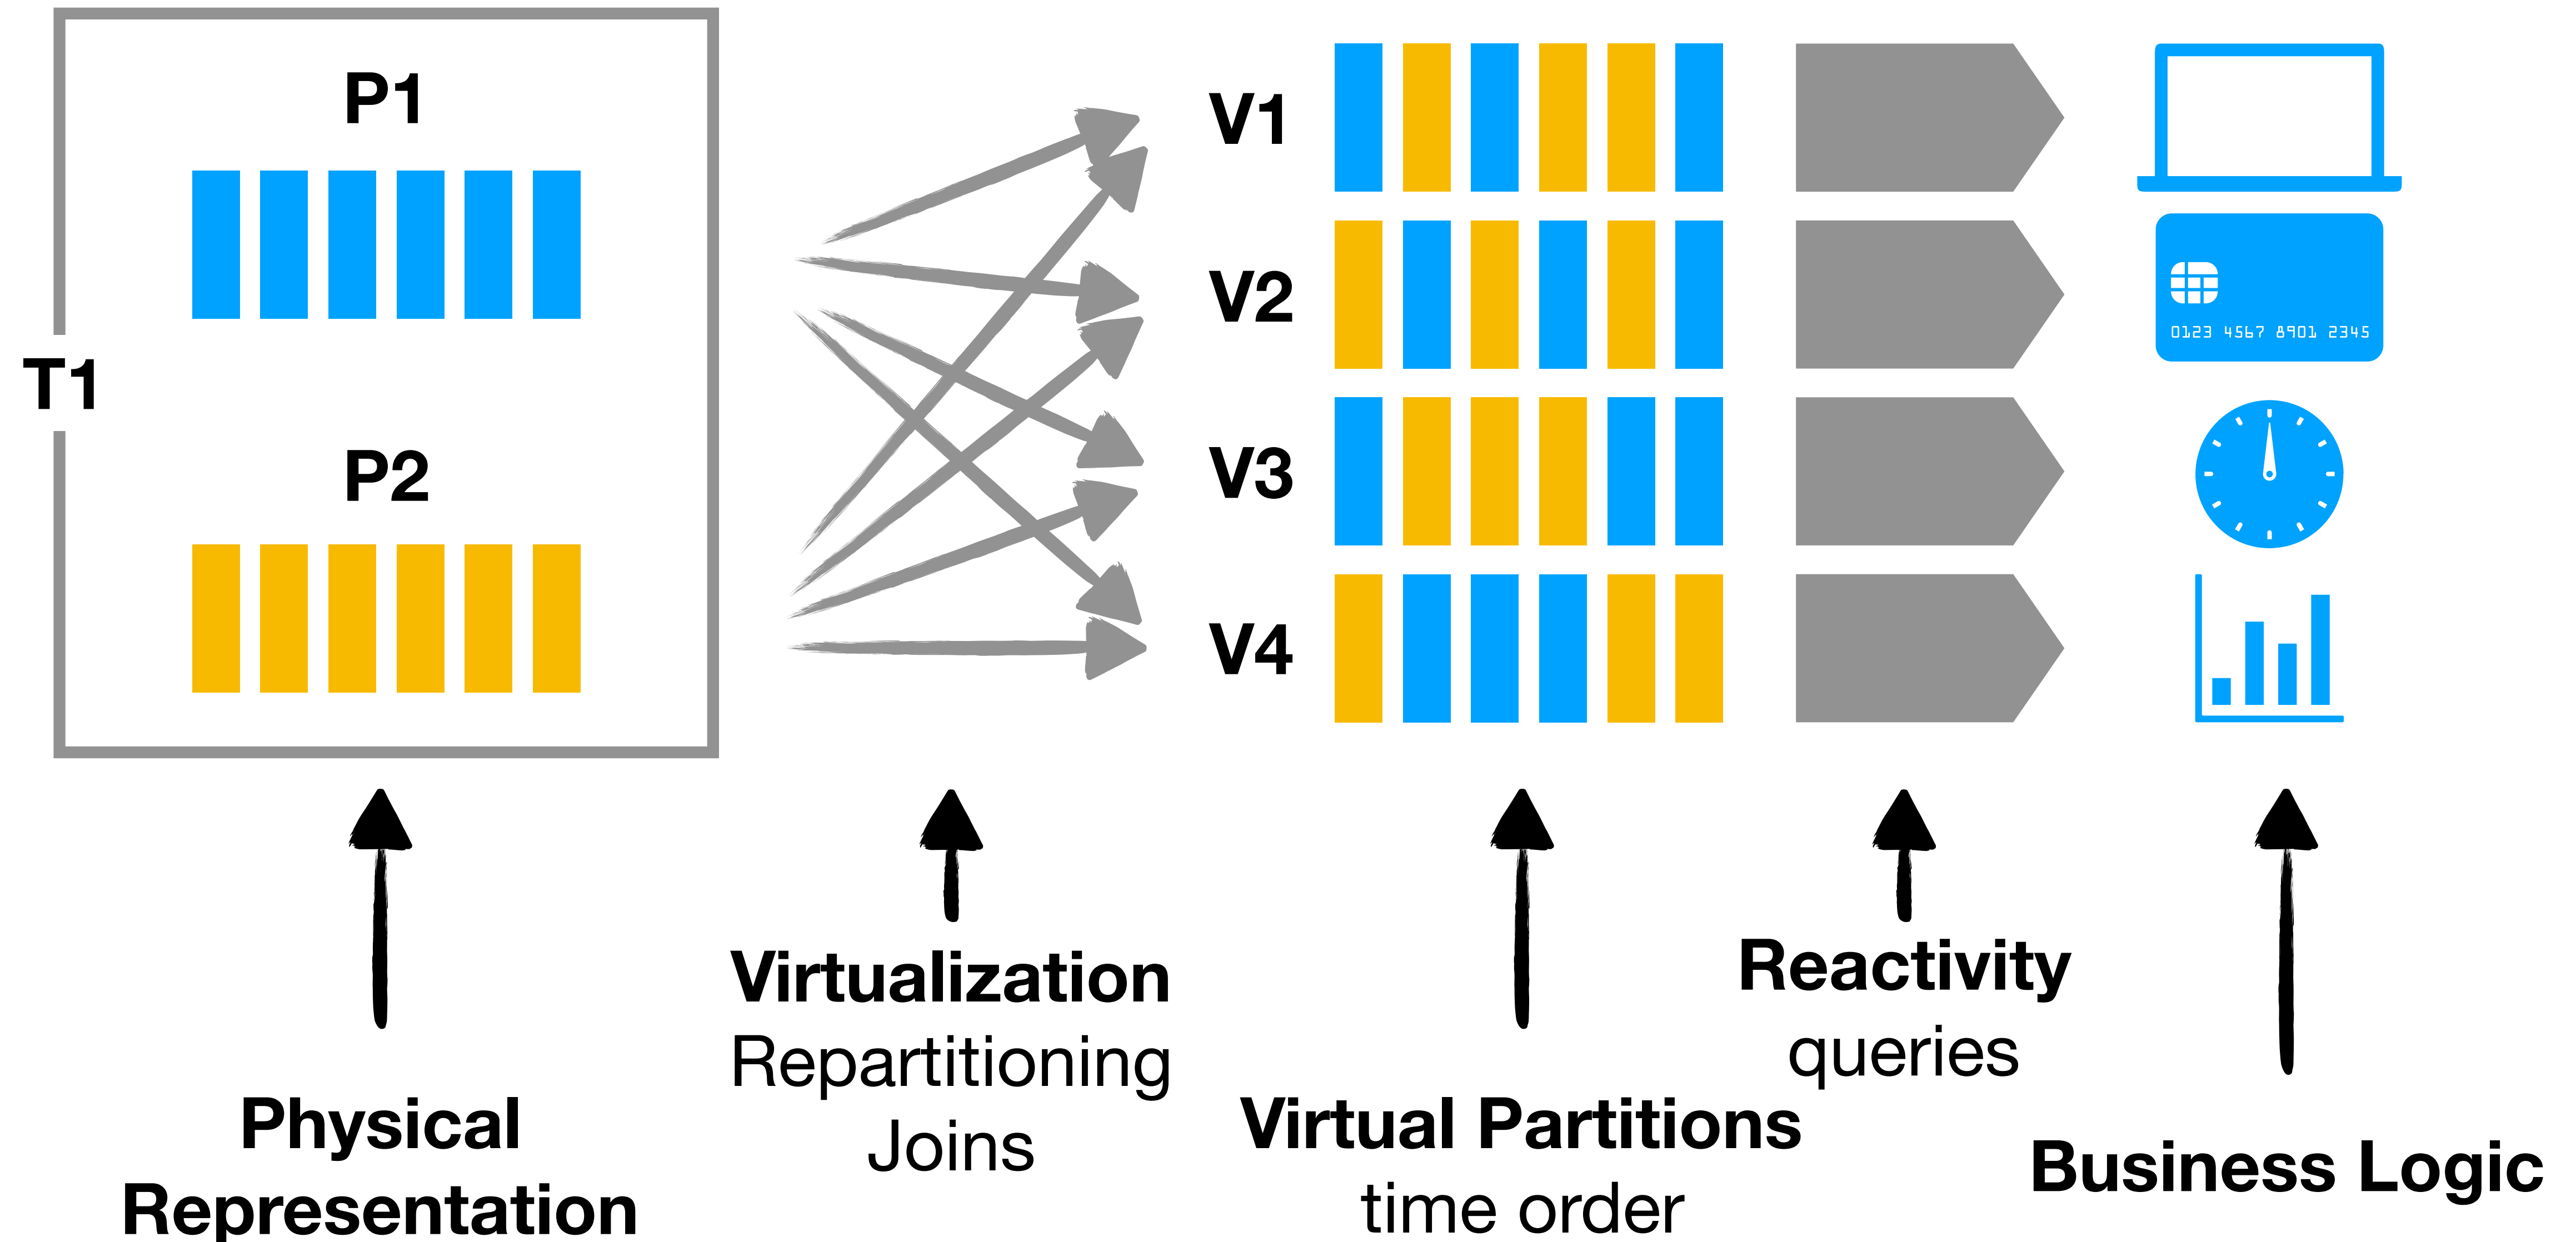


Use Case: Kafka Superpowers


(Partitions complete physical representation & use case)




Use Case: Kafka Superpowers




Stream Processing as Dataflow



Dataflow Parallelism




Dataflow Distribution



Correctness Troubles

DATA



SUM



Correctness Troubles



Correctness Troubles



Correctness Troubles

DATA

SUM

(3, t_0)

(5, t_1)

(1, t_0)



Timely Dataflow

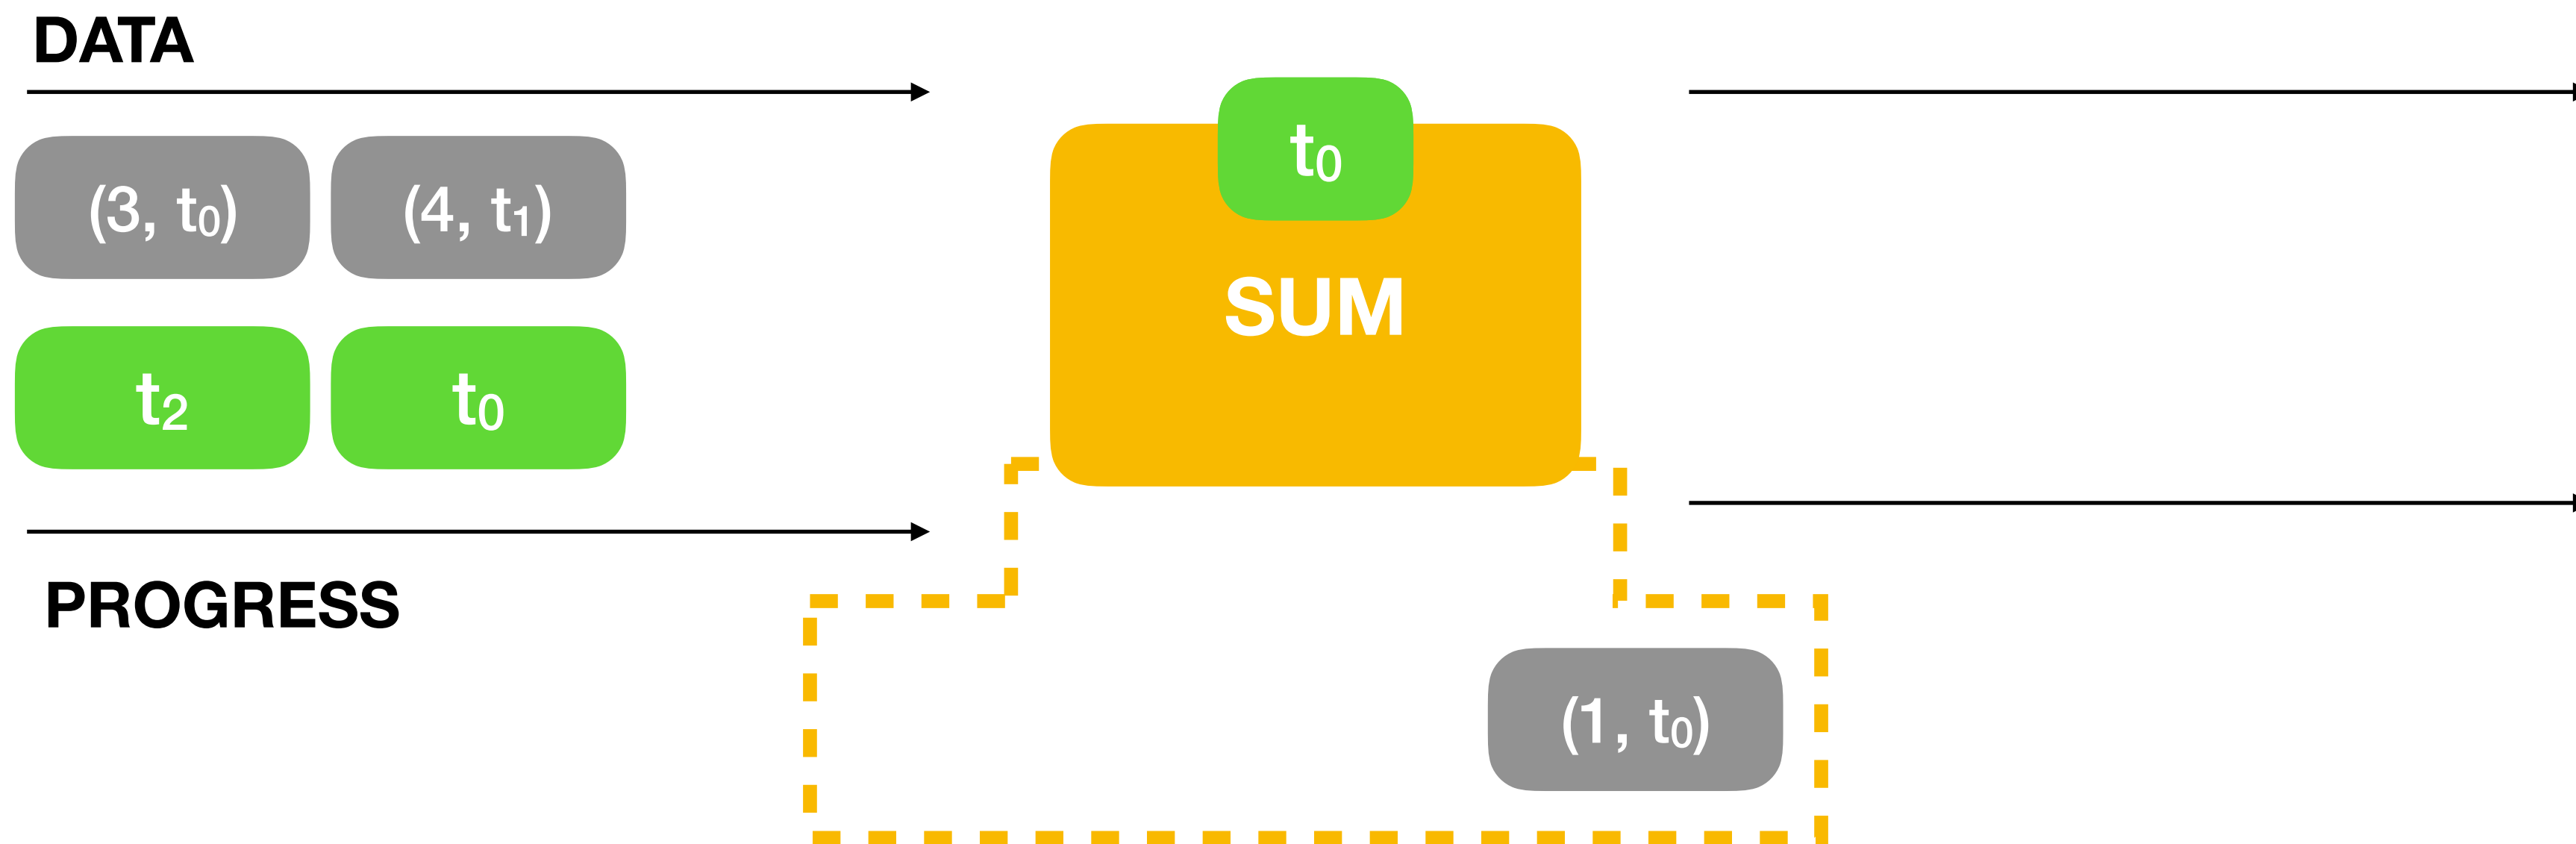
A low-latency runtime for
distributed cyclic dataflows

github.com/TimelyDataflow

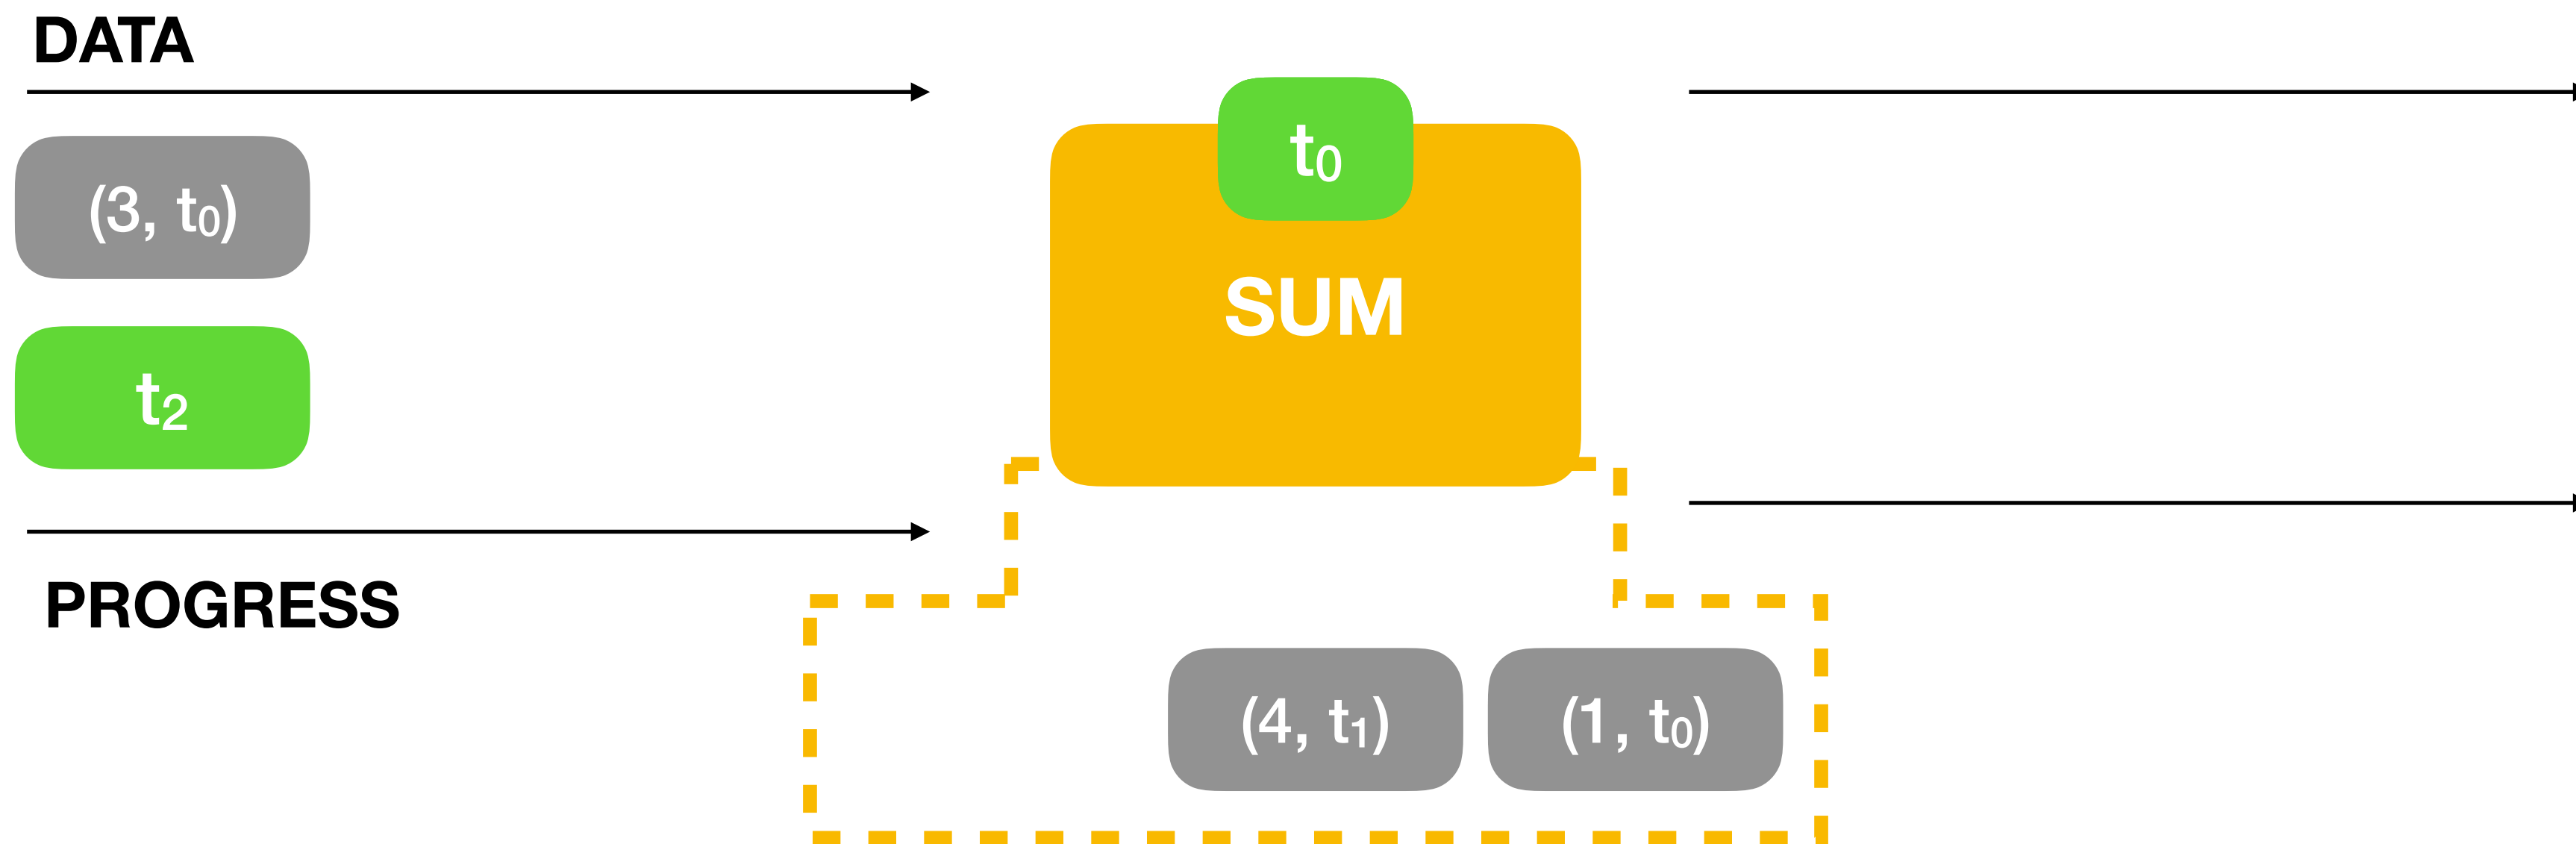
Correctness with Progress Tracking



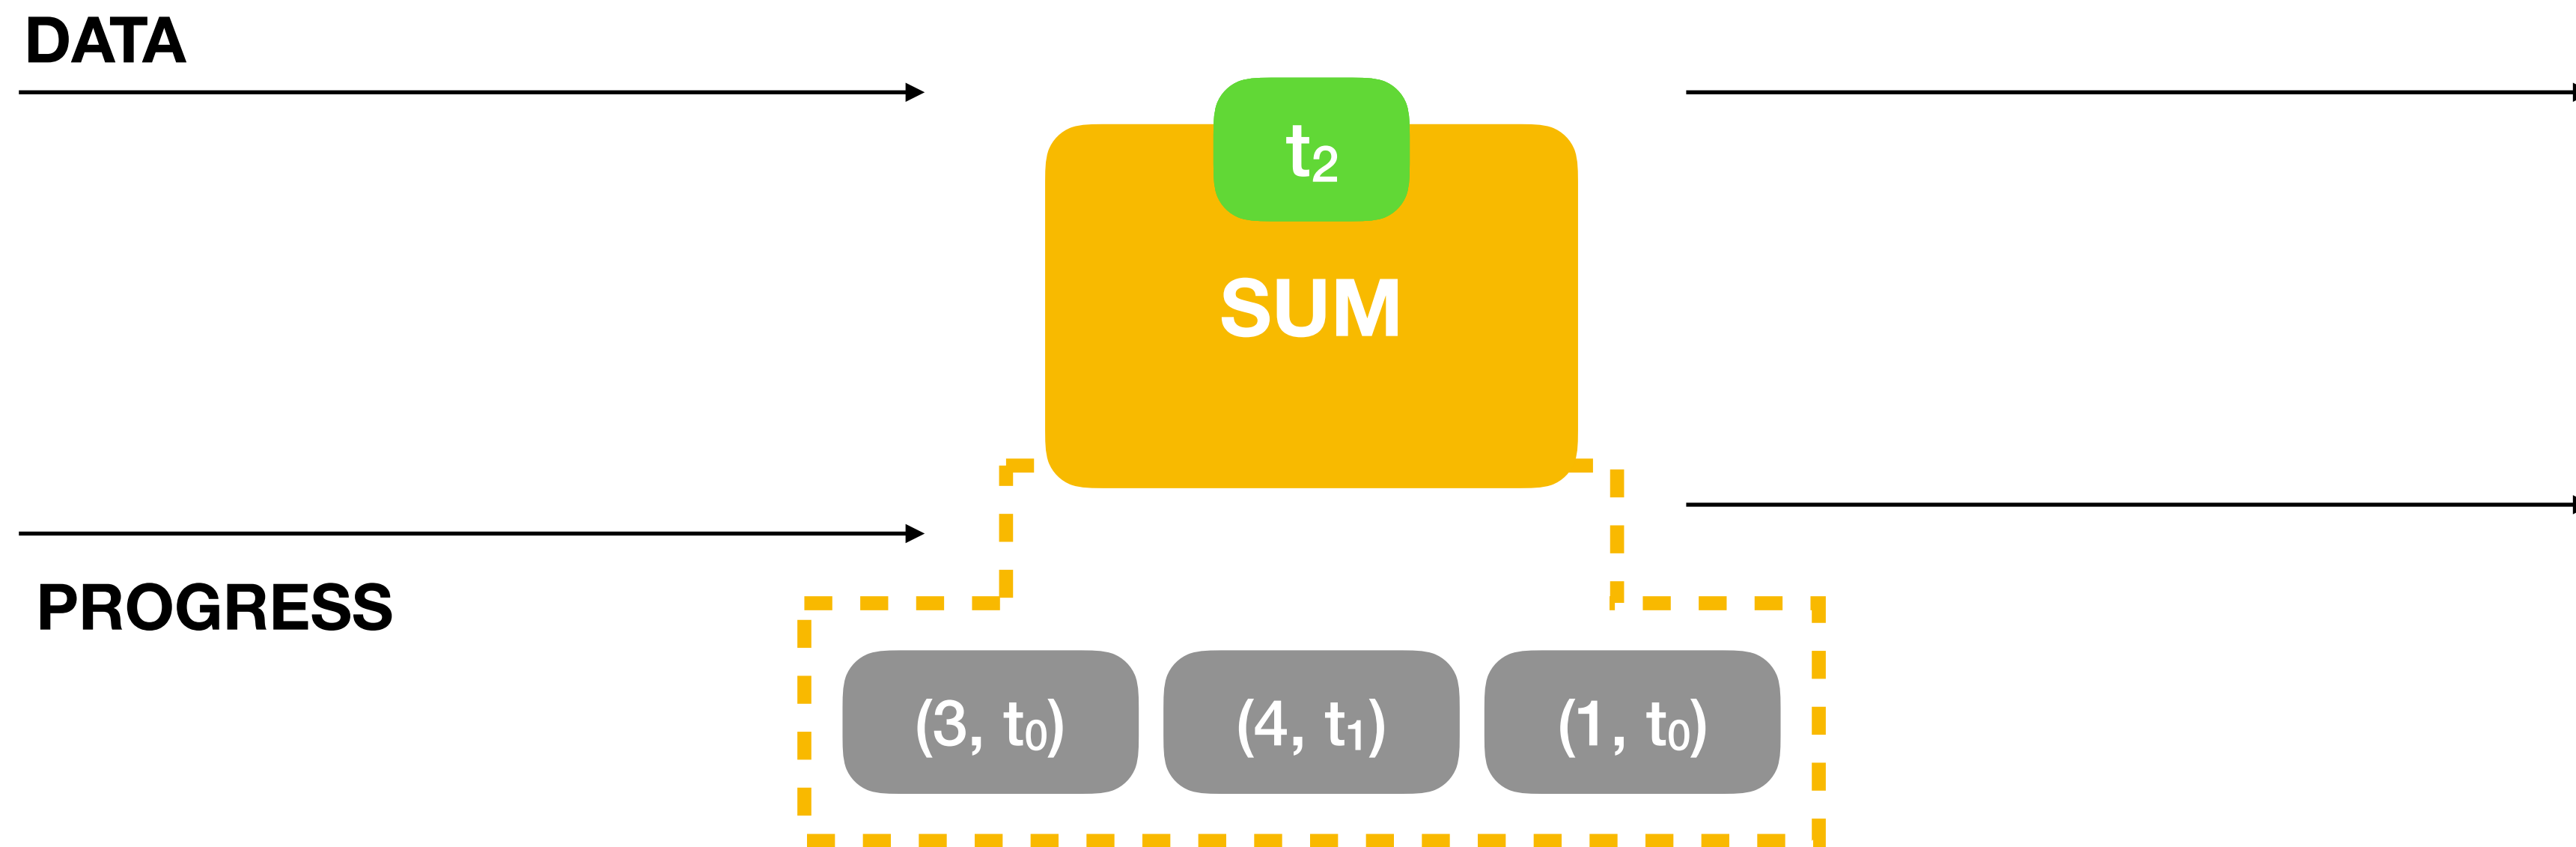
Correctness with Progress Tracking



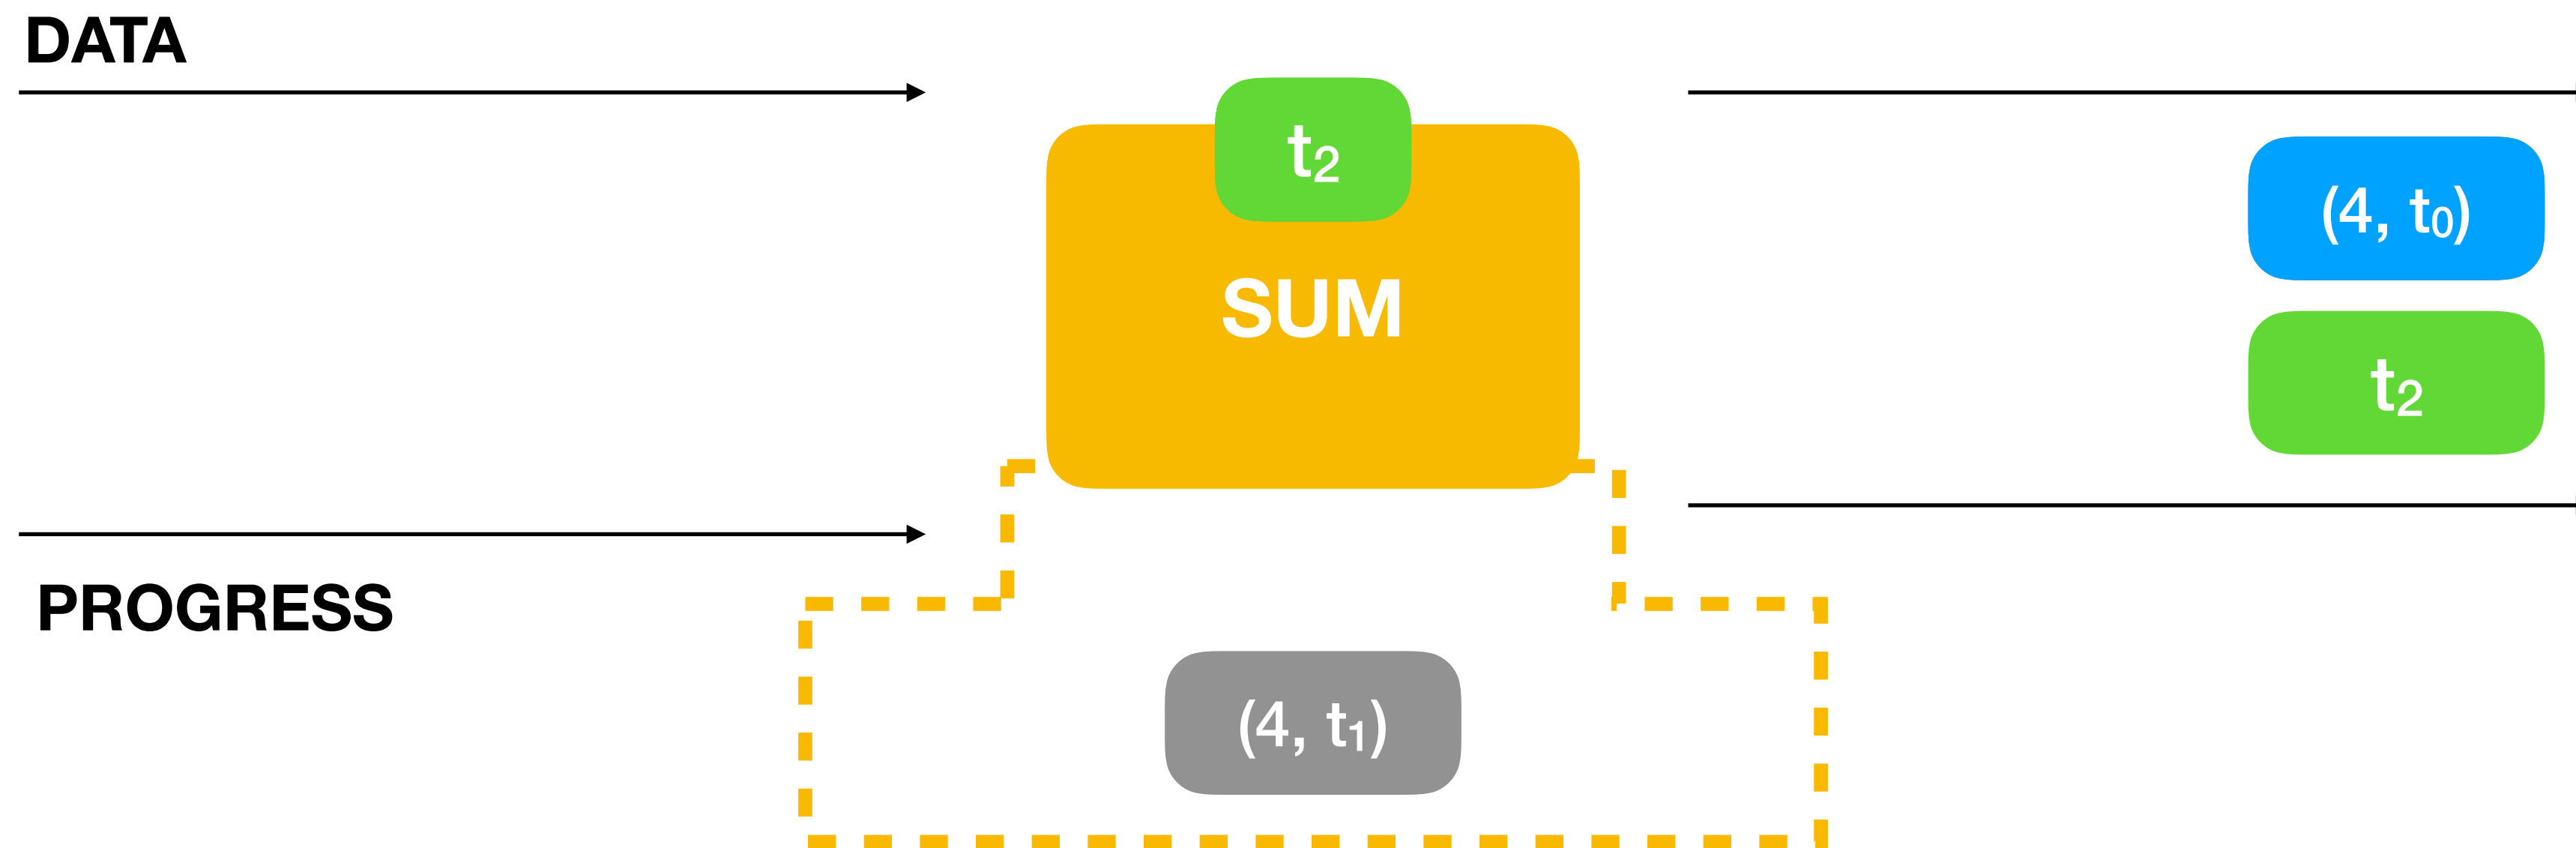
Correctness with Progress Tracking



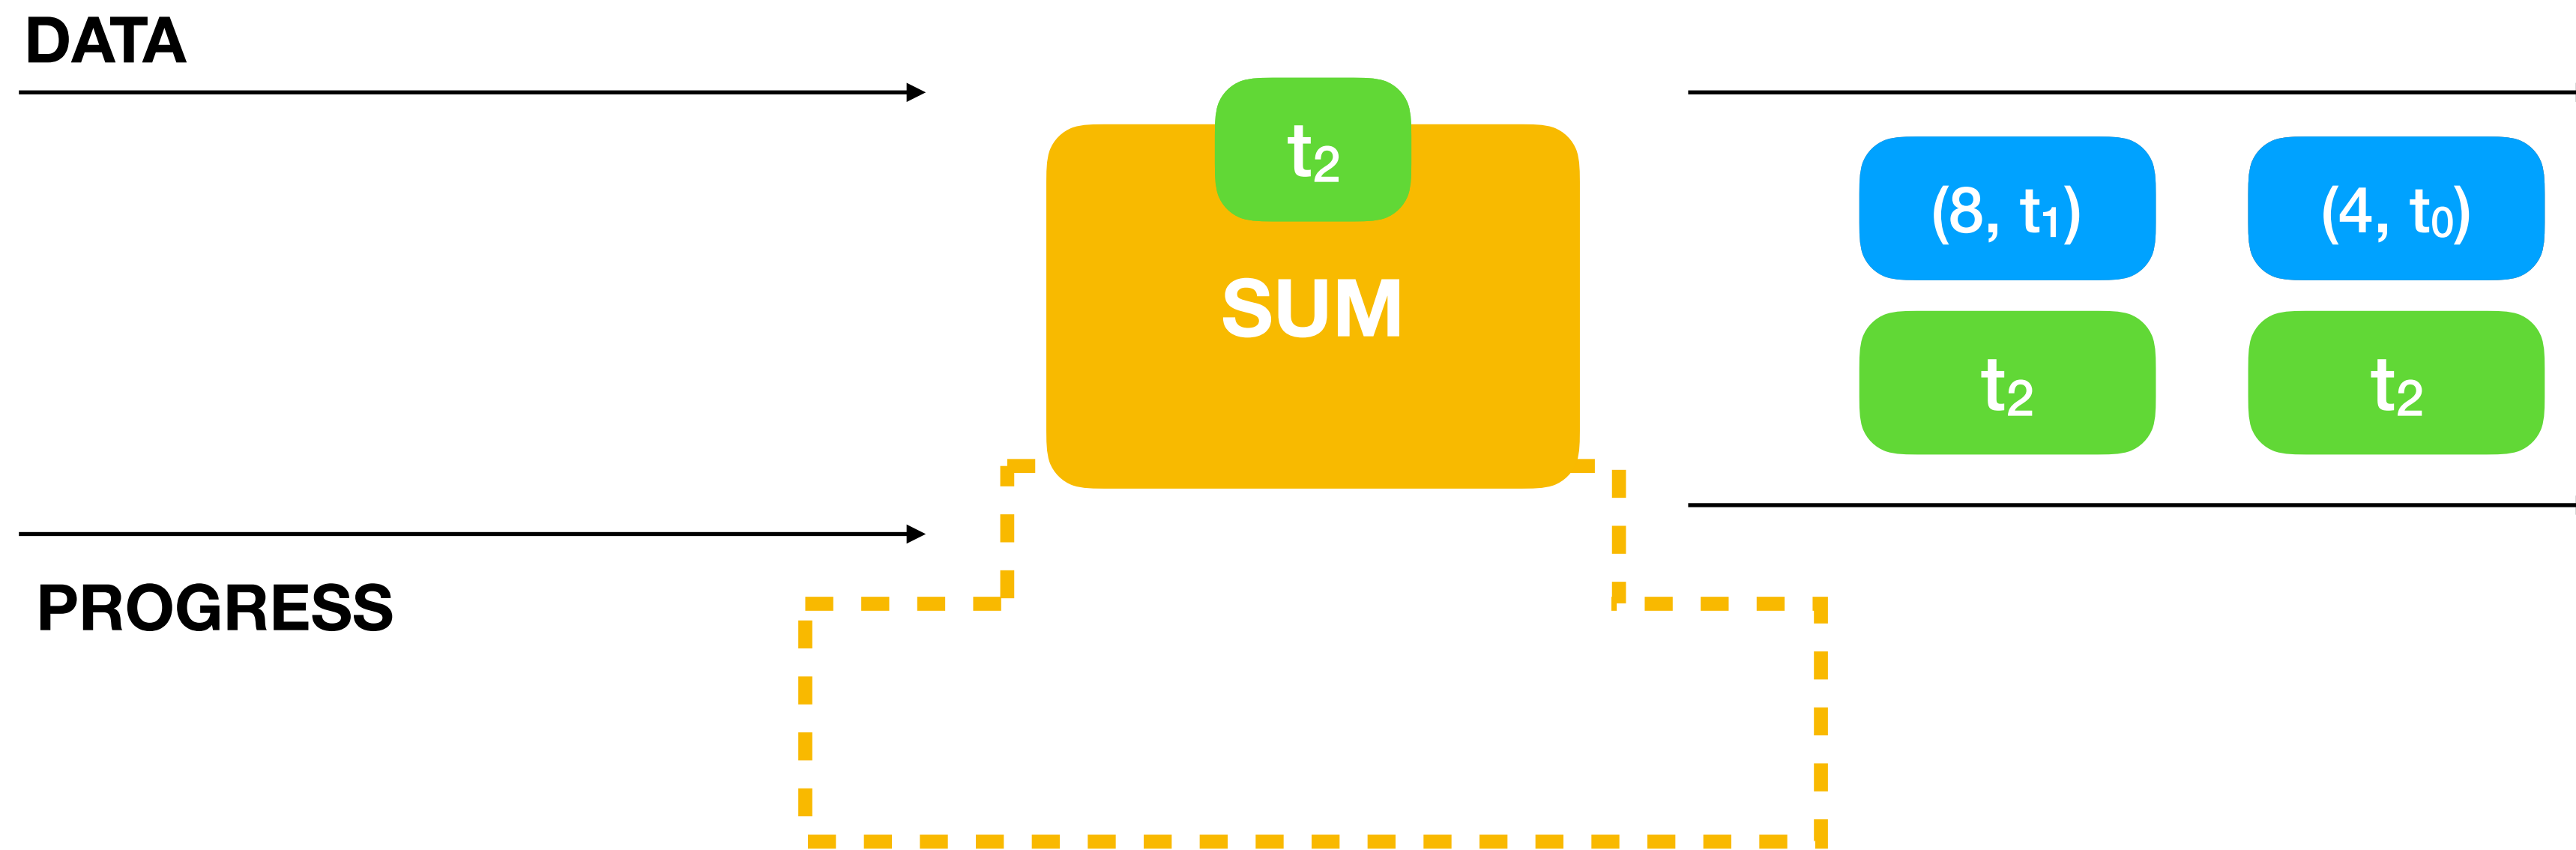
Correctness with Progress Tracking



Correctness with Progress Tracking

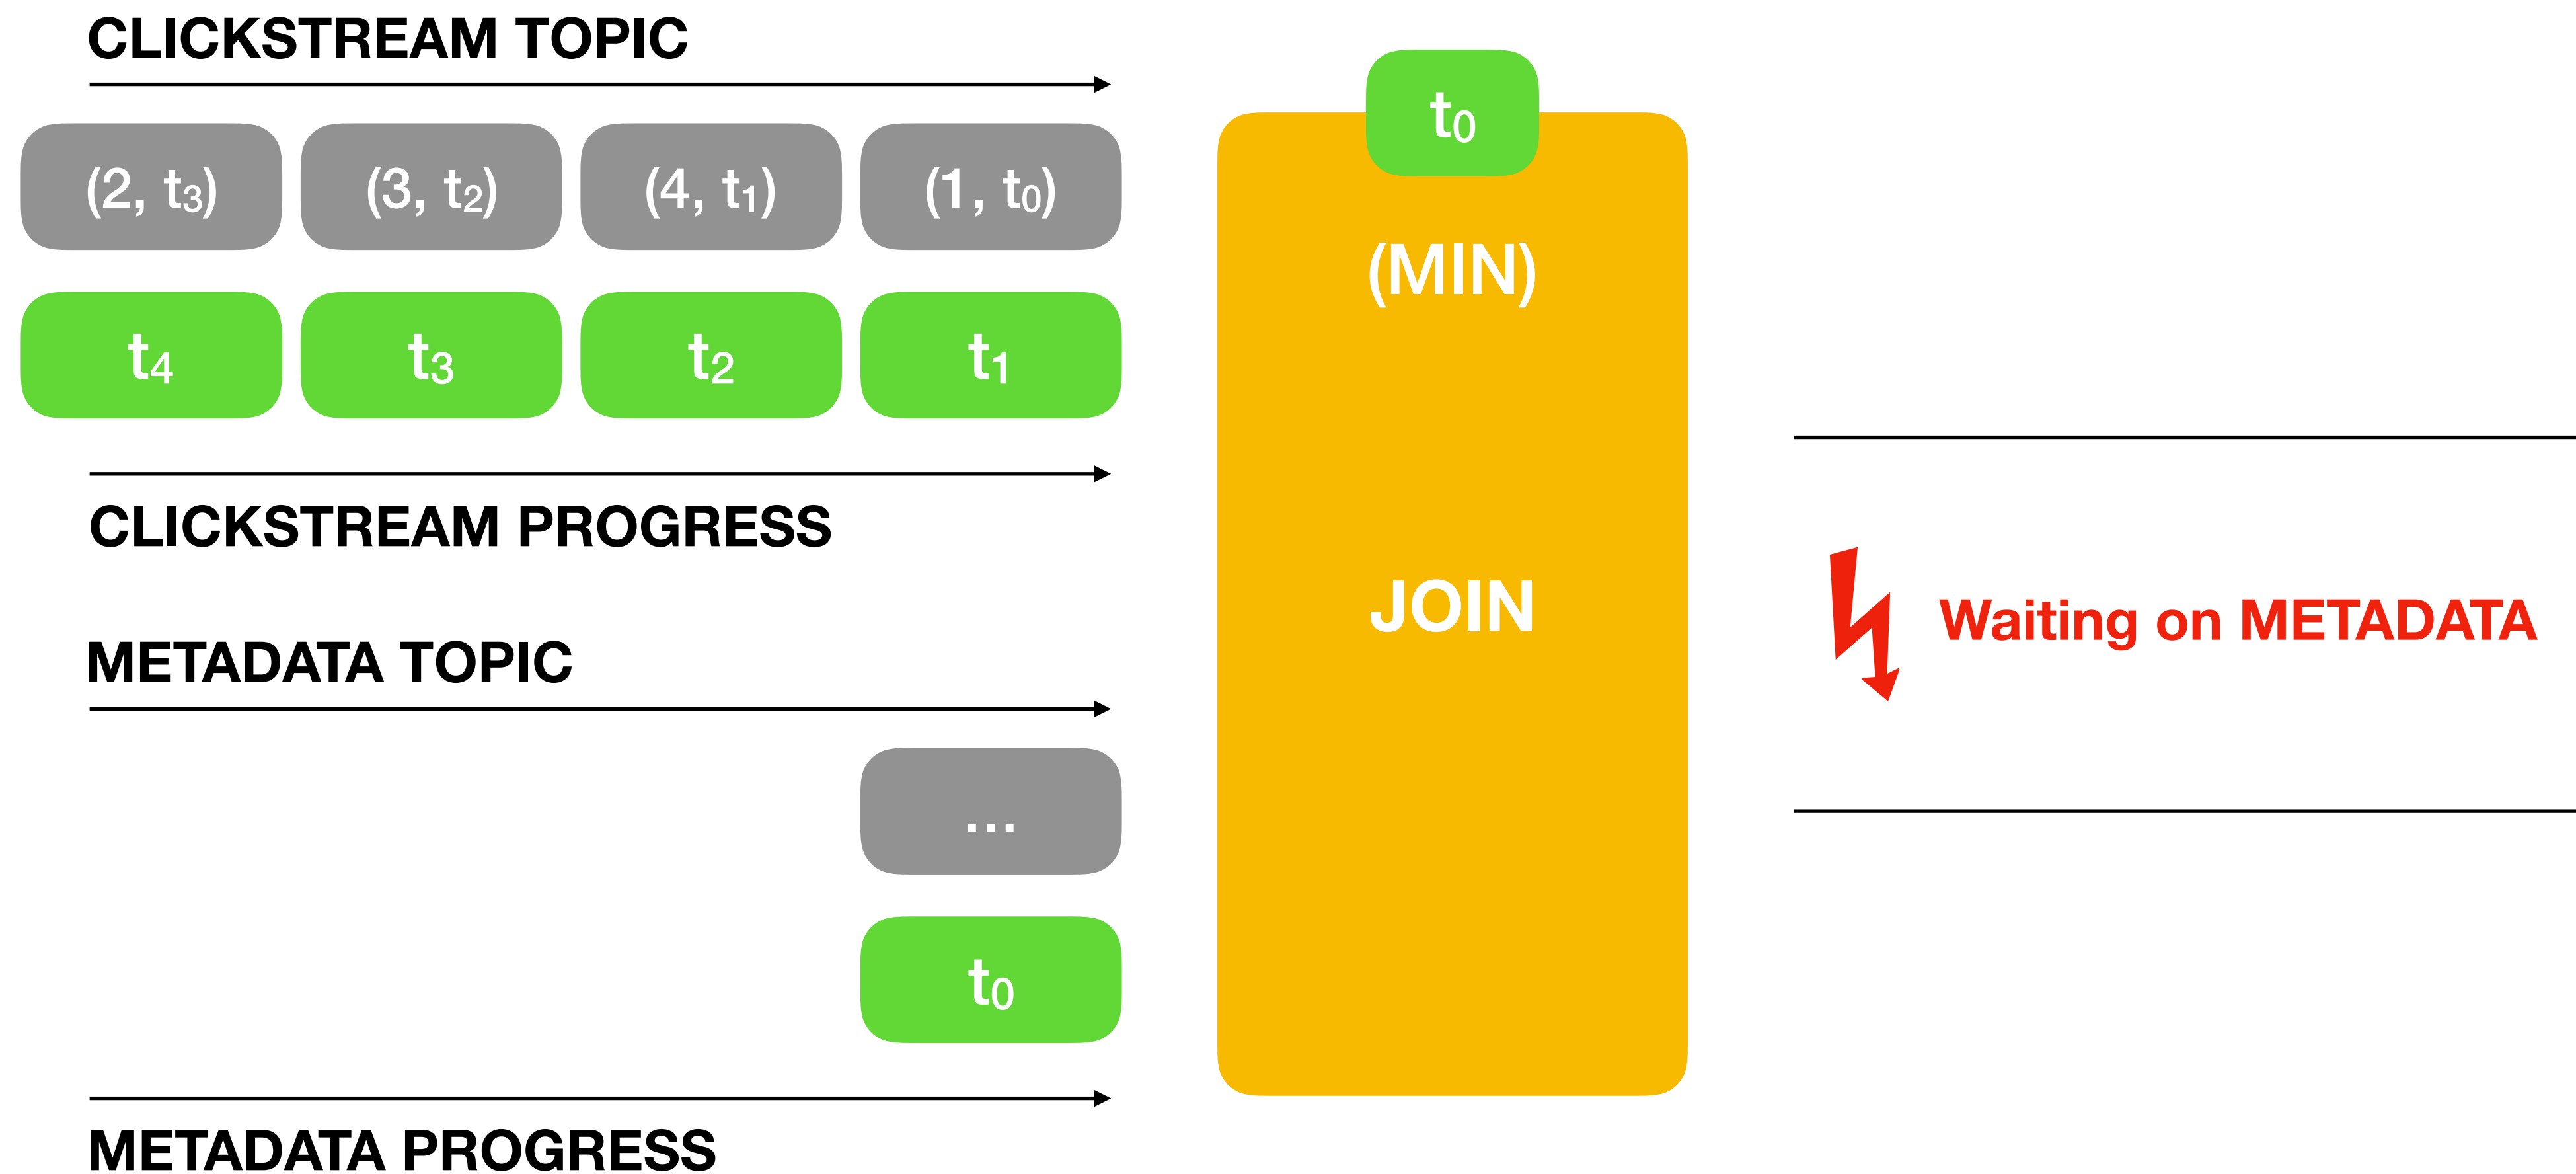


Correctness with Progress Tracking



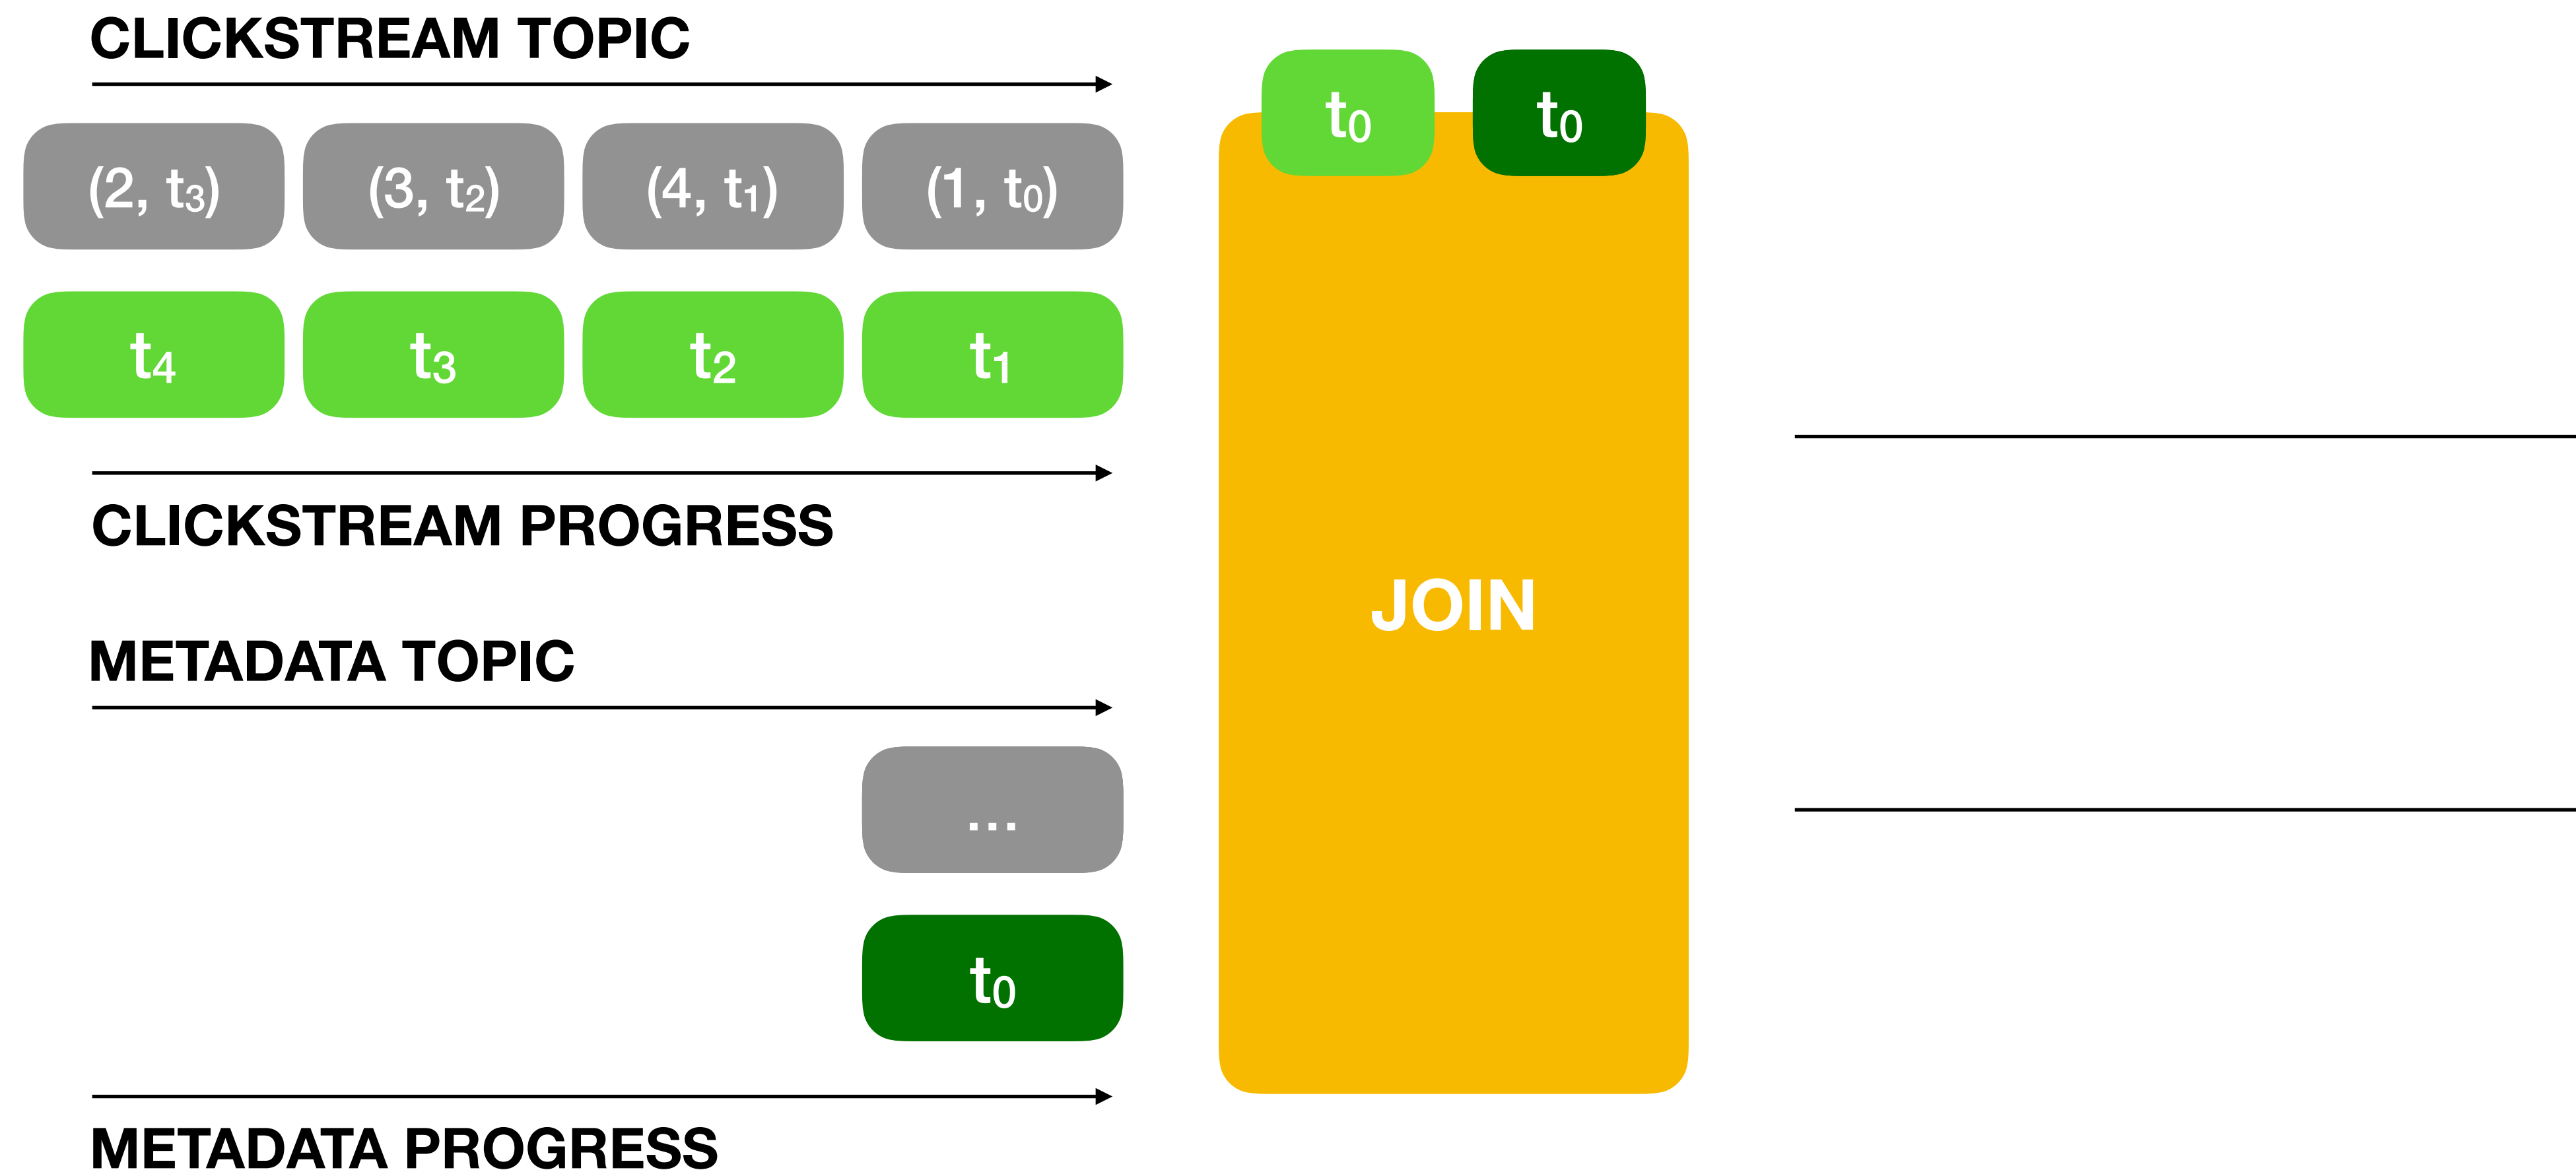
Progress Tracking... without Progress?

(data sources with different event frequencies)



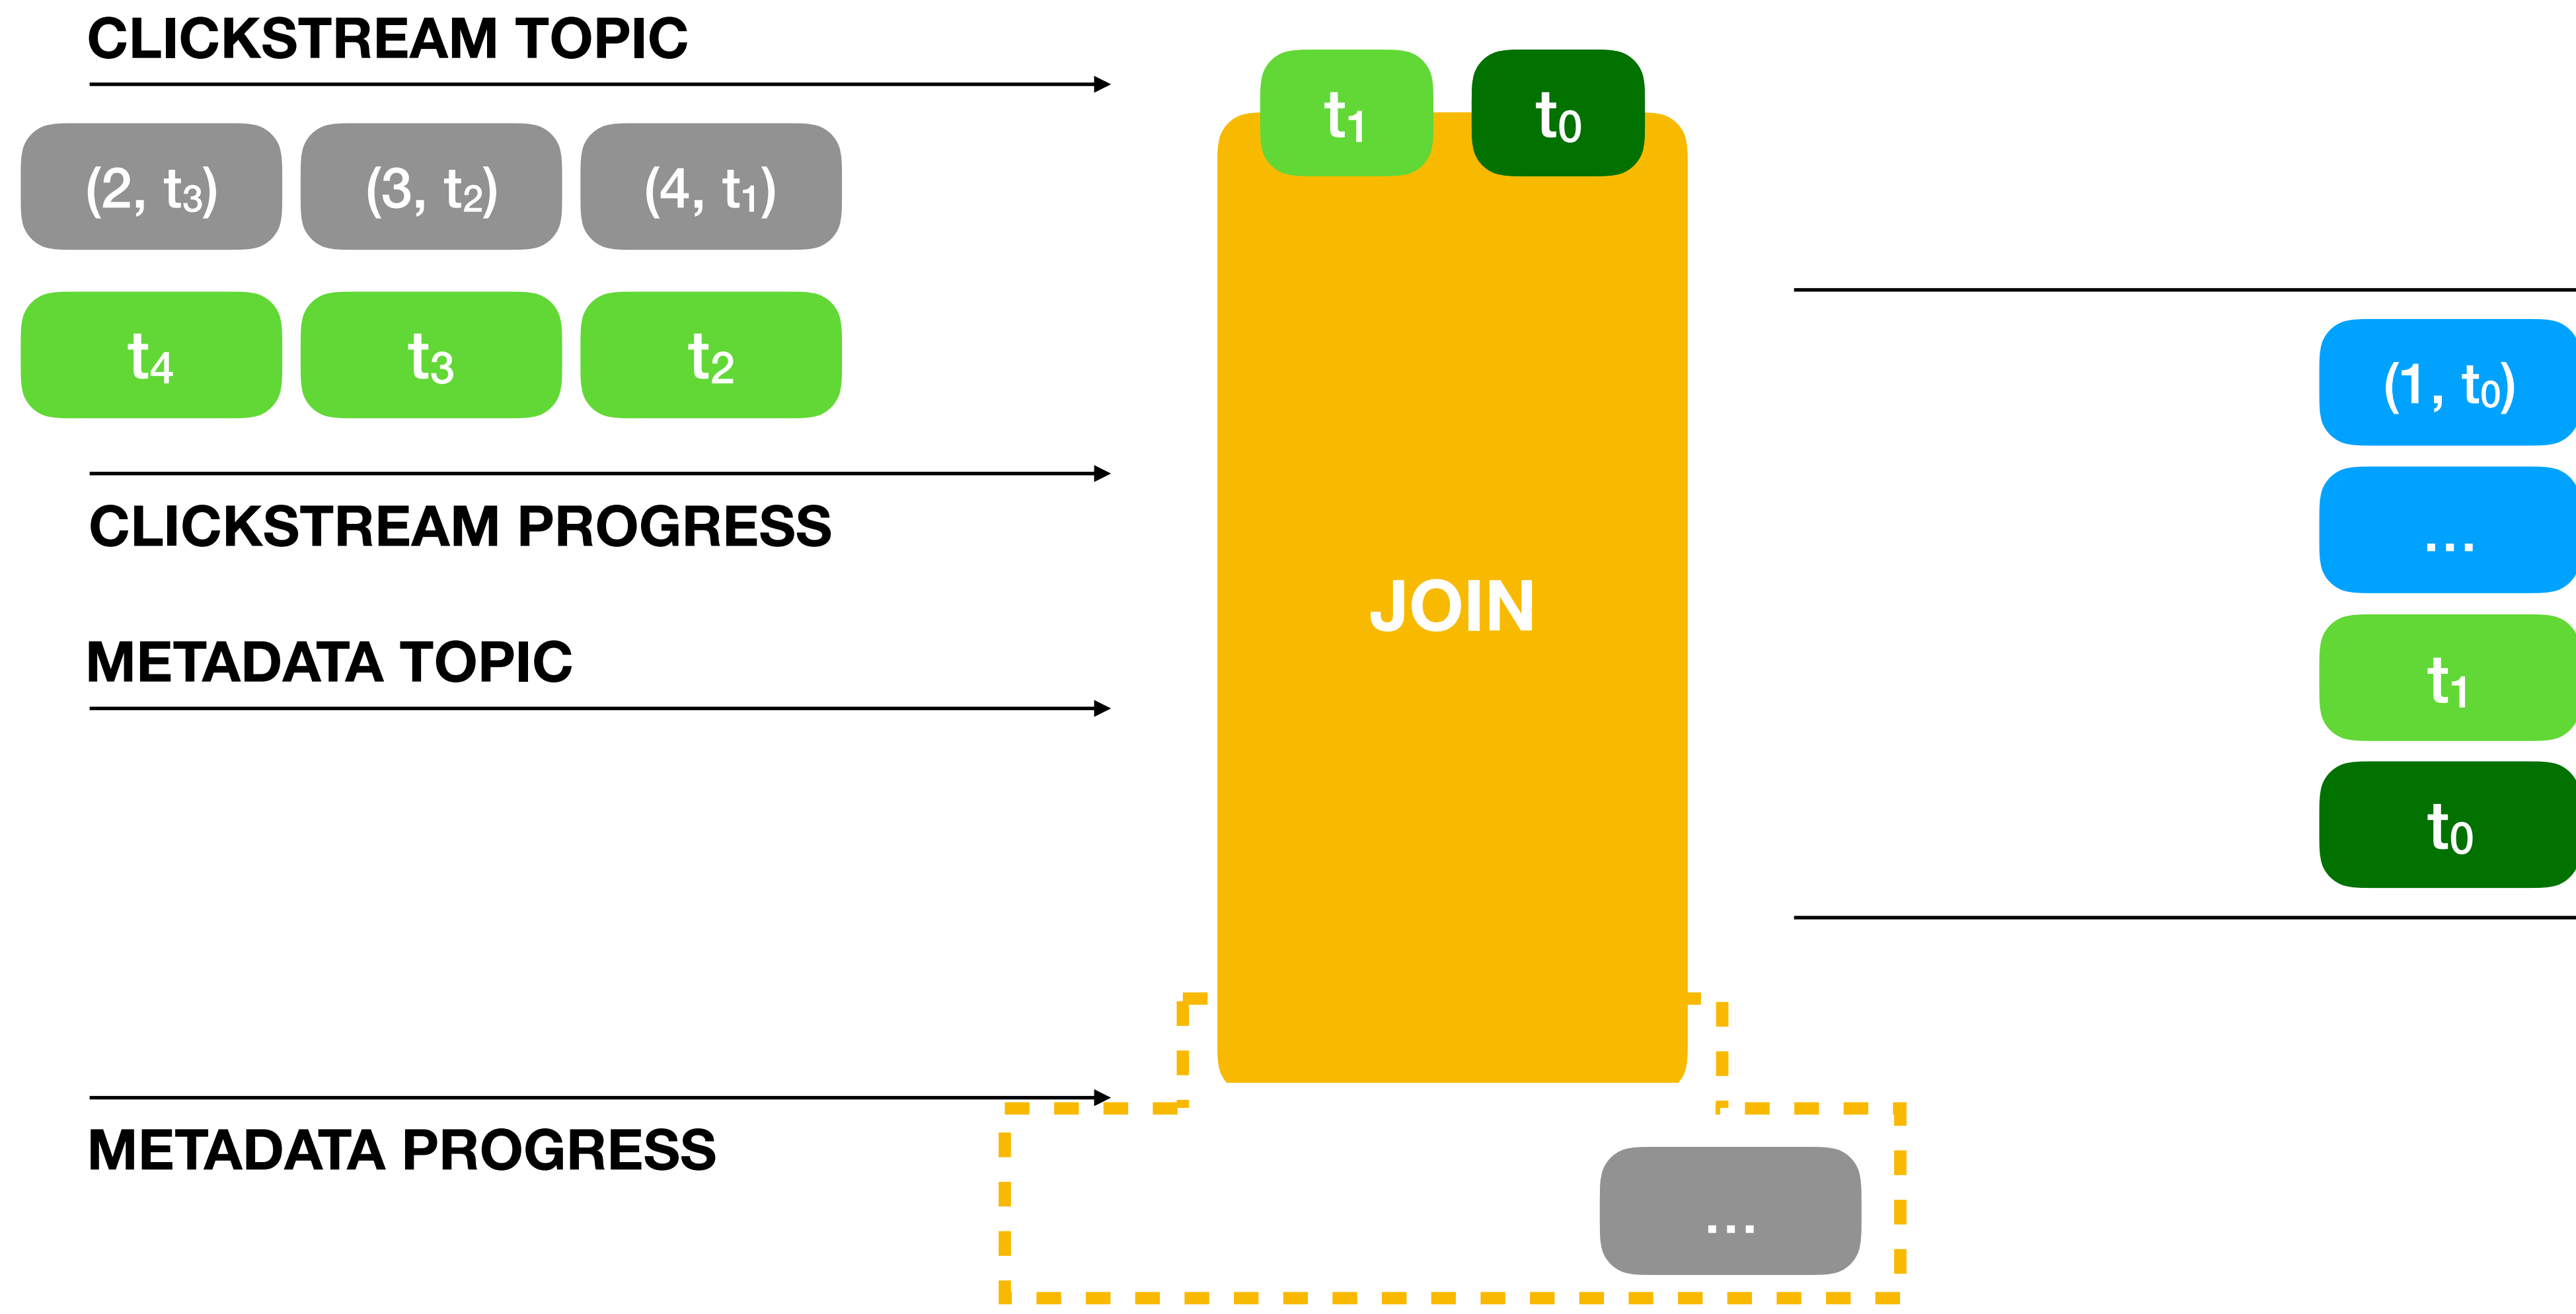
Multidimensional Progress Tracking

(track sources along independent timelines)



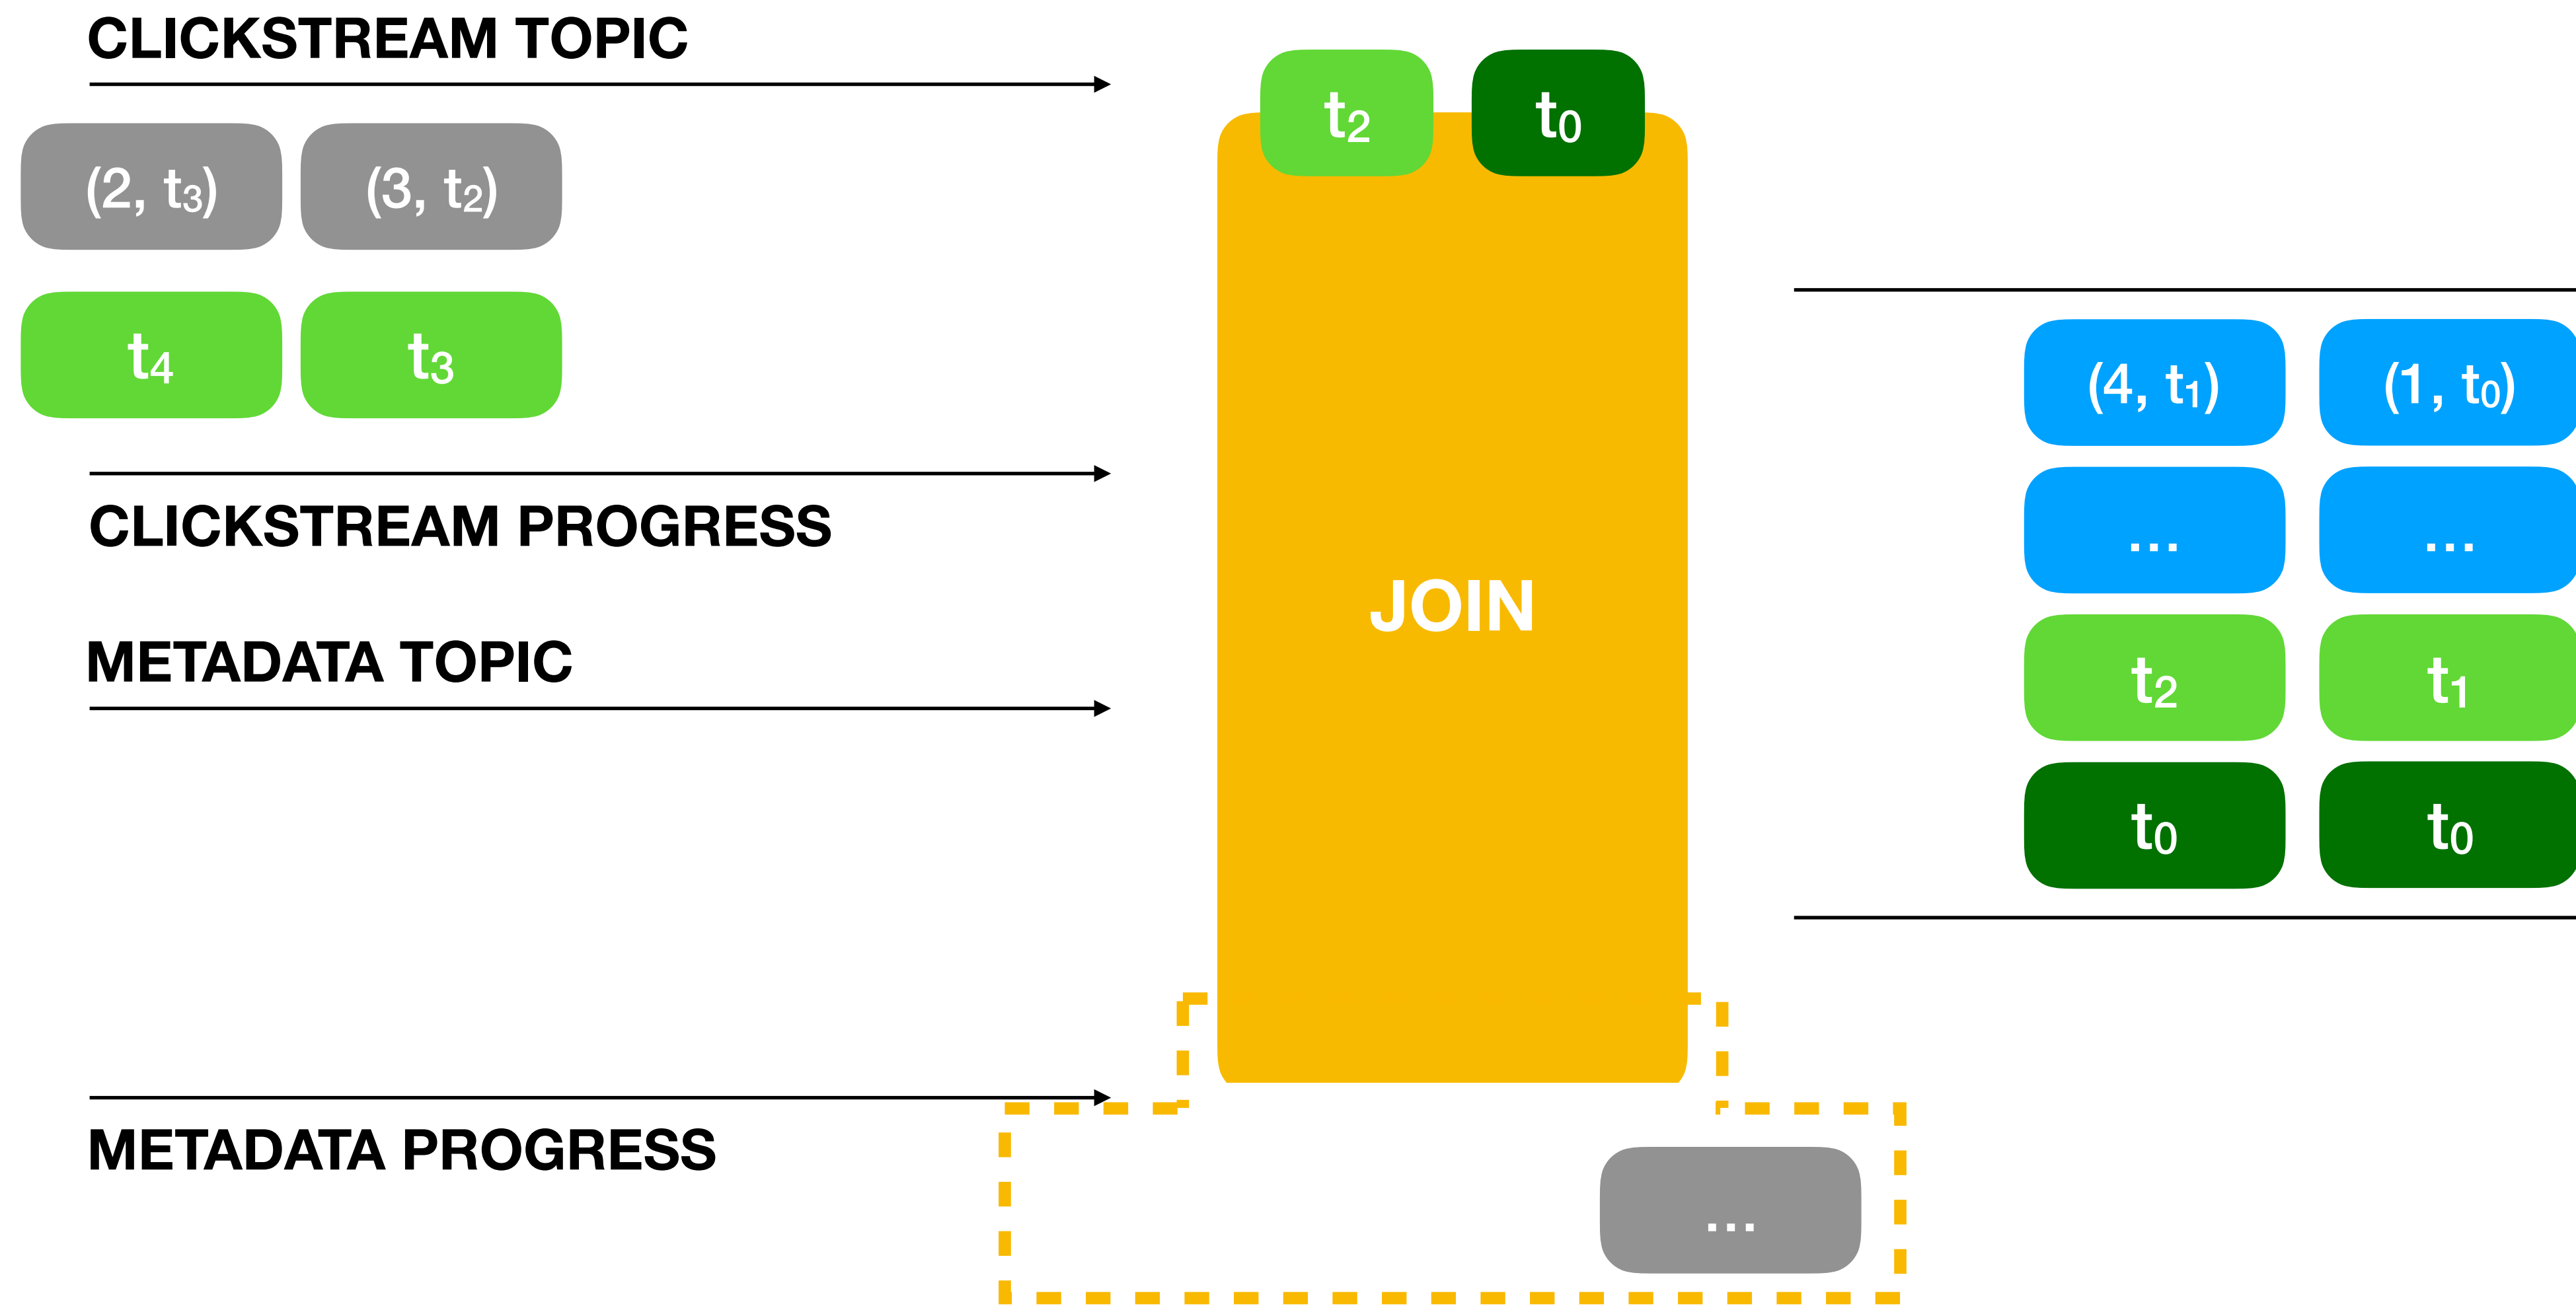
Multidimensional Progress Tracking

(track sources along independent timelines)



Multidimensional Progress Tracking

(track sources along independent timelines)



Creating Dataflows with Timely

```
fn main() {
    timely::execute_from_args(std::env::args(), |worker| {

        // Some computation
        let mut input = InputHandle::new();
        let probe = worker.dataflow(|scope|
            scope.input_from(&mut input)
                .exchange(|x| *x as u64 + 1)
                .inspect(move |x| println!("record {}", x))
                .probe()
        );

        for round in 0..100 {
            if worker.index() == 0 { (0..20).for_each(|i| input.send(i) ) }
            input.advance_to(round + 1);
            while probe.less_than(input.time()) { worker.step(); }
        }
    }).unwrap();
}
```

Creating Dataflows with Timely

```
fn main() {  
    timely::execute_from_args(std::env::args(), |worker| {  
  
        // Some computation  
        let mut input = InputHandle::new();  
        let probe = worker.dataflow(|scope|  
            scope.input_from(&mut input)  
                .exchange(|x| *x as u64 + 1)  
                .inspect(move |x| println!("record {}", x))  
                .probe()  
        );  
  
        for round in 0..100 {  
            if worker.index() == 0 { (0..20).for_each(|i| input.send(i) ) }  
            input.advance_to(round + 1);  
            while probe.less_than(input.time()) { worker.step(); }  
        }  
    }).unwrap();  
}
```

Creating Dataflows with Timely

```
fn main() {
    timely::execute_from_args(std::env::args(), |worker| {

        // Some computation
        let mut input = InputHandle::new();
        let probe = worker.dataflow(|scope|
            scope.input_from(&mut input)
                .exchange(|x| *x as u64 + 1)
                .inspect(move |x| println!("record {}", x))
                .probe()
        );


        for round in 0..100 {
            if worker.index() == 0 { (0..20).for_each(|i| input.send(i) ) }
            input.advance_to(round + 1);
            while probe.less_than(input.time()) { worker.step(); }
        }
    }).unwrap();
}
```


Creating Dataflows with Timely

```
fn main() {
    timely::execute_from_args(std::env::args(), |worker| {

        // Some computation
        let mut input = InputHandle::new();
        let probe = worker.dataflow(|scope|
            scope.input_from(&mut input)
                .exchange(|x| *x as u64 + 1)
                .inspect(move |x| println!("record {}", x))
                .probe()
        );

        for round in 0..100 {
            if worker.index() == 0 { (0..20).for_each(|i| input.send(i) ) }
            input.advance_to(round + 1);
            while probe.less_than(input.time()) { worker.step(); }
        }
    }).unwrap();
}
```



Running Dataflows with Timely


```
fn main() {  
    timely::execute_from_args(std::env::args(), |worker| {
```

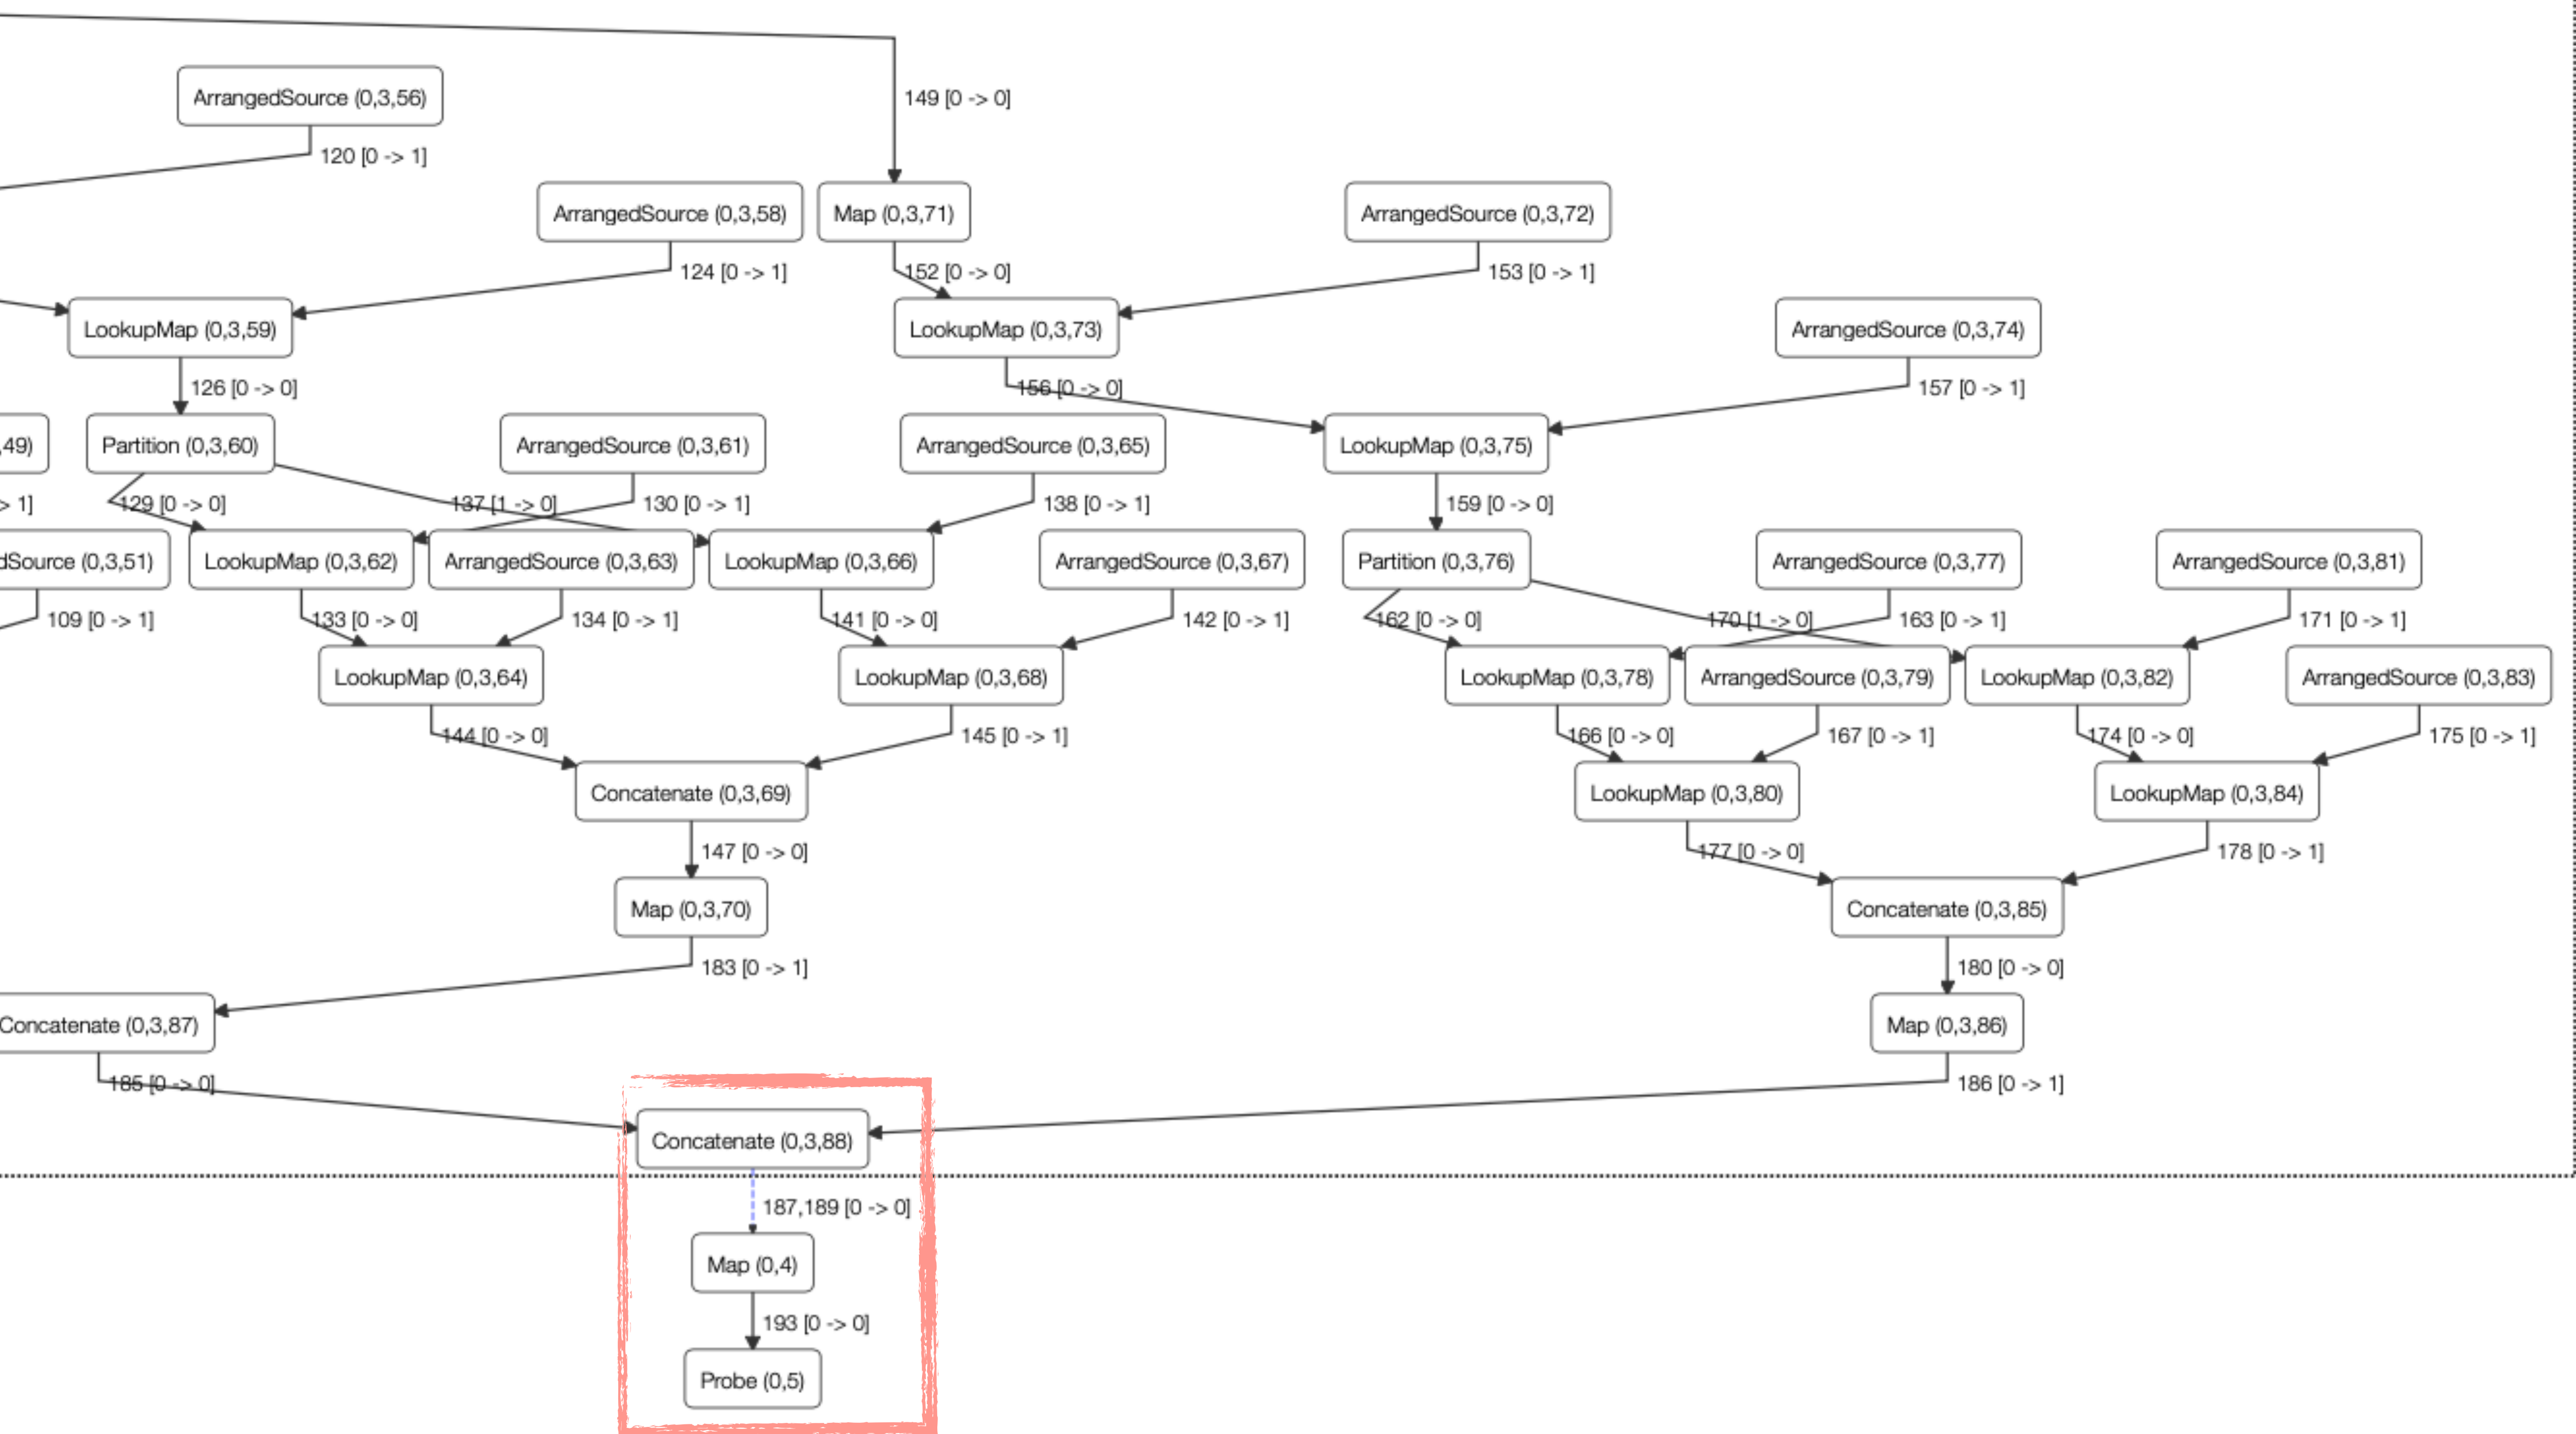
```
        // Some computation  
        let mut input = InputHandle::new();  
        let probe = worker.dataflow(|scope|  
            scope.input_from(&mut input)  
                .exchange(|x| *x as u64 + 1)  
                .inspect(move |x| println!("record {}", x))  
                .probe()  
        );
```

```
        for round in 0..100 {  
            if worker.index() == 0 { (0..20).for_each(|i| input.send(i)) }  
            input.advance_to(round + 1);  
            while probe.less_than(input.time()) { worker.step(); }  
        }
```

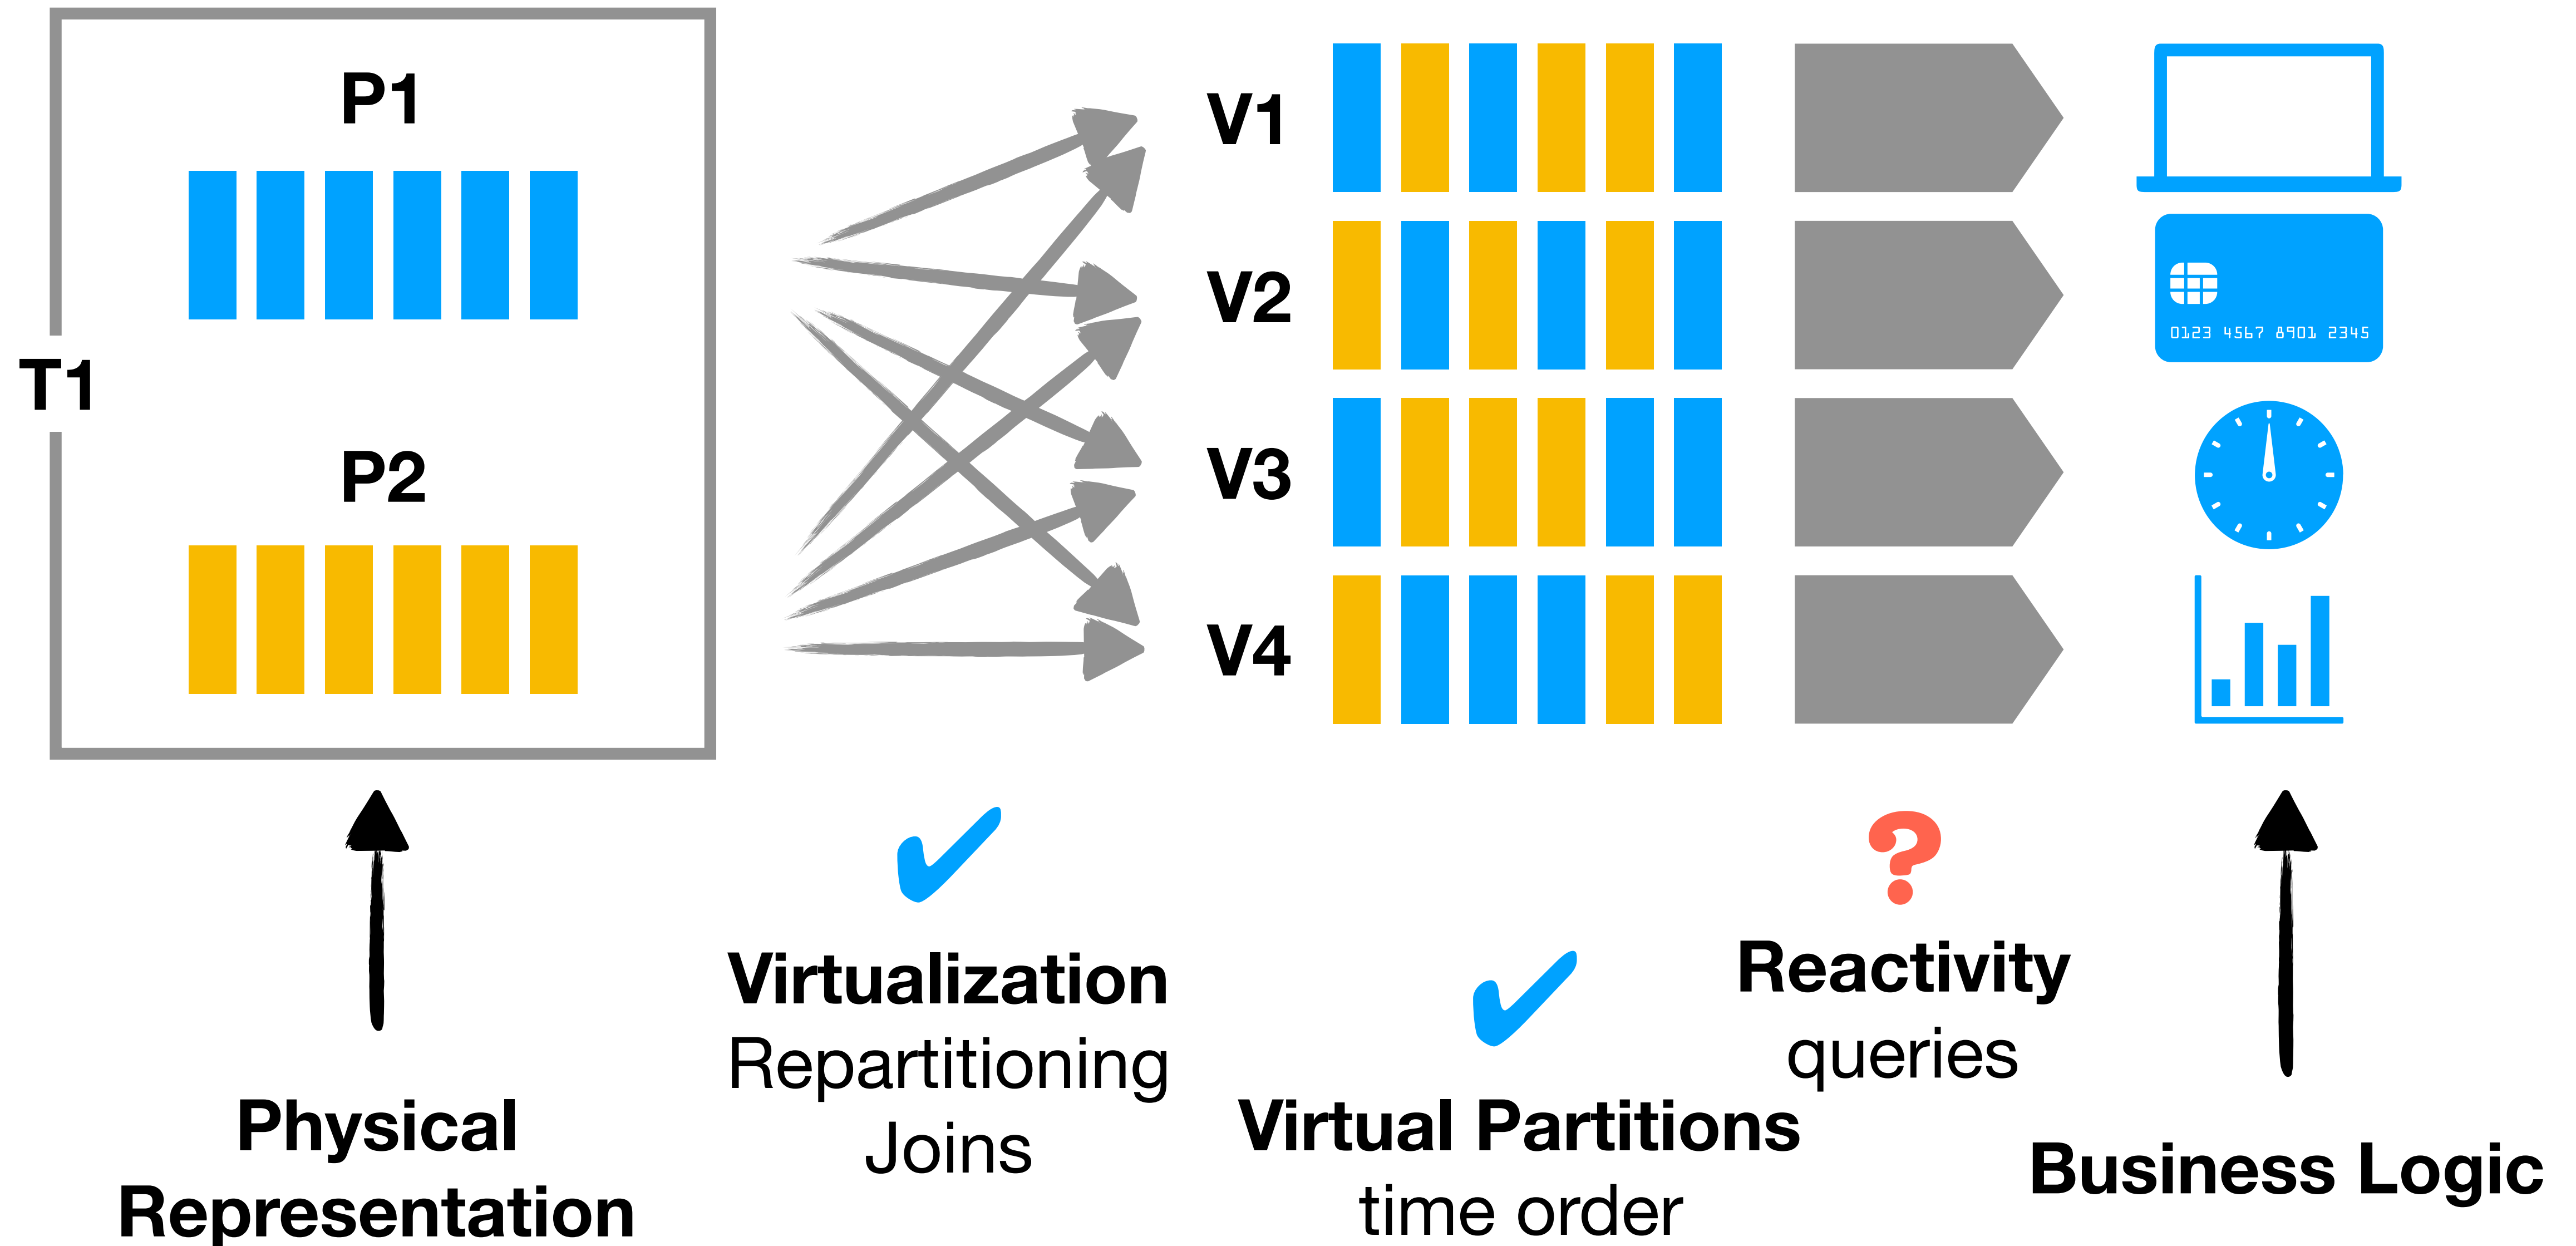
```
    }).unwrap();
```

```
}
```

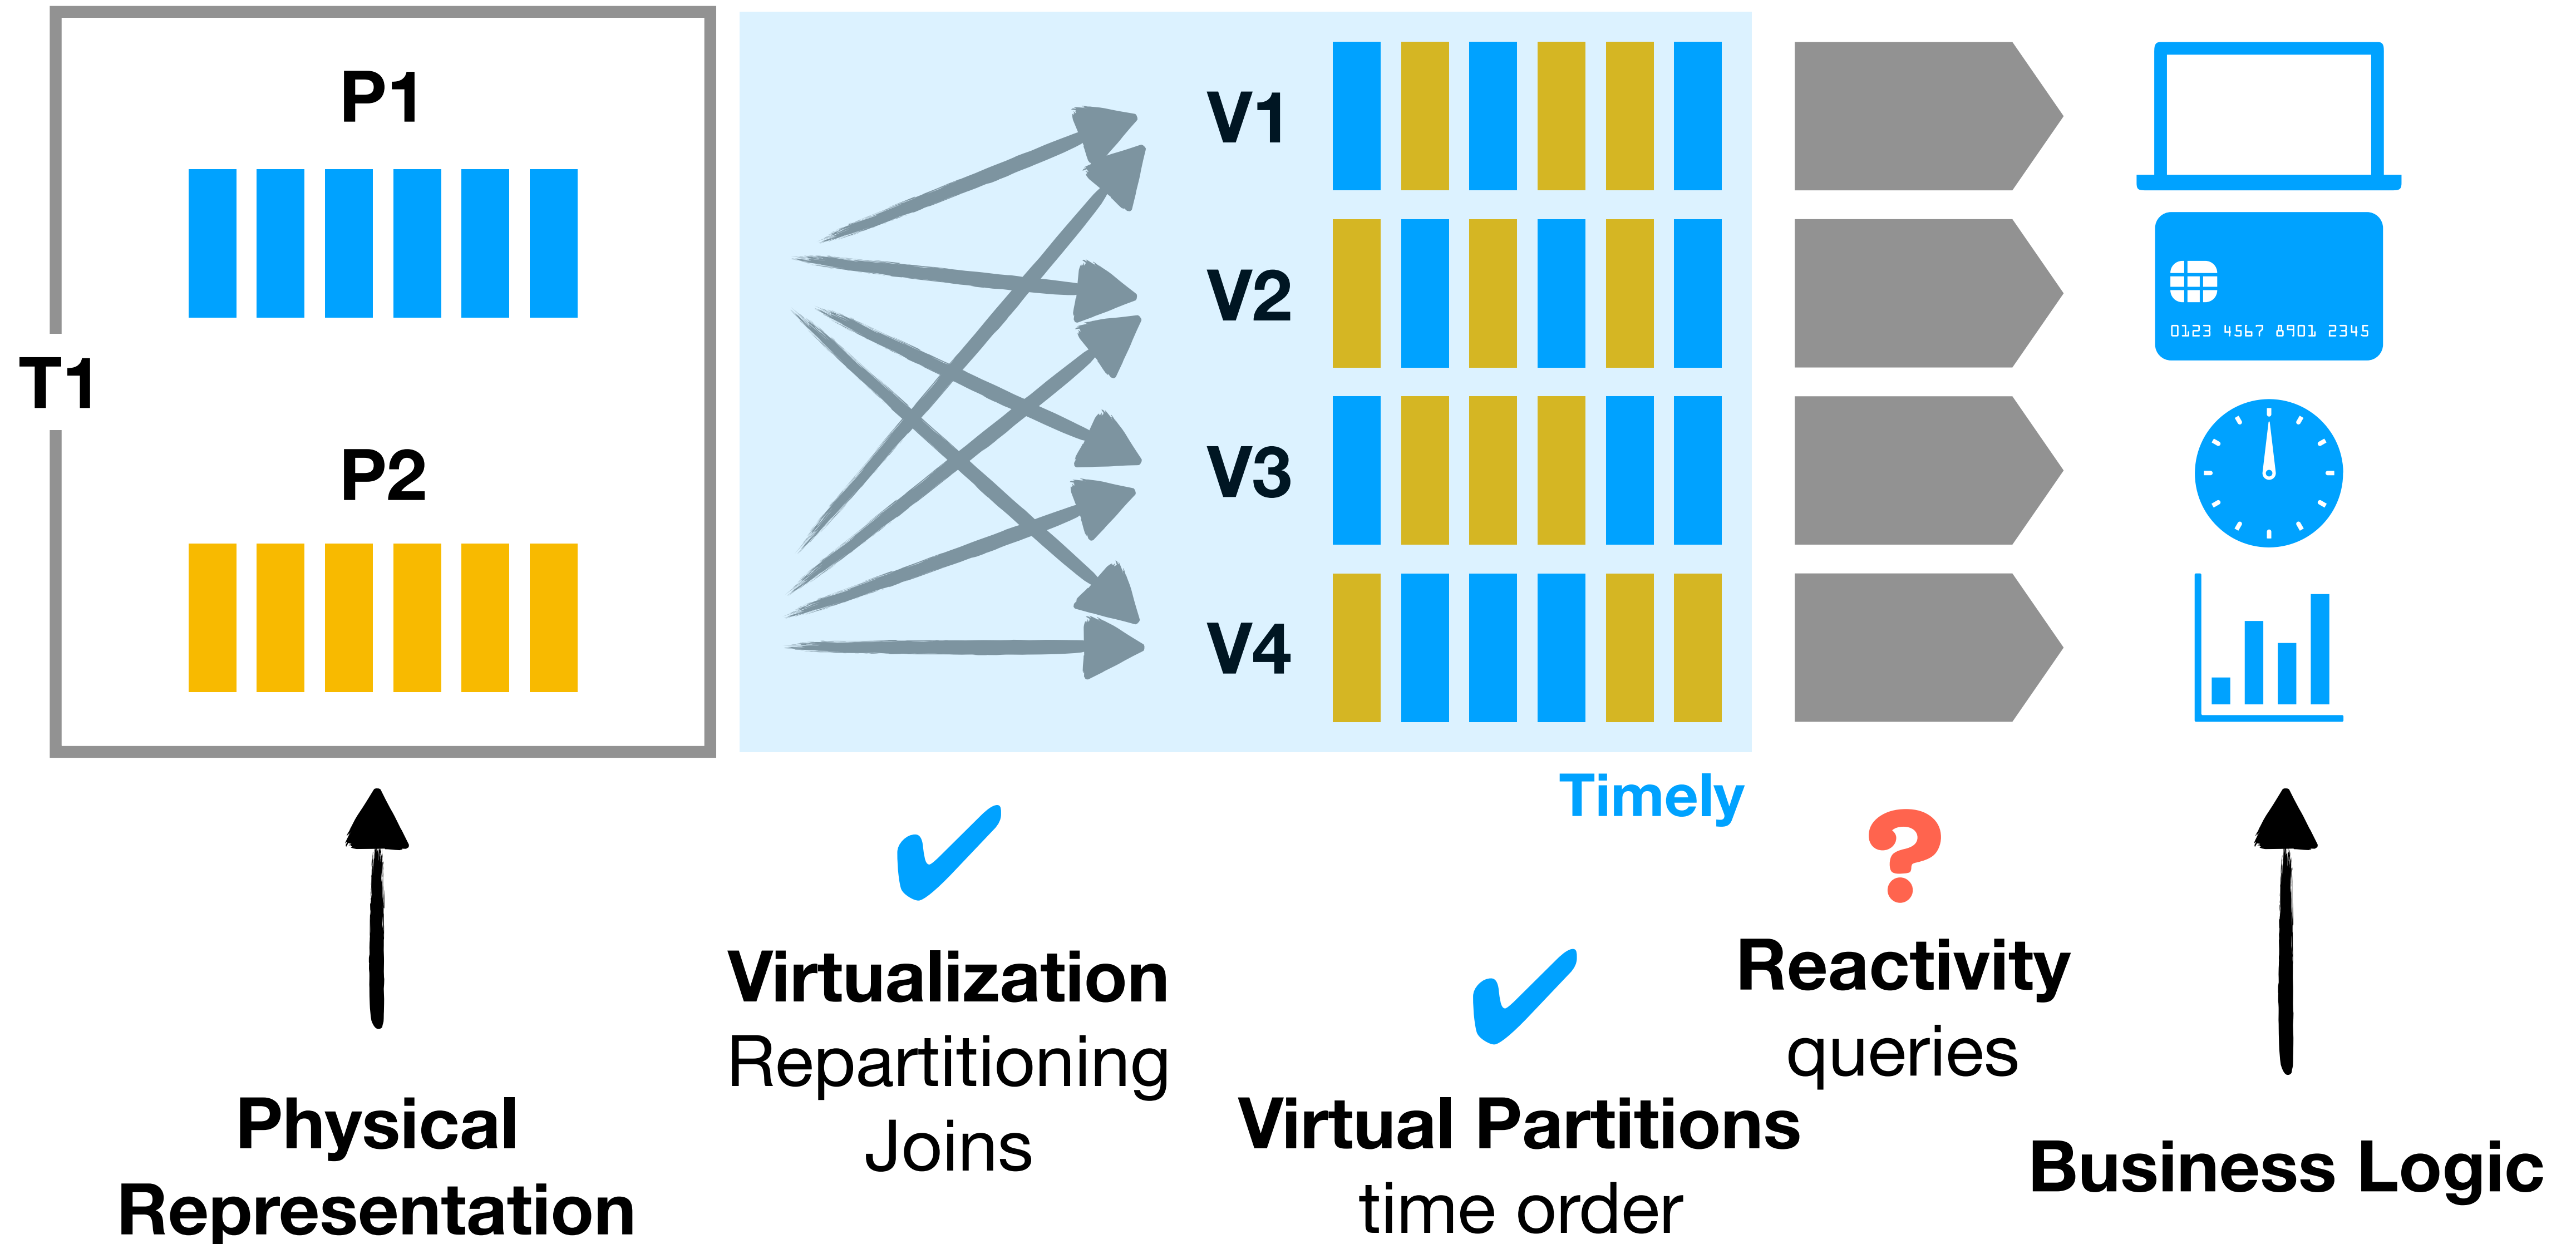





Kafka Superpowers




Kafka Superpowers




The Trifecta?




The Trifecta?




Recursive Graph Traversal



Recursive Graph Traversal



Recursive Graph Traversal



Recursive Dataflows


```
/// Breadth-First Search

let nodes = roots.map(|x| (x, 0));

nodes.iterate(|inner| {

  let edges = edges.enter(&inner.scope());
  let nodes = nodes.enter(&inner.scope());

  inner.join(&edges, |_k,l,d| (*d, l+1))
    .concat(&nodes)
    .reduce(|_, s, t| t.push((*s[0].0, 1)))
})
```



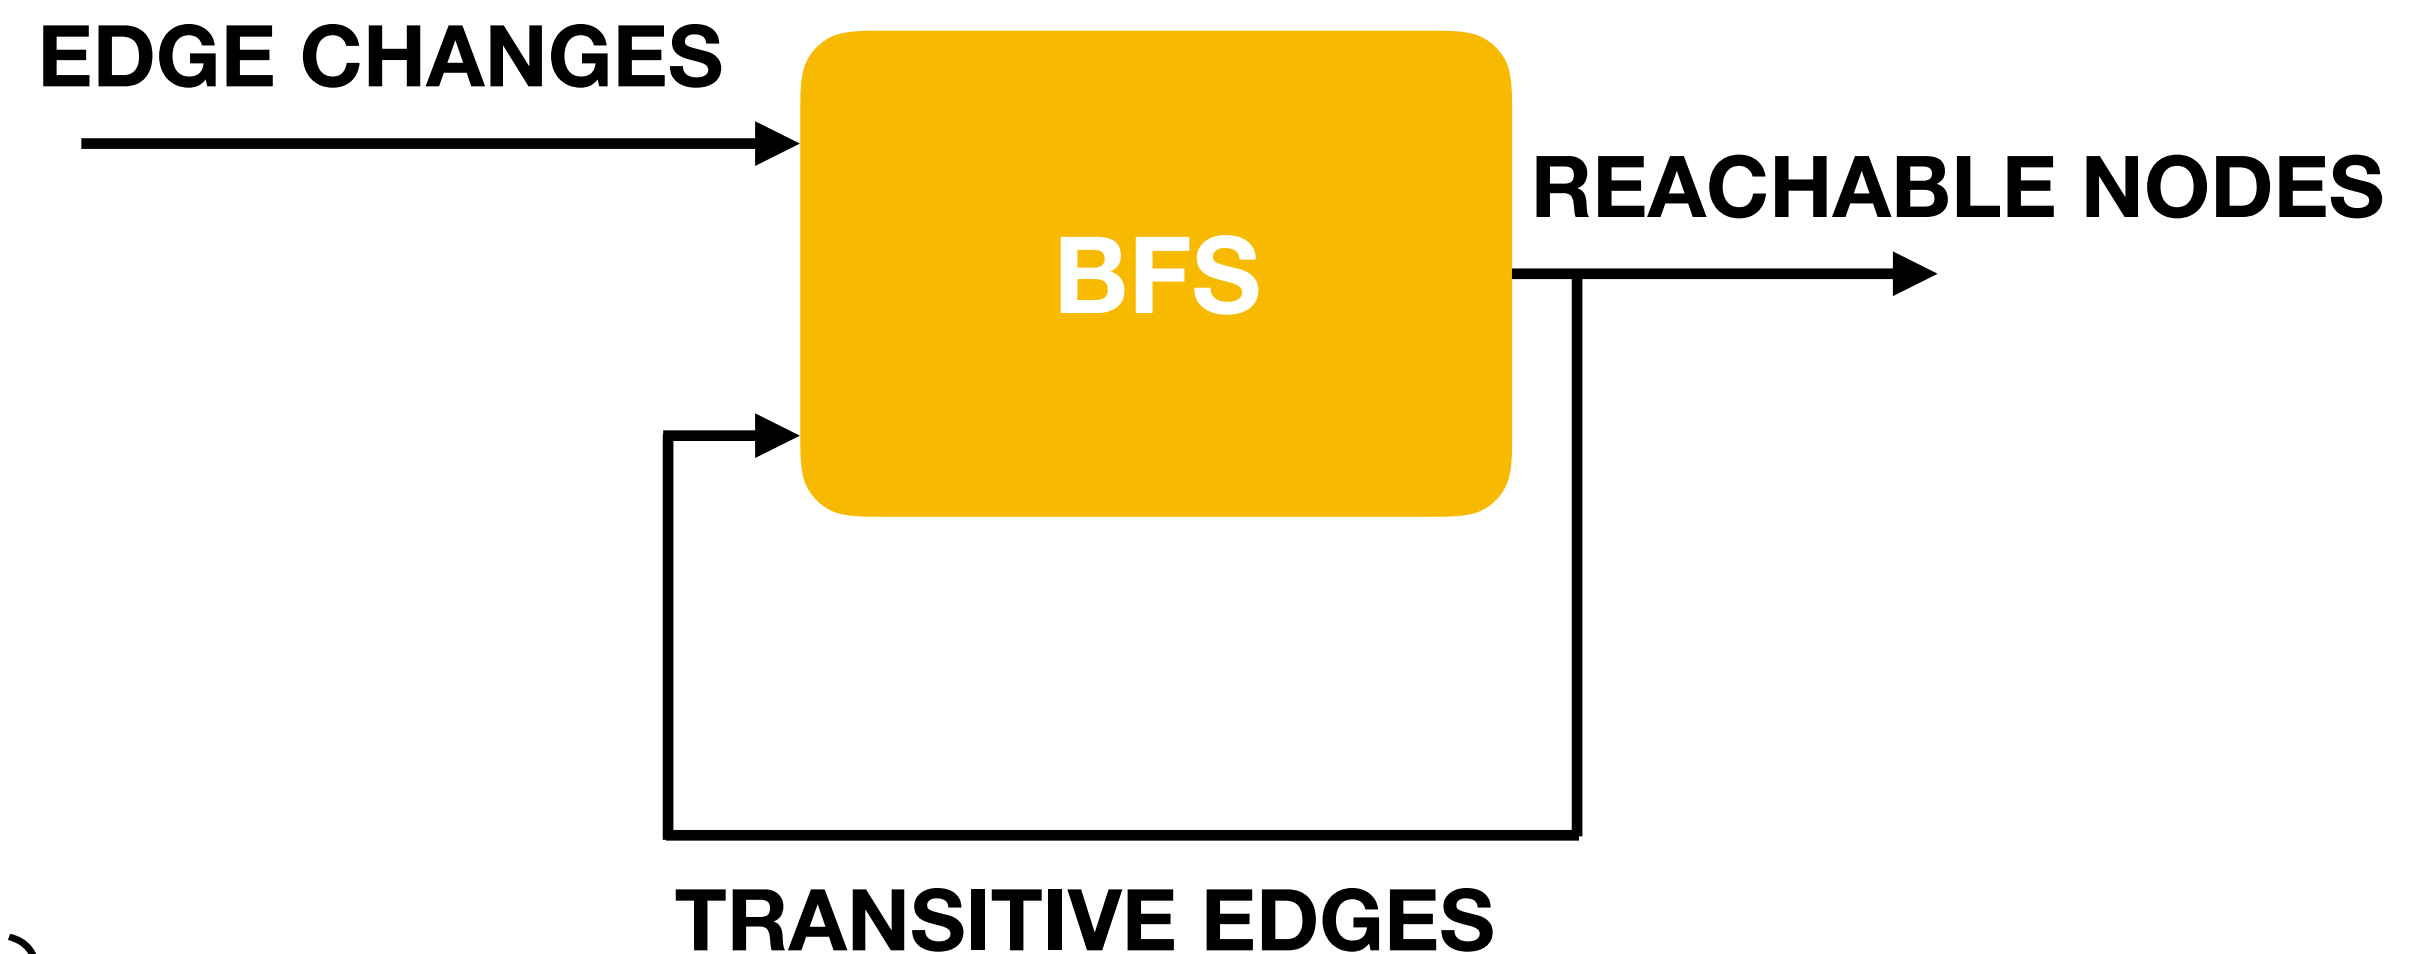
Recursive Dataflows

```
/// Breadth-First Search

let nodes = roots.map(|x| (x, 0));
nodes.iterate(|inner| {

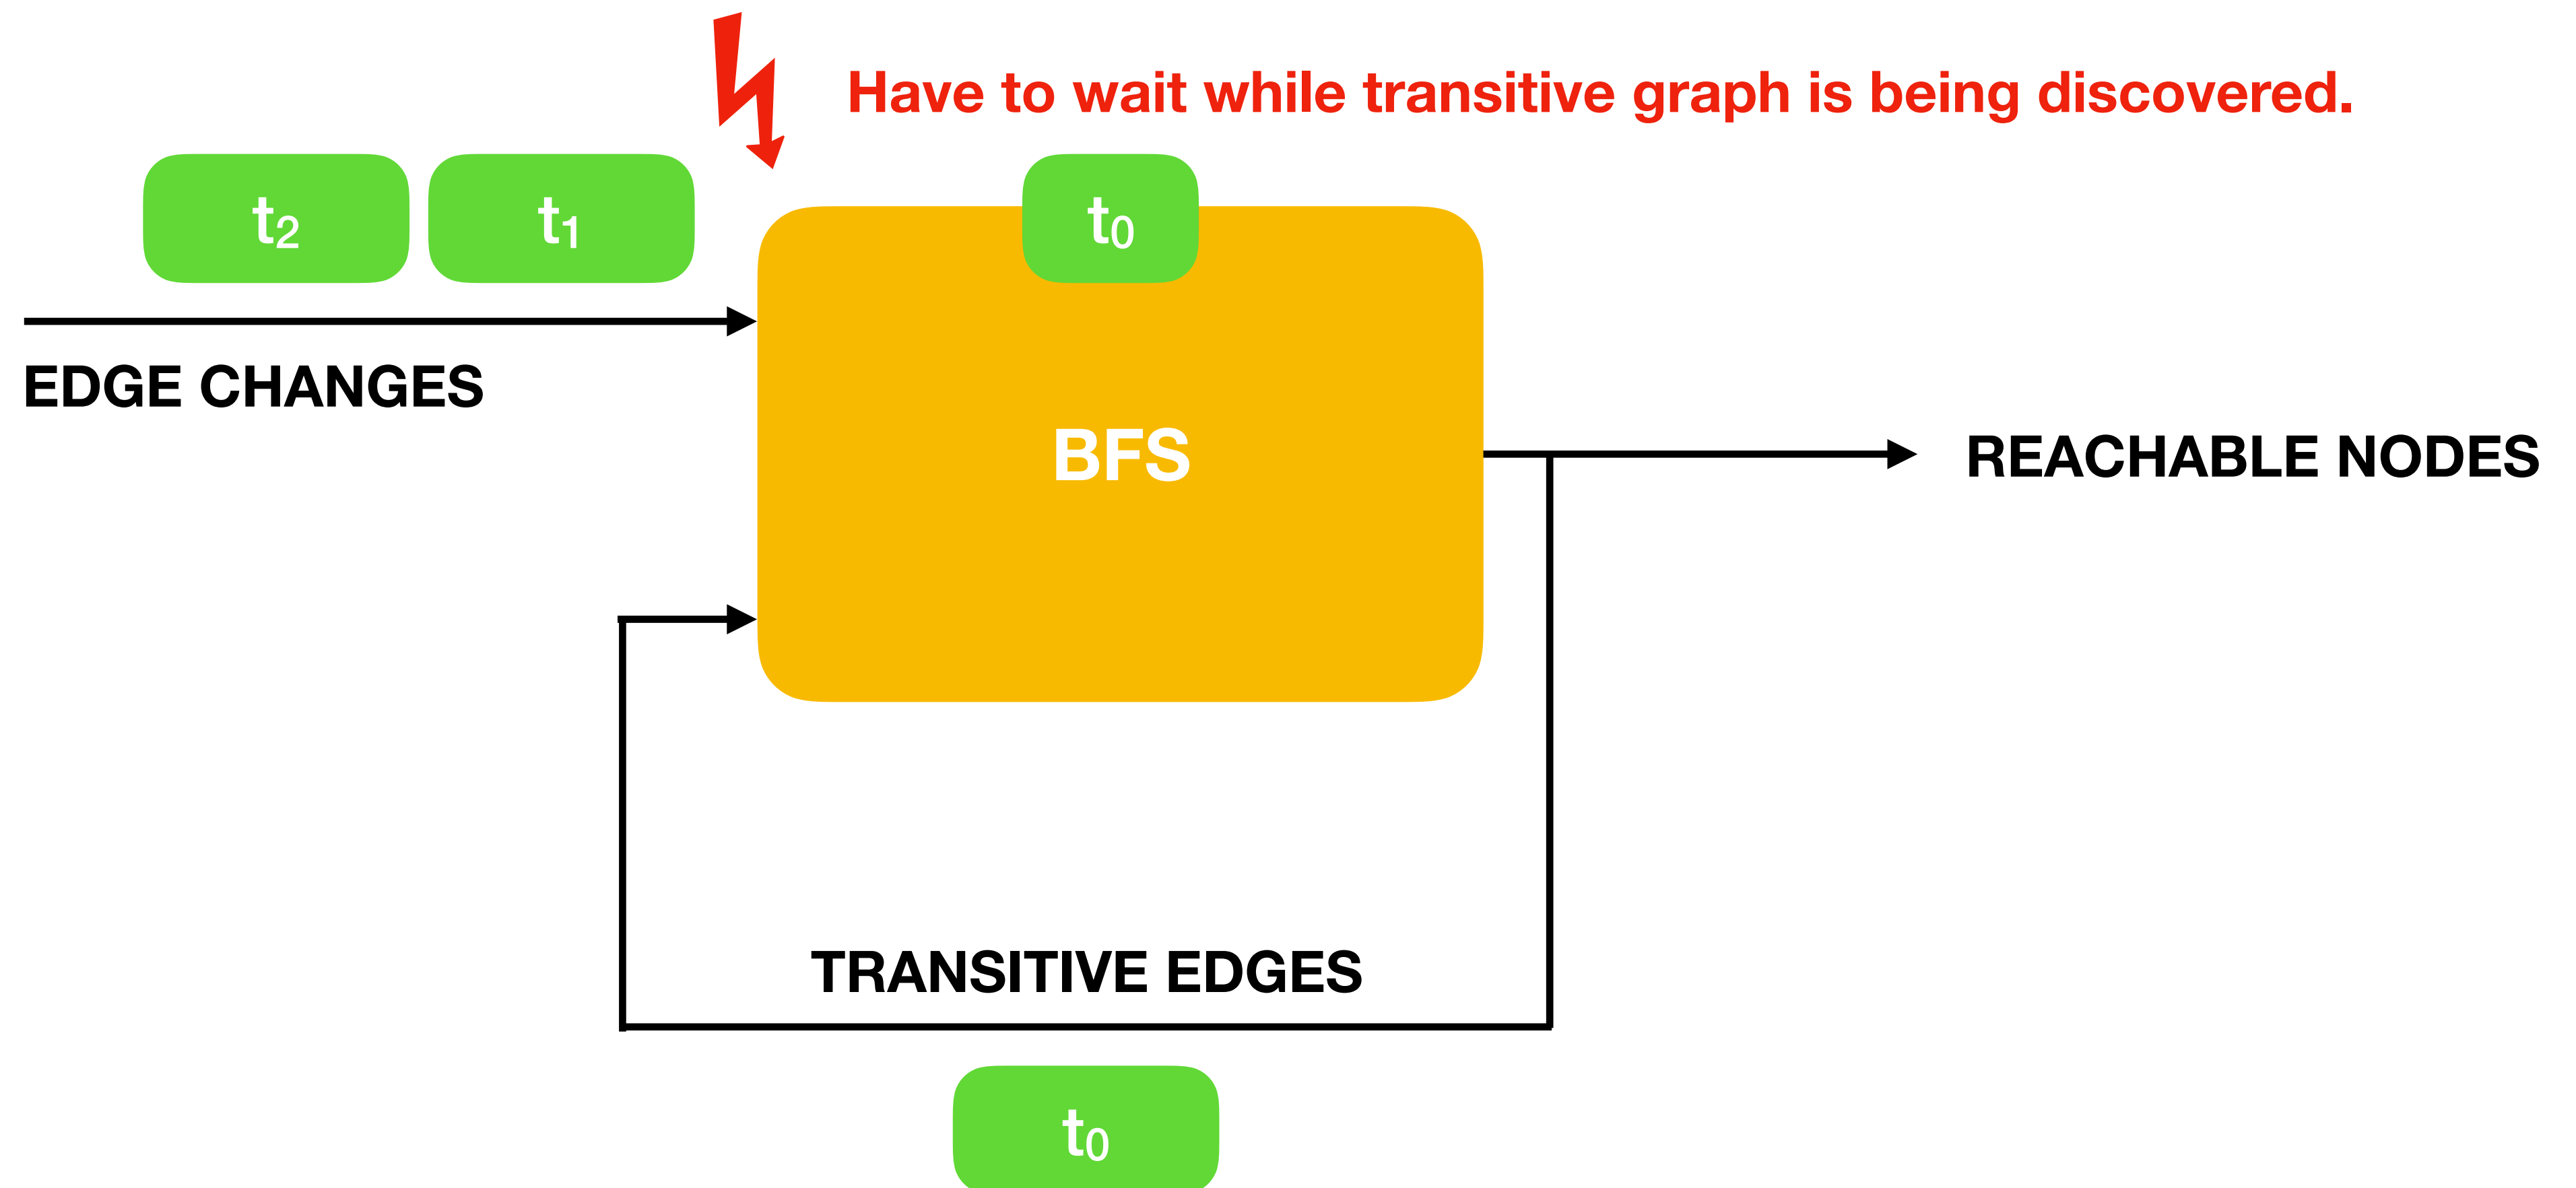
  let edges = edges.enter(&inner.scope());
  let nodes = nodes.enter(&inner.scope());

  inner.join(&edges, |_k,l,d| (*d, l+1))
    .concat(&nodes)
    .reduce(|_, s, t| t.push((*s[0].0, 1)))
})
```



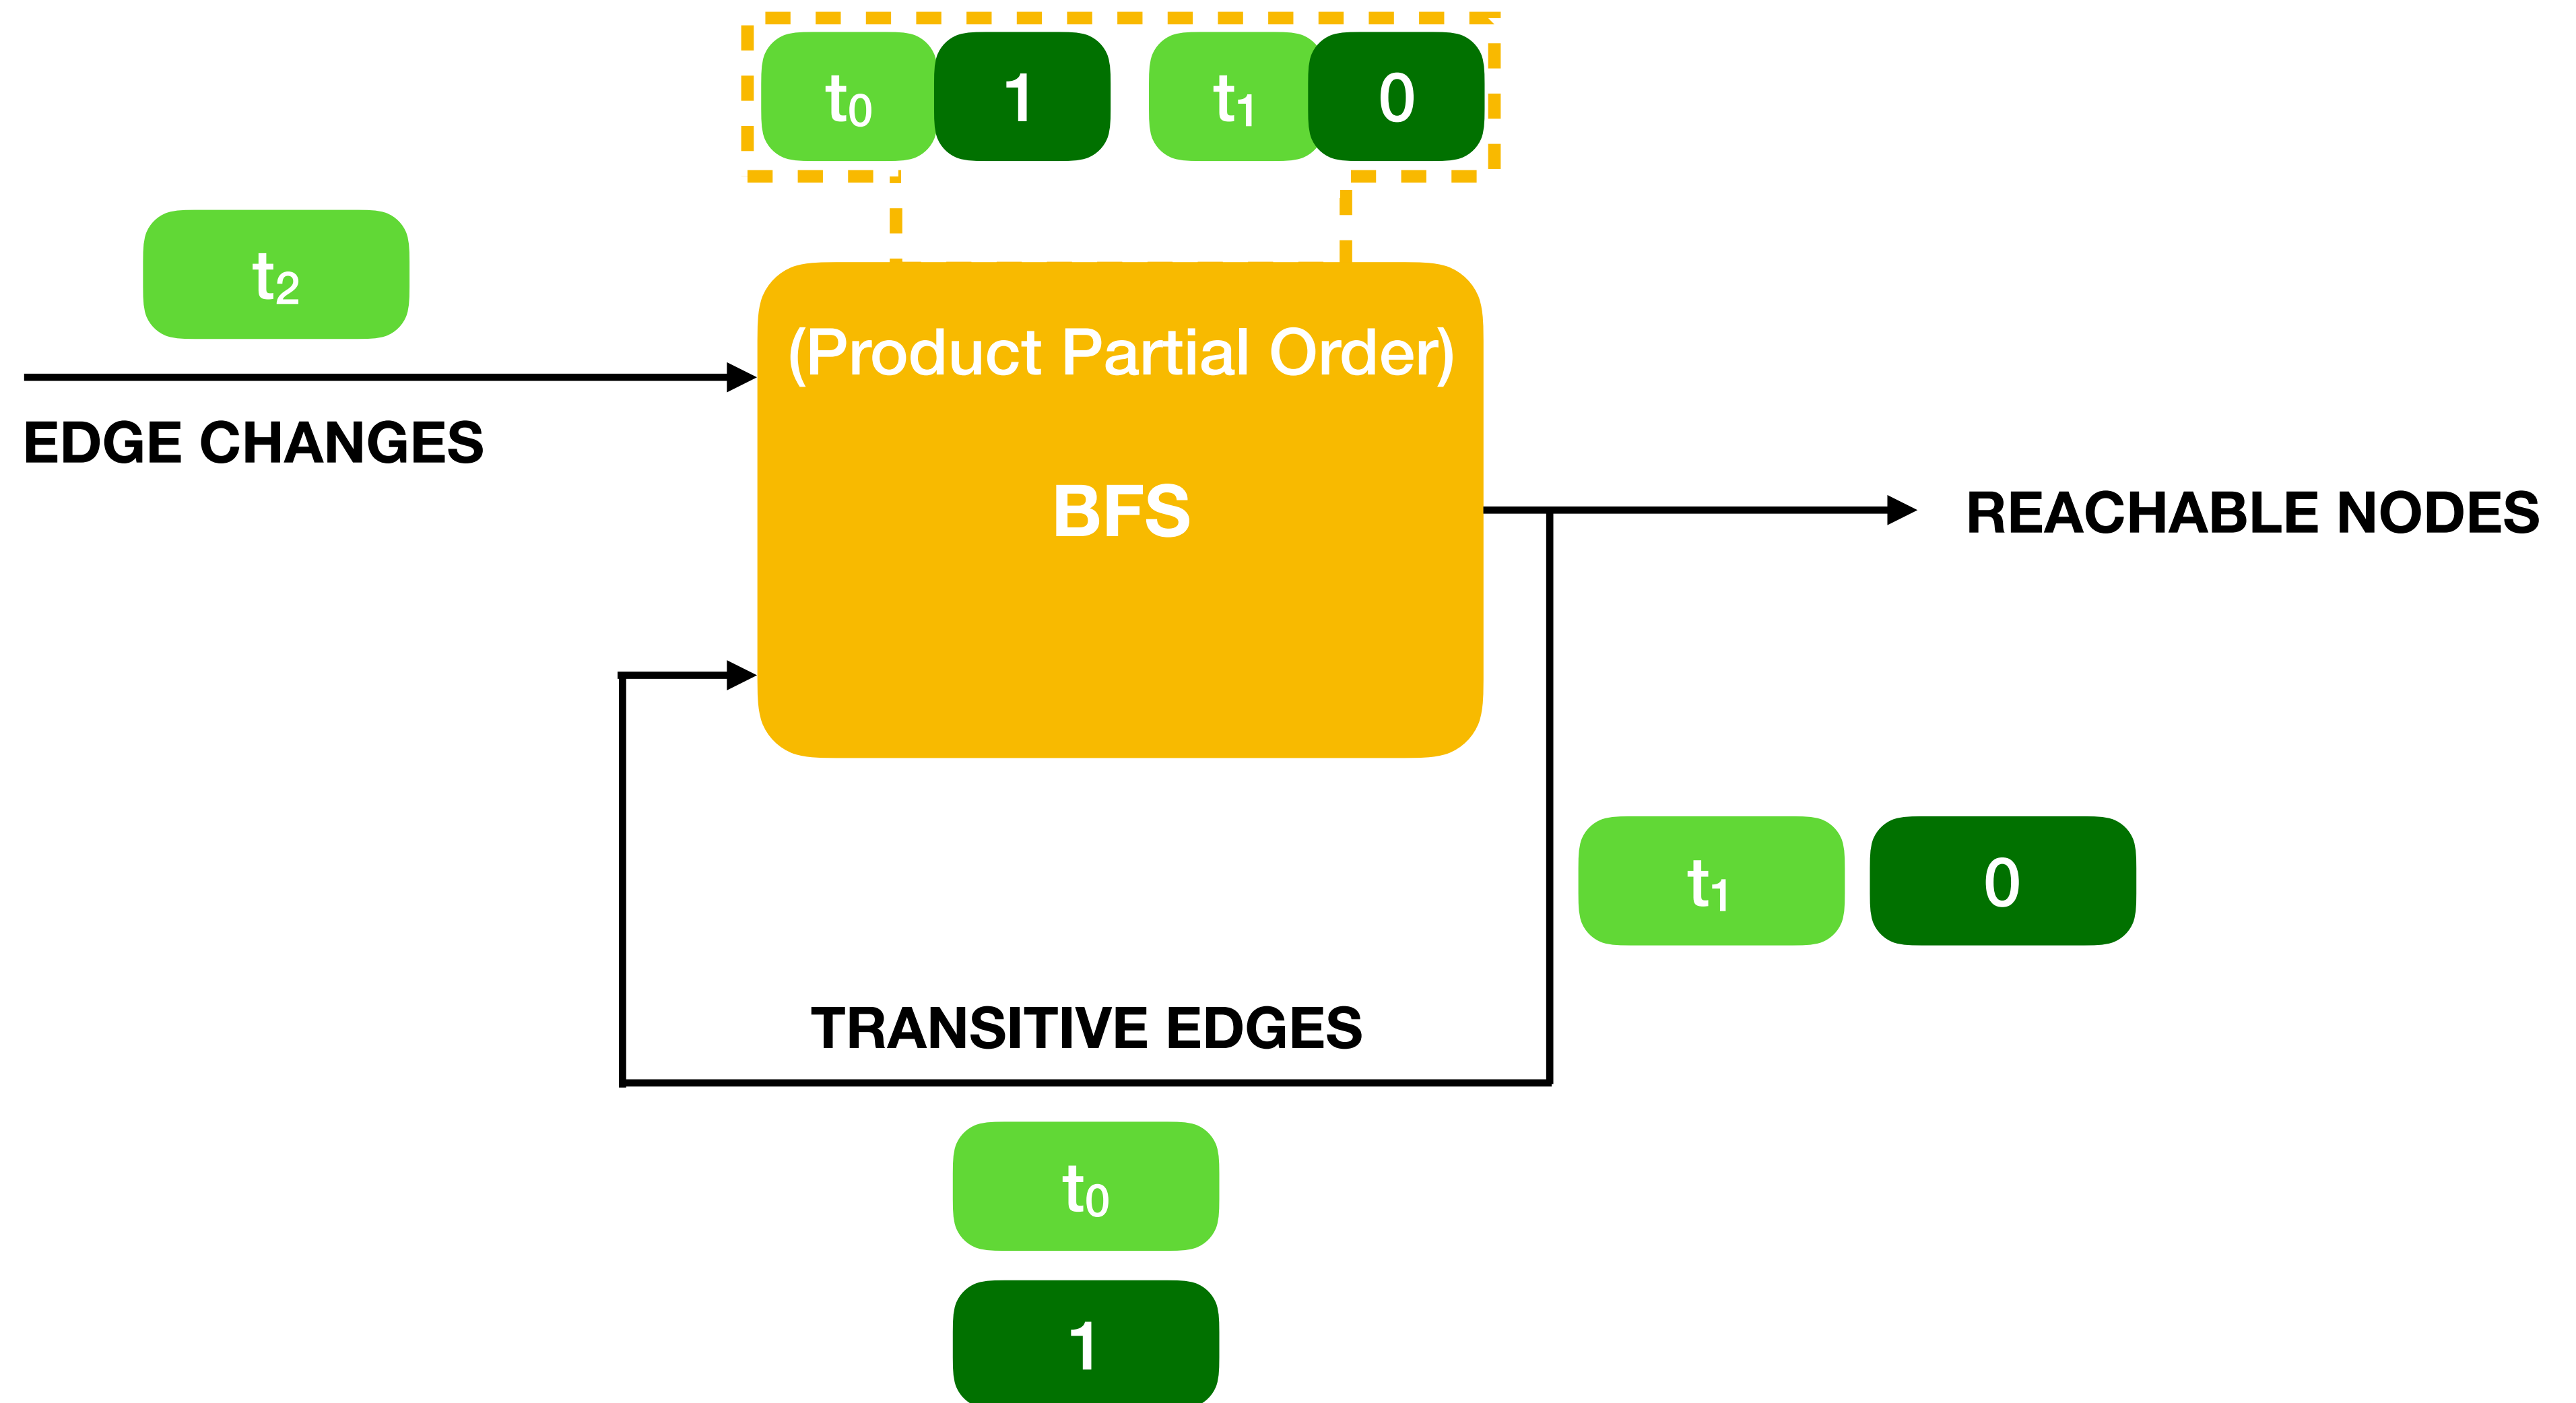
Progress Tracking... with Loops?

(have to finish iterating before we can handle next input)



Multidimensional Progress Tracking

(track iteration depth separately)



Lexicographical Order (Join)

(visibility for t₂ t₂)

	t ₀	t ₁	t ₂	t ₃
t ₀	✓	✓	✓	✓
t ₁	✓	✓	✓	✓
t ₂	✓	✓	✓	✗
t ₃	✗	✗	✗	✗


Product Partial Order (Iteration)

(visibility for t_2 2)

	0	1	2	3
t_0	✓	✓	✓	✗
t_1	✓	✓	✓	✗
t_2	✓	✓	✓	✗
t_3	✗	✗	✗	✗

Multidimensional Progress Tracking

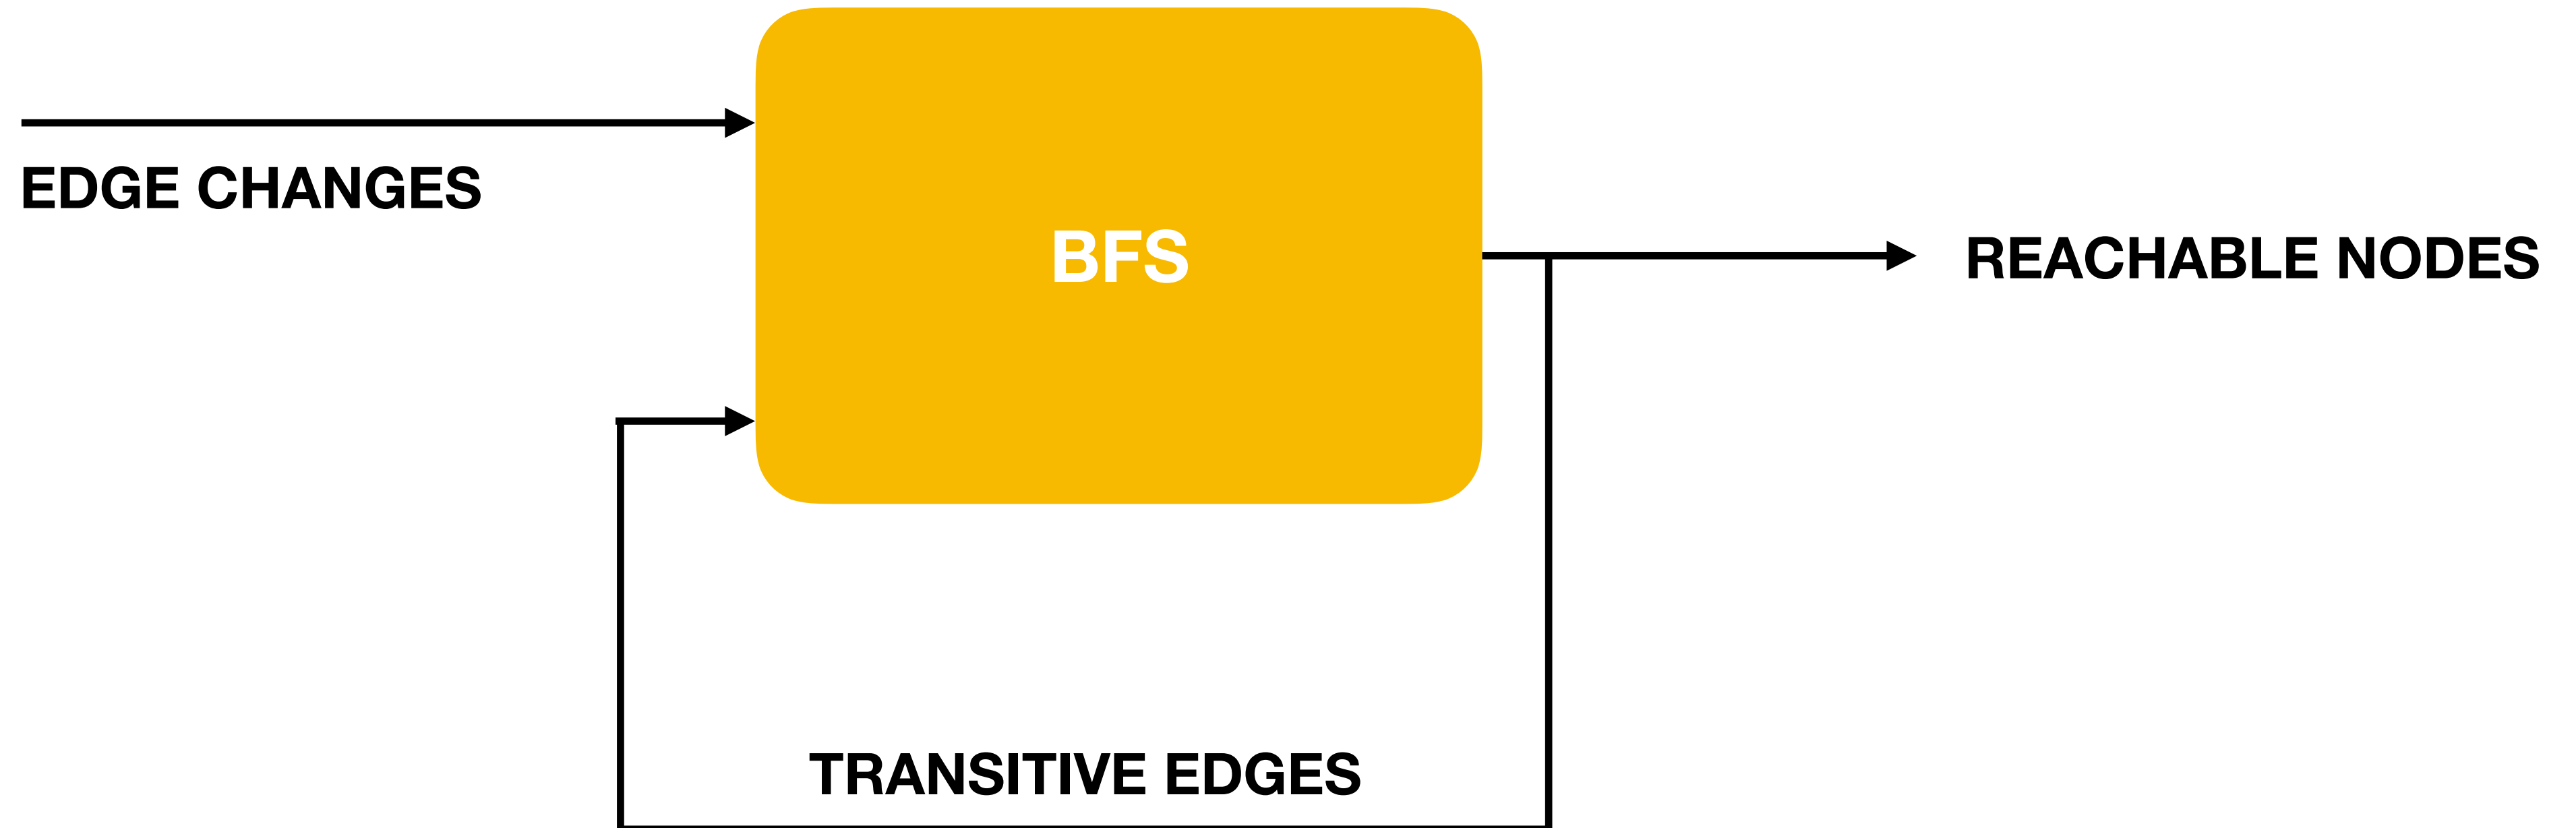
(track iteration depth separately)



Incremental Execution?



Have to start from scratch for every transaction?



Differential Dataflow

Iterative, incrementalized operators for Timely

github.com/TimelyDataflow

Performance

Connected	cores	livejournal	orkut
GraphX	128	59s	53s
Socialite	128	54s	78s
Myria	128	37s	57s
BigDatalog	128	27s	33s
Differential	1, 2	20s, 11s	43s, 26s
update	1, 2	98us, 109us	200us, 216us

Streaming & Relational Queries

Declarative Differential Dataflows (3DF)


```
/// BFS

let nodes = roots.map(|x| (x, 0));

nodes.iterate(|inner| {

  let edges = edges.enter(&inner.scope());
  let nodes = nodes.enter(&inner.scope());


  inner.join_map(&edges, |_k,l,d| (*d, l+1))
    .concat(&nodes)
    .reduce(|_, s, t| t.push((*s[0].0, 1)))
})
```



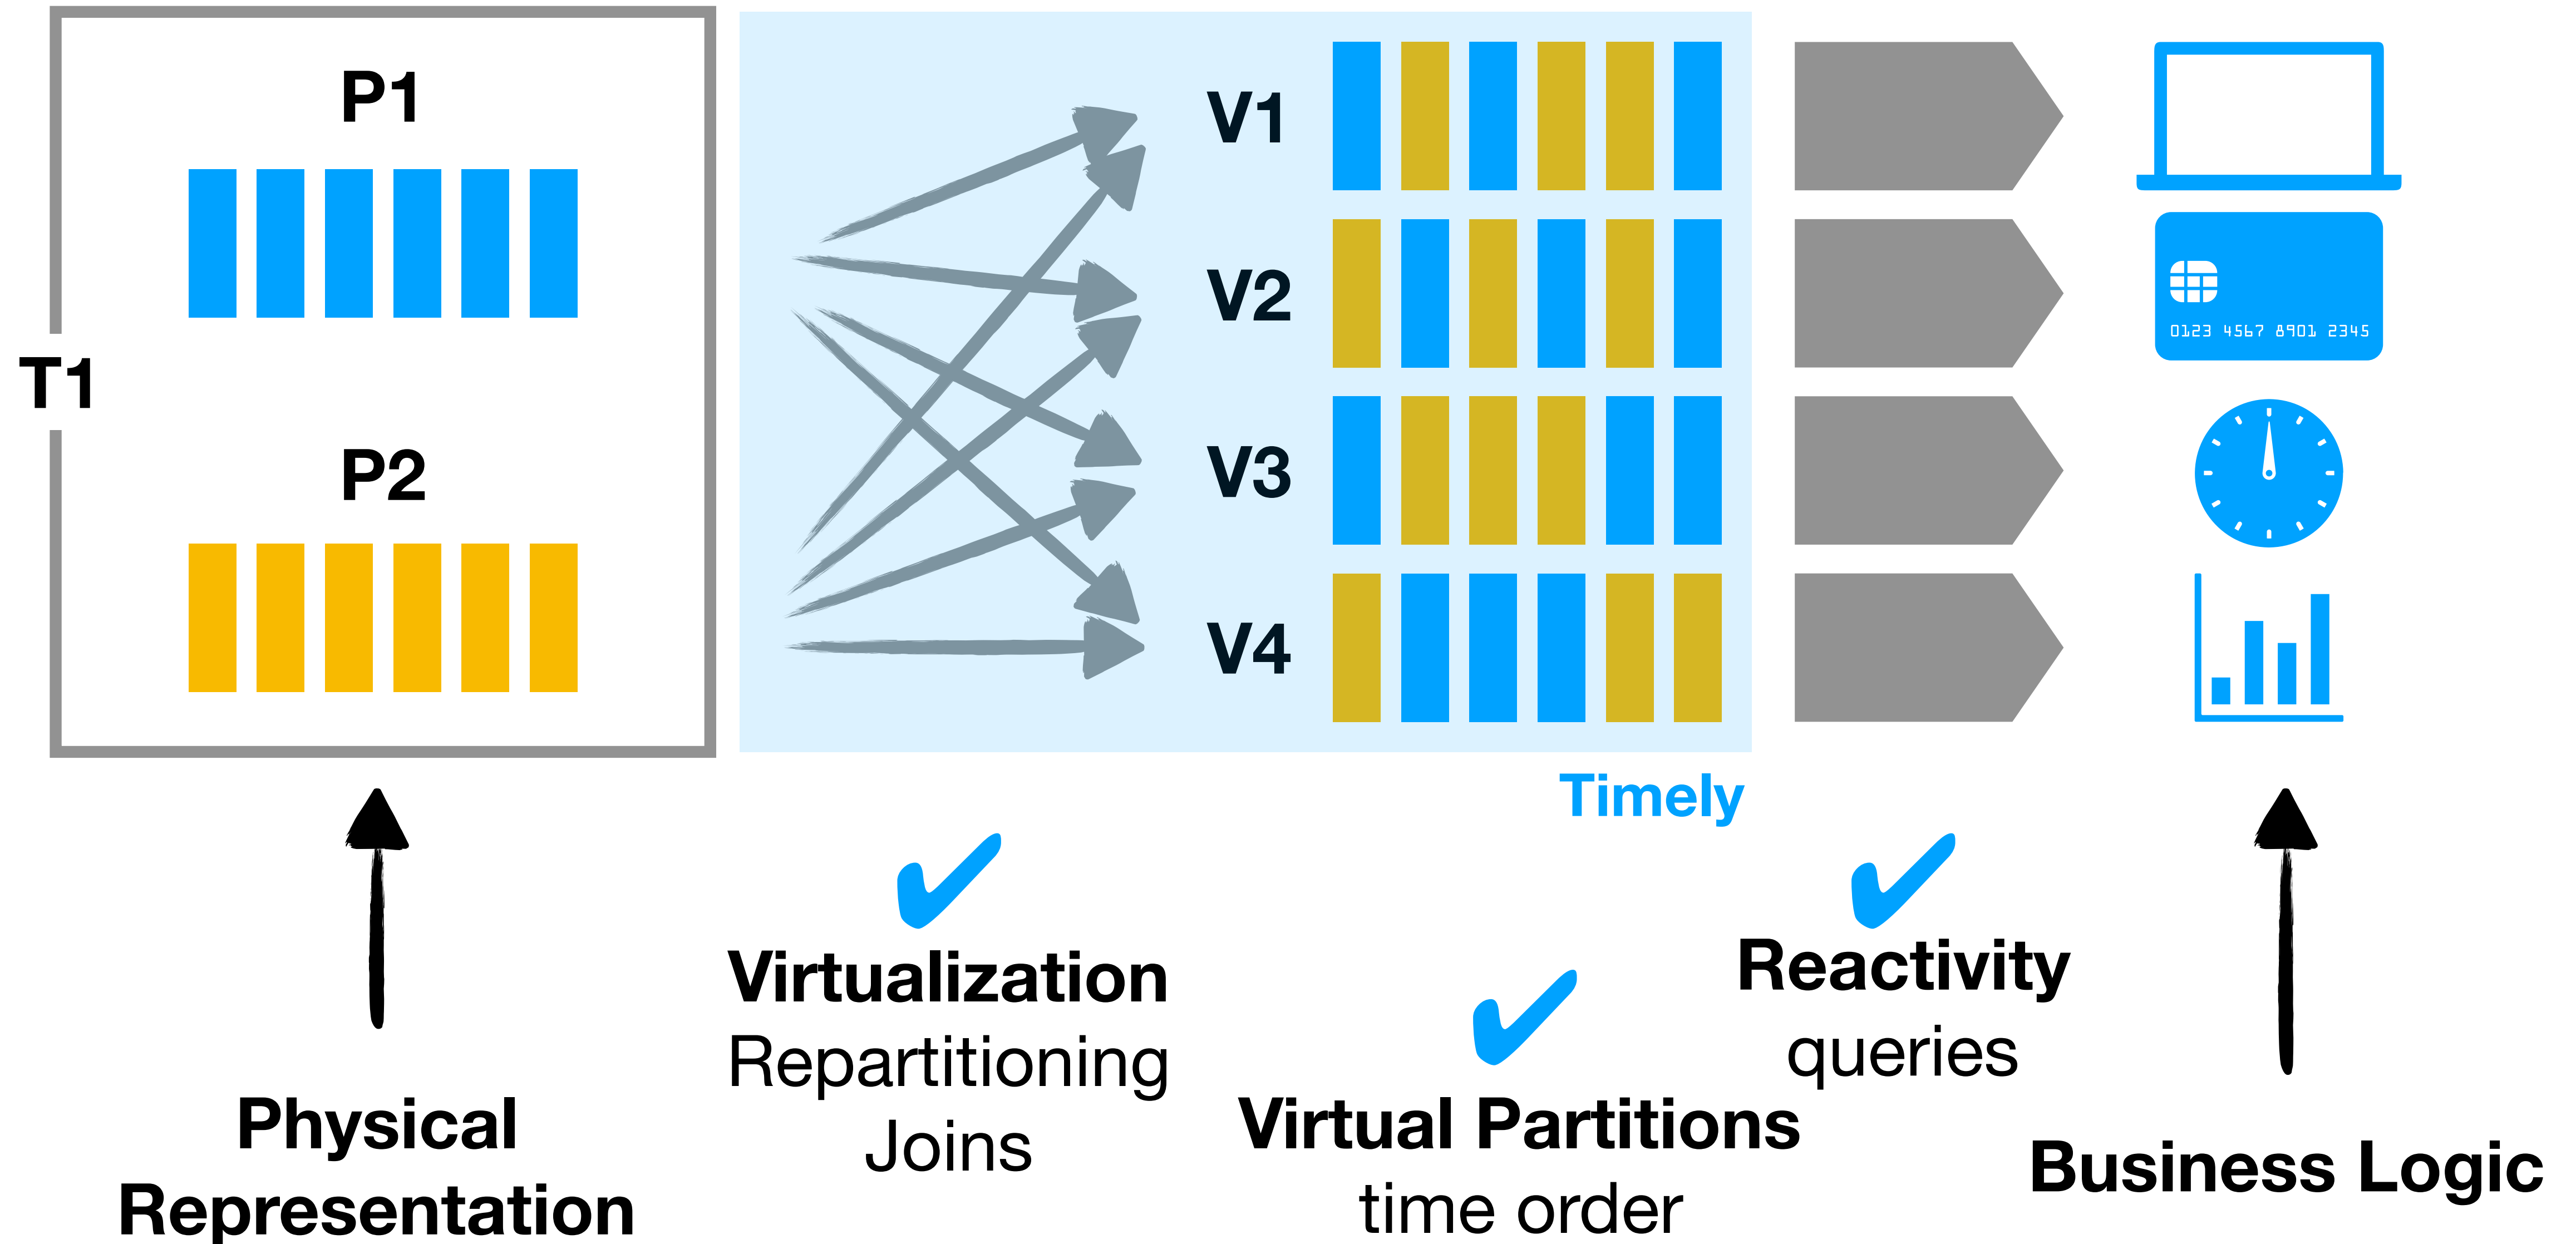
```
[[(bfs ?from ?to)
  [?from :edge ?to]]

[(bfs ?from ?to)
  [?from :edge ?hop]
  (bfs ?hop ?to)]]
```

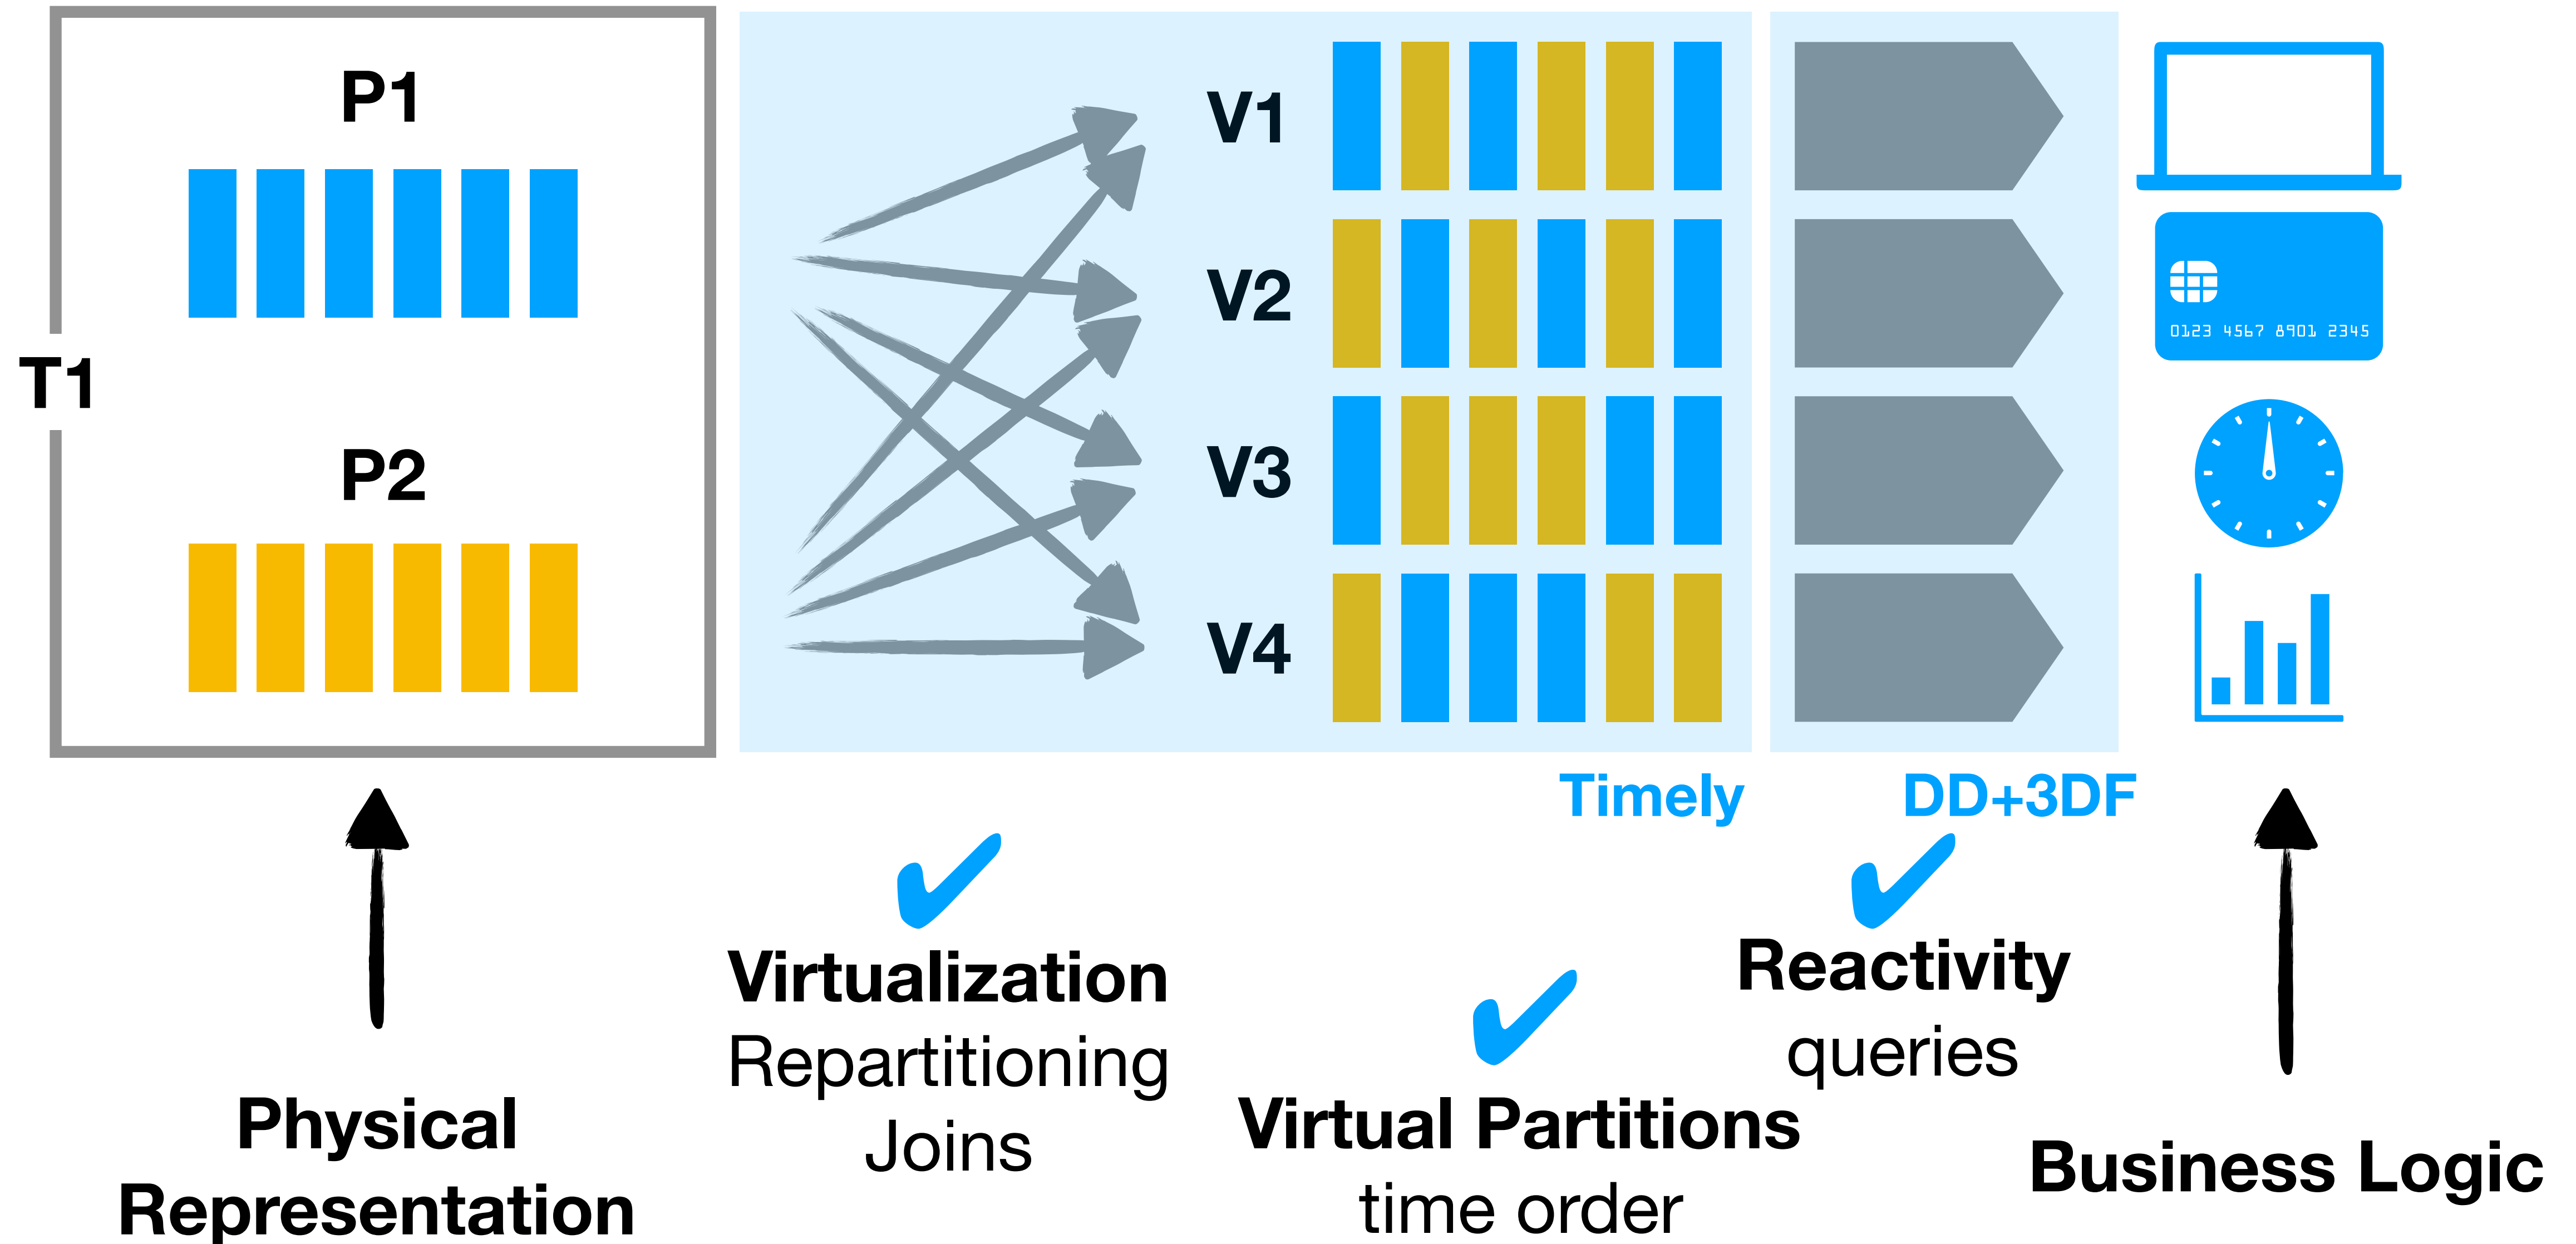
The Trifecta!



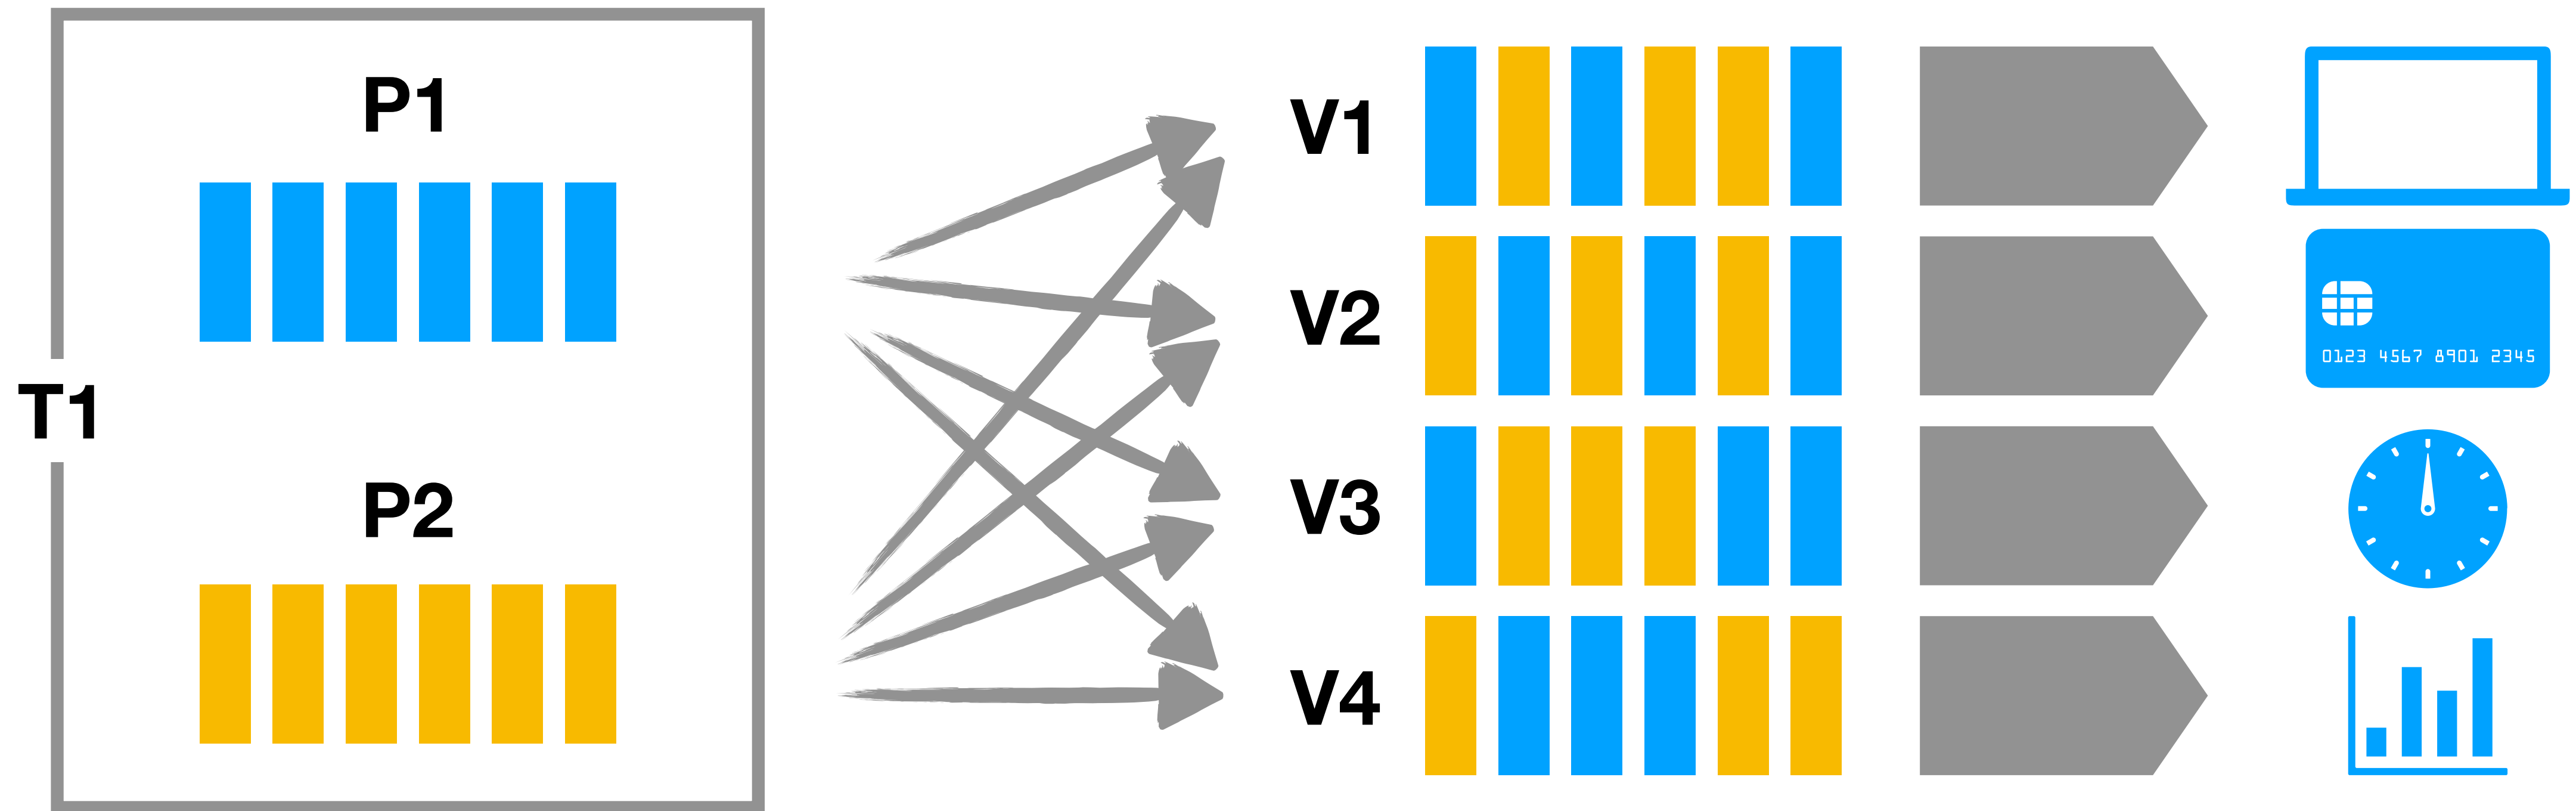
Kafka Superpowers



Kafka Superpowers



Kafka Superpowers



clockworks.io/kplex

Timely as a Programming Model

3DF

(Streaming Relational Queries)

Differential Dataflow

(Iterative Incrementalized Operators)

Timely Dataflow

(Dataflows w/ Multidimensional Progress Tracking)

Sources



clockworks

www.clockworks.io

{david, malte, moritz, niko}@clockworks.io

Repositories

- Timely: github.com/TimelyDataflow
- ST2: github.com/li1/snailtrail
- 3DF: github.com/comnik/declarative-dataflow
- Differential FAQ: github.com/eoxxs/differential-aggregate-query

Papers

- Naiad (Timely Dataflow): <http://dl.acm.org/citation.cfm?doid=2517349.2522738>
- Differential Dataflow: <http://michaelisard.com/pubs/differentialdataflow.pdf>, arxiv.org/abs/1812.02639
- SnailTrail: hdl.handle.net/20.500.11850/228581

Talks

- Reactive Datalog for Datomic (clojure/conj 2018): clockworks.io/2018/12/01/conj-talk.html
- Across Time and Space (BobKonf 2019): clockworks.io/2019/03/22/across-time-space.html

Blog Posts

- frankmcsherry.org
- Incremental Functional Aggregate Queries: clockworks.io/2019/07/06/Incremental-Functional-Aggregate-Queries.html
- Dataflows you can't refuse: clockworks.io/2019/02/10/dataflows-you-cant-refuse.html
- Reactive Datalog with Vega: clockworks.io/2018/11/25/reactive-datalog-with-vega.html
- Incremental Datalog with Differential Dataflows: clockworks.io/2018/09/13/incremental-datalog.html