


# Out of the Tar Pit

by Ben Moseley and Peter Marks

Nikolas Göbel

niko@clockworks.io

**ETH** zürich



Systems Group



**All problems with software stem from  
our limited ability to understand systems.**

**Complexity destroys  
the ability to understand a system.**

**Complexity :=**

*Number of things we have to keep in  
mind concurrently.*

# Methods of Understanding

black-box testing

informal reasoning



more errors found

less errors created

**State** := *What is the case?*

**Control Logic** := *When is it the case?*

**Behaviour** := *State over time.*

# Complexity Causes

- (1) possible states ~ scenarios to consider
- (2) control logic ~ interactions to consider


# 2nd Order Complexity

- State contaminates whatever it touches
- Over-specified control causes mismatch w/ runtime
- In general: all language powers will seep in eventually




**(1)**


**Essential vs Accidental**




**Problem Domain**  
*(Essential Complexity)*





**Possible Behaviours**  
*(Accidental Complexity)*













**(2)**

**Reducing Accidental State**

# Immutability



De-couples state from behaviour



Control is still a problem



Doesn't reduce the amount of state

# Functions



Referential transparency



No hidden state, everything visible in parameters



Control is still a problem



How much state can we stuff into parameters?

# Transactions



All the benefits of functions



Atomicity allows us to ignore intermediate states



All state pushed into a parameter (the database)



Control is still a problem

# Derived State (1)

```
{(x :mail/read-at 2019-01-01)
 (x :mail/read? true)
 (y :mail/read-at 2019-02-02)
 (y :mail/read? true)
 (user :unread-mails 2)}
```

# Derived State (2)

`(x :mail/read-at 2019-01-01) -> (x :mail/read? true)`



**essential**



**derived**

# Derived State (2)

```
{(x :mail/read-at 2019-01-01)  
 (y :mail/read-at 2019-02-02)  
 (user :unread-mails 2)}
```

```
(defn is-read? [state id]  
  (contains? state (id :mail/read-at? _)))
```

# Derived State (3)

`{x | (x :mail/read? true)} -> (user :unread-mails <count>)`



derived



derived



# Derived State (4)

```
{(x :mail/read-at 2019-01-01)  
 (y :mail/read-at 2019-02-02)}
```

```
(defn unread-count [state id]  
  (->> (user-mails state id)  
        (remove is-read?)  
        (count)))
```

# Reify Underlying Records (1)

```
{(player :position/x 5)  
 (player :position/y 4)}
```

# Reify Underlying Records (2)

```
[(player :move/up)  
 (player :move/up)  
 (player :move/right)  
 (player :move/up)]
```

Reification can expose new functional dependencies.

# Object-Orientation?

- squinting a bit, methods are per-object transactions
- problem 1: cross-object transactions
- problem 2: object identity

# Out Of The State Pit

- (1) Use immutable state wherever possible.
- (2) Model essential state in relations.
- (3) From those, derive all other state as views.

Note: Some more distinctions in the paper.

*Essential derived vs accidental derived*, mutable derived vs immutable derived.

**(3)**

**Reducing Accidental Control**

# Overspecified Control

Most languages force us to deal with control (e.g. sequencing and branching).



Impedance mismatches throughout the stack



Concurrency and distribution. Please no.

# Logic Programming



Declarative: *what?* not *how?*



Ordering follows entirely from data dependencies



Programs fully reversible out-of-the-box



Not really declarative in practice (e.g. Prolog)



No efficient implementations



# Relational Model



Declarative views: join, select, project, ...



Data logically independent of physical layout




Efficient and extensible



Programs only reversible w.r.t references

# Language Fragments



# Language Power

We are used to classifying time and space complexity.  
Classifying **descriptive complexity** and picking  
an appropriate language fragment can get us...



Control-free, declarative programs



Efficient, general-purpose implementations



Fine-grained concurrency and distribution for free

# Derived State – Revisited

```
{(x :mail/read-at 2019-01-01)  
 (y :mail/read-at 2019-02-02)}
```

```
[[ (mail/read? ?mail)  
   [?mail :mail/read-at _]]]
```

```
[(unread-count ?user ?count)  
 [?user :user/mail ?mail]  
 (not (mail/read? ?mail))  
 [(count ?mail) ?count]]]
```

# Out Of The Control Pit

- (1) Express derivations in appropriate language fragment.
- (2) Prefer dataflow to control flow.

**The relational model and its extensions provide a sweet-spot!**

**(4)**

**Feasibility**

# Out Of The Tar Pit

- (1) Use immutable state wherever possible.
- (2) Model essential state in relations.
- (3) From those, derive all other state as views.
- (4) Express derivations in appropriate language fragment.
- (5) Prefer dataflow to control flow.

# (1) Use immutable state wherever possible.

- first-class language support: Clojure, Elm
  - DataScript, ImmutableJS, Mori
  - append-only logs, Kafka
- ✓ Immutable data structures are efficient and proven in practice.
- ✓ Immutability as a concept is used on all levels of the stack.

(2) Model essential state in relations.

(3) From those, derive all other state as views.

(4) Express derivations in appropriate language fragment.

(5) Prefer dataflow to control flow.



(1) Use immutable state wherever possible.

(2) Model essential state in relations.

(3) From those, derive all other state as views.

- problem: consistent, incremental view maintenance

✓ works today, in the small: React + Redux

? goal: query power of DBs + performance of stream-processors

(4) Express derivations in appropriate language fragment.

(5) Prefer dataflow to control flow.

# Differential Dataflow

An expressive programming framework that quickly updates computations when its inputs change.

- describe functional-relational dataflow computations (map, reduce, join, iterate, filter, ...)
- efficient, incrementalized execution as inputs change
- distributable out-of-the-box (laptop -> cluster)

# 3DF (Declarative Differential Dataflow)

- describe your views in Datalog
- views are incrementally maintained
- feeds from durable data sources

[github.com/comnik/clj-3df](https://github.com/comnik/clj-3df)

[github.com/comnik/declarative-dataflow](https://github.com/comnik/declarative-dataflow)

- (1) Use immutable state wherever possible.
- (2) Model essential state in relations.
- (3) From those, derive all other state as views.
- (4) Express derivations in appropriate language fragment.
- (5) Prefer dataflow to control flow.

- previous attempts: tuple stores, coordination languages
- lots of R&D dollars to build on: databases, query optimization, ...
- CALM work by Hellerstein and Alvaro
- relational extensions: probabilistic, temporal, constraint solving, ...

 It's all there for the taking!

# 2nd Order Benefits



History and Analytics are first-class



Concurrency and distribution out-of-the-box



Tracing and Provenance, Meta-Queries



Efficient runtimes, coordination free distribution

# Adoption?

- Log-based architectures, event-sourcing
- Stream Processing
- Web Frontends, real-time sync front $\leftrightarrow$ back
- Distributed correctness, formal verification

# Into The Tar Pit

- State
- Conditionals
- Objects and Methods
- Classes and Inheritance
- Pattern matching
- Actors, Futures, Async/Await, Promises, ...
- ORMs
- ...

# Out Of The Tar Pit

- (1) Use immutable state wherever possible.
- (2) Model essential state in relations.
- (3) From those, derive all other state as views.
- (4) Express derivations in appropriate language fragment.
- (5) Prefer dataflow to control flow.



# Pointers

- “Simple Made Easy”, by Rich Hickey
- “I See What You Mean”, by Peter Alvaro
- [github.com/TimelyDataflow](https://github.com/TimelyDataflow)
- [github.com/comnik/declarative-dataflow](https://github.com/comnik/declarative-dataflow)
- [github.com/tonsky/datascript](https://github.com/tonsky/datascript)
- “The Declarative Imperative”, by Joseph Hellerstein