

Reactive Datalog for Datomic

by Nikolas Göbel

**“We want to make reactive systems that don't poll.
And we want those systems to get a consistent
view of the world.”**

– Rich Hickey, “Deconstructing the Database”

Nikolas Göbel

niko@clockworks.io

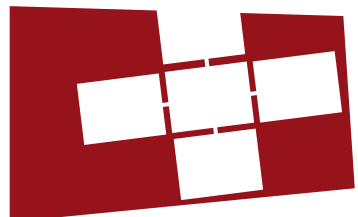
in collaboration with:

Frank McSherry (ETH)

David Bach (Clockworks)

Malte Sandstede (Clockworks)

ETH zürich



Systems Group



Reactive Systems?

- live-updating web applications
- alerting and real-time dashboards
- stream processors
- rule engines

Reactive Systems?

The arrival of **data coordinates code**, not the other way around.



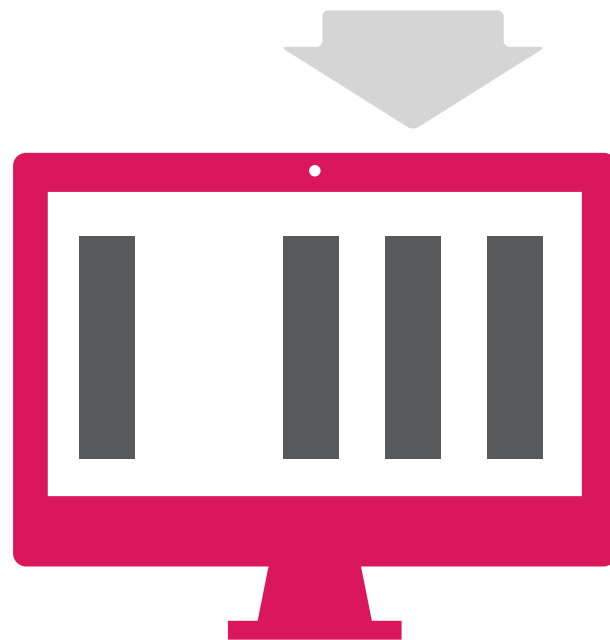
Source of Truth



Accessible Data

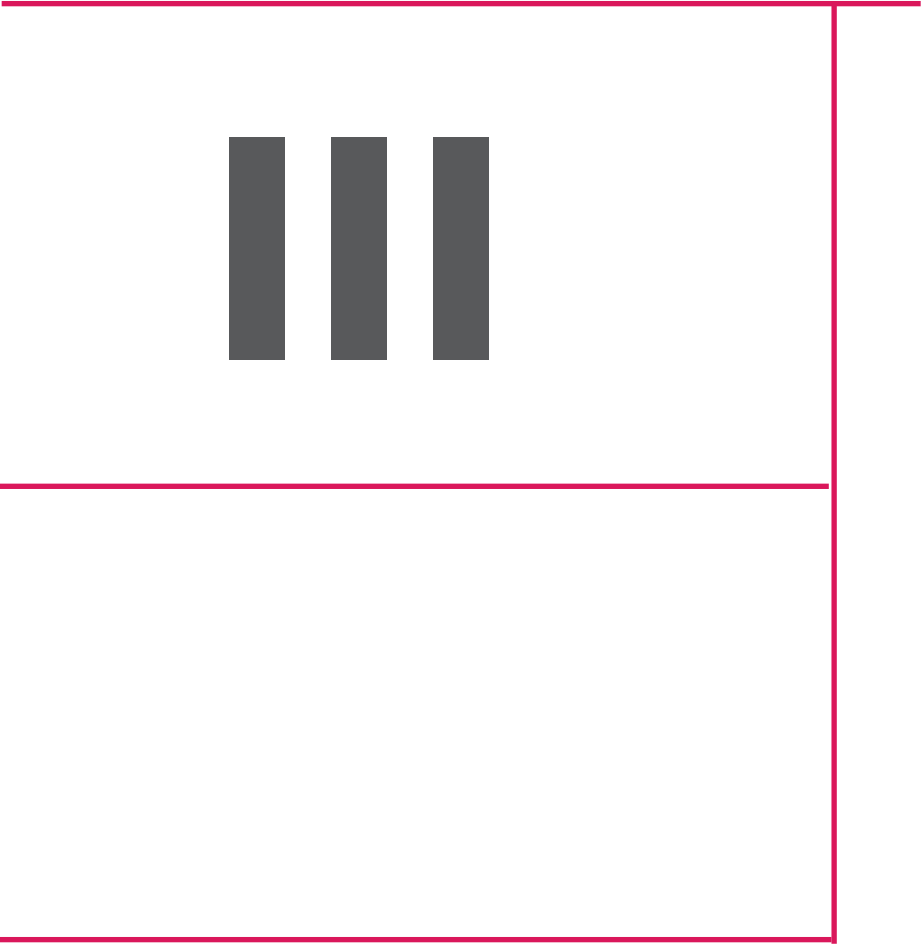


Active View



Data Coordinates Code

- “run a query” => “subscribe to a query”
- re-run to incorporate changes => changes propagate



Datalog



:person/...

:loan/...

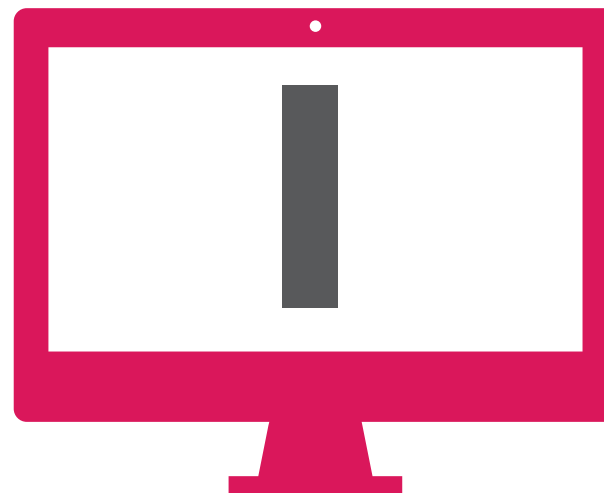
:branch/...



conjbank/opt-ins



conjbank/debts



<session>/receivables

3DF

- reactive Datalog evaluation on a distributed, data-parallel stream processor
- feeds from durable data sources (Datomic!)
- propagates only changes

github.com/comnik/clj-3df

github.com/comnik/declarative-dataflow

transact

SOURCE OF TRUTH



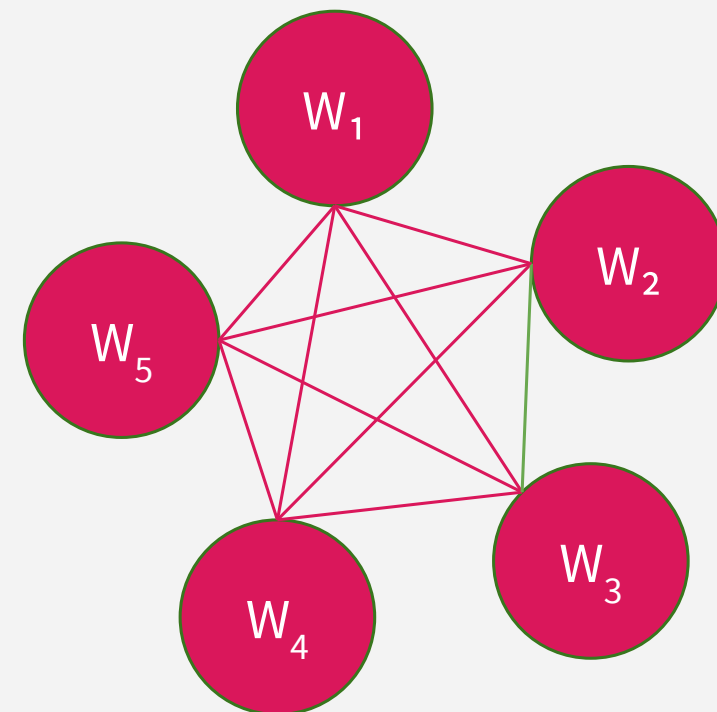
read tx-log

Datalog

Query Plan

register

push output changes



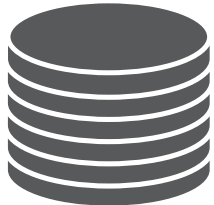
Client / Consumer

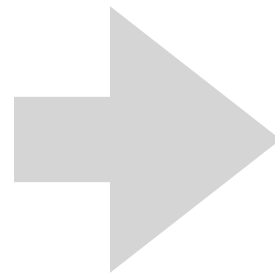
3DF

Requirements

- data coordinates code: work like a stream processor
- transparent: notion of change should not leak into Datalog
- reactive iteration w/ arbitrary retractions (recursion!)

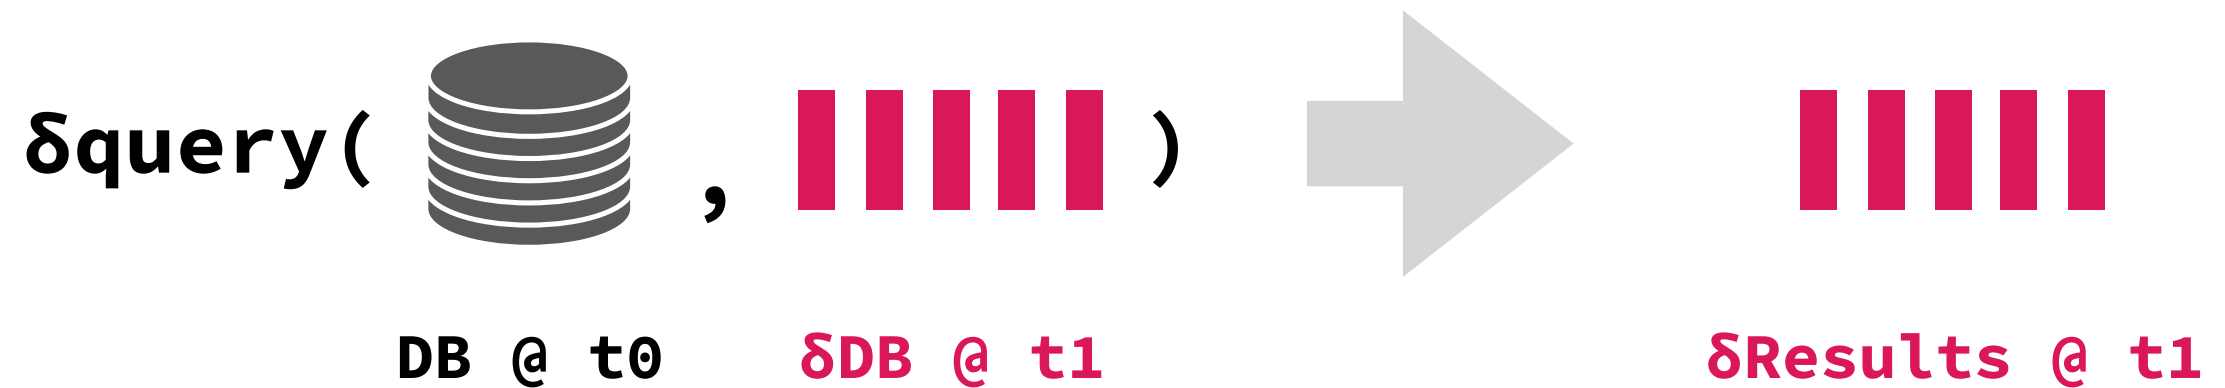
Computing w/ Collections

query( **)**
DB @ t0




Results @ t0

Incremental Computation



Naive δ query



```
[:find ?device
```

```
  :where
```

```
    [?device :settings/speed ?target]
```

```
    [?device :device/speed ?speed]
```

```
    [(< ?speed ?target)]]
```



```
[ :find ?device
```

```
  :in $ ?tx
```

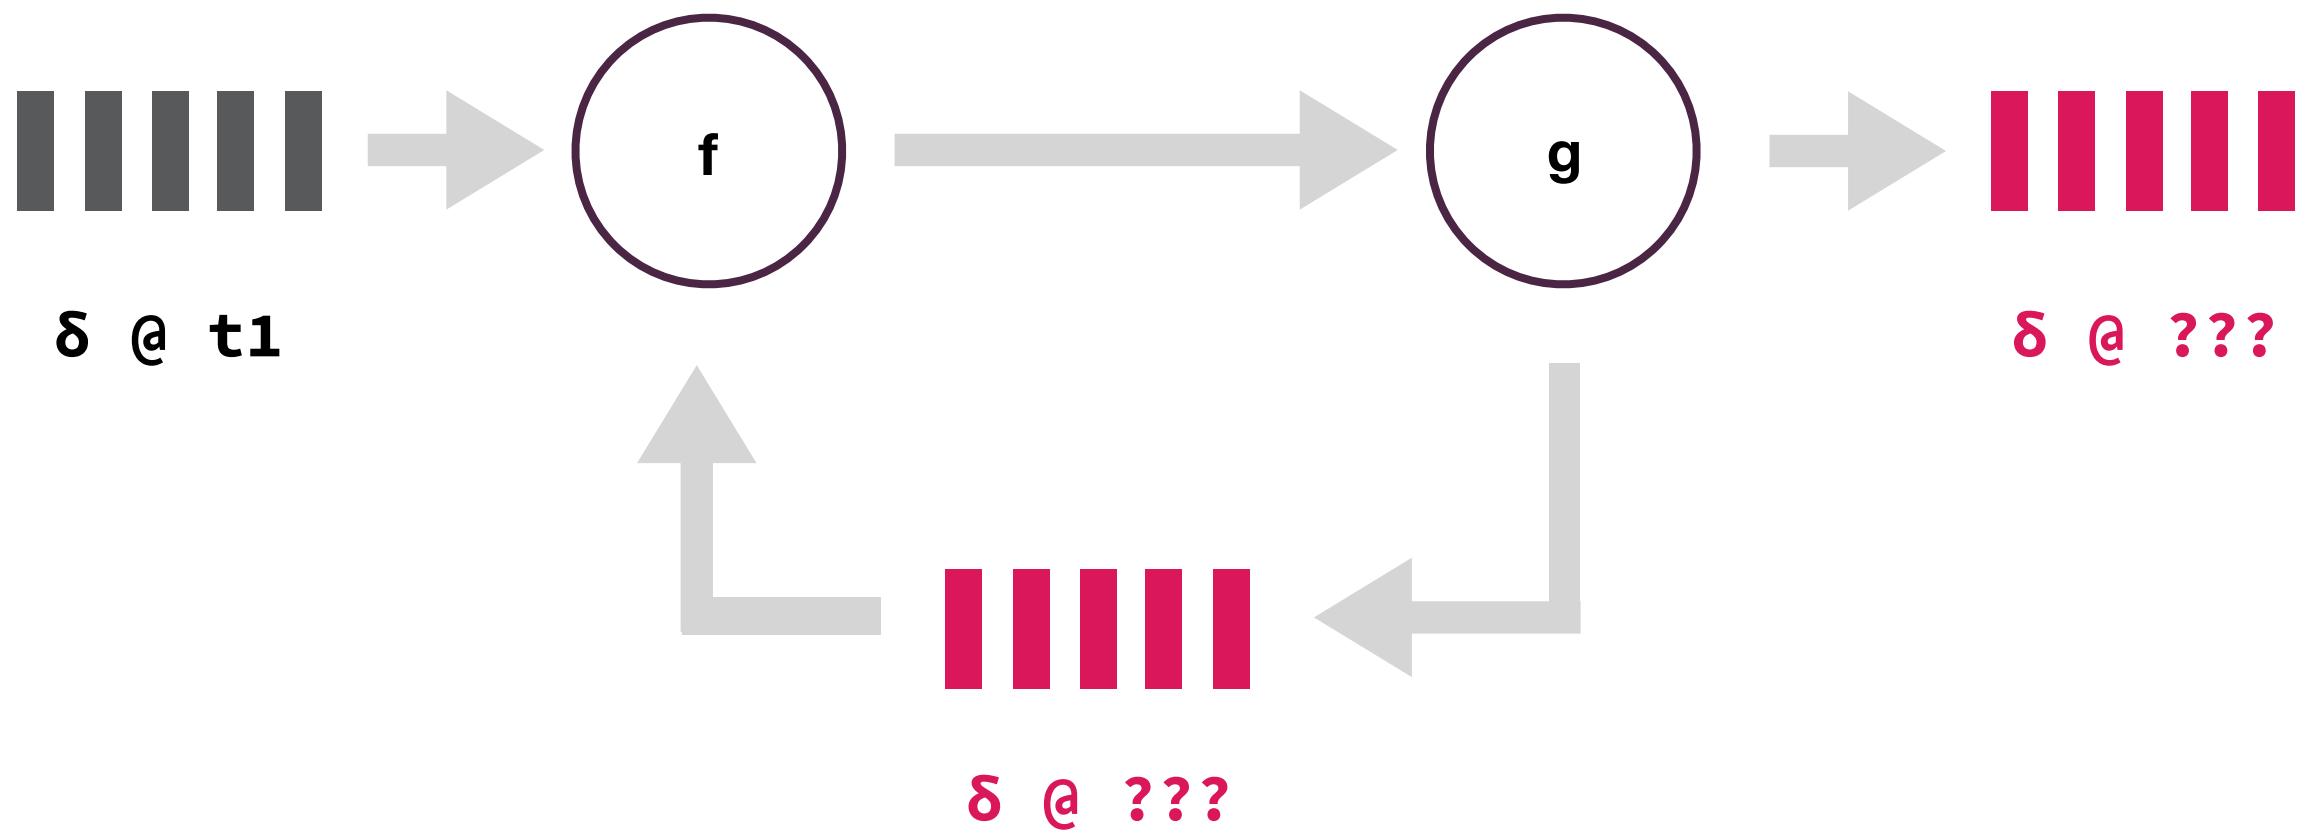
```
  :where
```

```
    [$ ?device :settings/speed ?target]
```

```
    [ ?tx ?device :device/speed ?speed]
```

```
    [( < ?speed ?target)]]
```

Reactive Iteration?



Differential Dataflow

“a data-parallel programming framework designed to **quickly respond to arbitrary changes** in input”

github.com/frankmcsherry/differential-dataflow

frankmcsherry.org

github.com/frankmcsherry/timely-dataflow

Persistent Collections

:edge

(a b) +1 t0	
(c d) +1 t0	← (#{[a b] [c d]} t0)
(c d) -1 t1	
(c e) +1 t1	← (#{[a b] [c e]} t1)

Partially-ordered Times

(bfs :edge a)

(a b) +1 t0		
(c d) +1 t0		
(b c) +1 t1	(a c) +1 t1.1	(a d) +1 t1.2
		(c d) -1 t2

Partially-ordered Times

(bfs :edge a)

(a b) +1 (t0 0)		
(c d) +1 (t0 0)		
(b c) +1 (t1 0)	(a c) +1 (t1 1)	(a d) +1 (t1 2)
(c d) -1 (t2 0)		

Partially-ordered Times

(bfs :edge a)

(a b) +1 (t0 0)		
(c d) +1 (t0 0)		
(b c) +1 (t1 0)	(a c) +1 (t1 1)	(a d) +1 (t1 2)
(c d) -1 (t2 0)		

Partially-ordered Times

(bfs :edge a)

(a b) +1 (t0 0)		
(c d) +1 (t0 0)		
(b c) +1 (t1 0)	(a c) +1 (t1 1)	(a d) +1 (t1 2)
(c d) -1 (t2 0)		

Partially-ordered Times

(bfs :edge a)

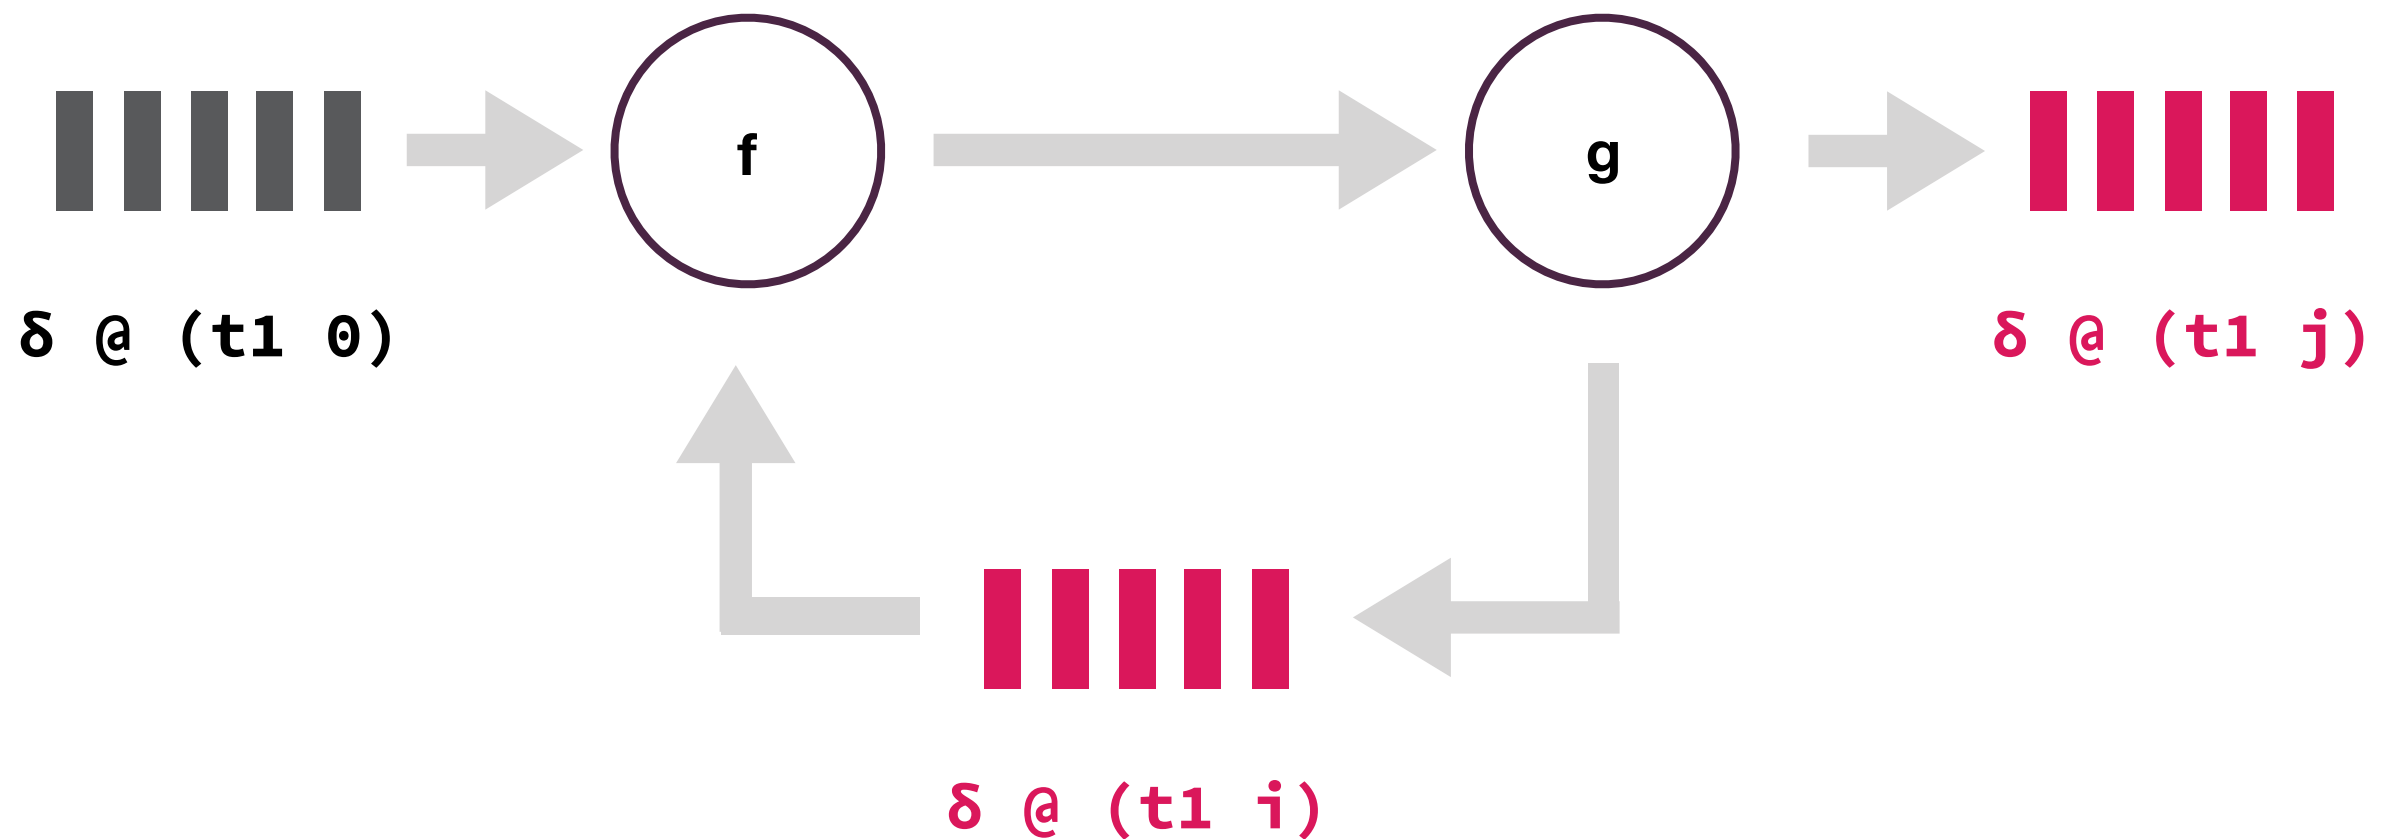
(a b) +1 (t0 0)		
(c d) +1 (t0 0)		
(b c) +1 (t1 0)	(a c) +1 (t1 1)	(a d) +1 (t1 2)
(c d) -1 (t2 0)	-	

Partially-ordered Times

(bfs :edge a)

(a b) +1 (t0 0)		
(c d) +1 (t0 0)		
(b c) +1 (t1 0)	(a c) +1 (t1 1)	(a d) +1 (t1 2)
(c d) -1 (t2 0)	-	(a d) -1 (t2 2)

Reactive Iteration!



```
/// BFS
```

```
let nodes = roots.map(|x| (x, 0));
```

```
nodes.iterate(|inner| {
```

```
    let edges = edges.enter(&inner.scope());
```

```
    let nodes = nodes.enter(&inner.scope());
```

```
    inner.join_map(&edges, |_k,l,d| (*d, l+1))
```

```
        concat(&nodes)
```

```
        group(|_, s, t| t.push((*s[0].0, 1)))
```

```
})
```

Reactive Systems!

- ~~live-updating web applications~~
- alerting and real-time dashboards
- stream processors
- rule engines



Loans



Defaults



Alerting



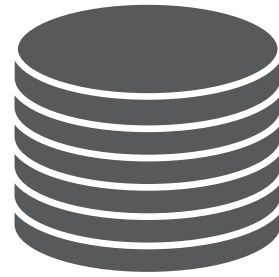
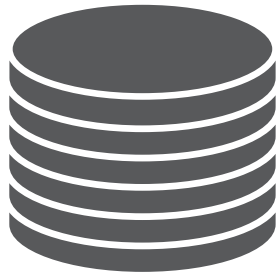
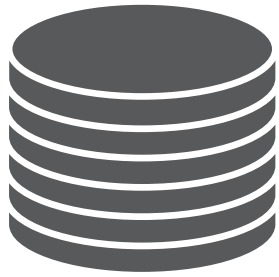
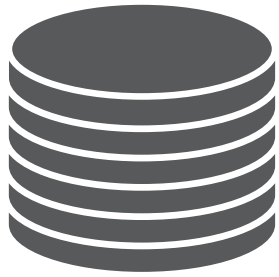
User balance



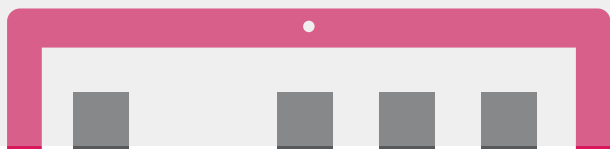
Overdue



Admin Dashboard



**Databases /
Services /
Filesystems**



Notes on Performance

- incremental computation pegs system load to rate of change, rather than to granularity \times # of users
- queries not limited to a single peer
- timestamps support nanosecond precision (no windows!)
- drop down to hand-built Differential Dataflow if needed

Future Work

- scalability and performance isolation for large # of queries
- bullet-proofing the drop-in Datomic integration
- Beta release

**“We want to make reactive systems that don't poll.
And we want those systems to get a consistent
view of the world.”**

– Rich Hickey, “Deconstructing the Database”

3DF lets you...

- ...leverage Datalog to build reactive systems and distributed stream-processing jobs
- ...scale queries beyond the limits of a single peer
- ...continue using Datomic for everything else
- ...integrate with hand-written Differential Dataflow

Interested? Let's talk!

@NikolasGoebel | niko@clockworks.io